



QDMA Subsystem for PCI Express Product Guide (PG302)

Introduction

- Features
- IP Facts

Overview

- QDMA Architecture
- Limitations
- Applications
- Licensing and Ordering

Product Specification

- Standards
- Performance and Resource Utilization
- Minimum Device Requirements
- QDMA Operations
- Port Descriptions
- Register Space

Designing with the Subsystem

- General Design Guidelines
- Clocking
- Tandem Configuration

Design Flow Steps

- Customizing and Generating the Subsystem
- Constraining the Subsystem
- Simulation
- Synthesis and Implementation

Example Design

- Available Example Designs
- Example Design Registers
- Customizing and Generating the Example Design

Test Bench

- Architecture
- Scaled Simulation Timeouts
- Test Selection
- Waveform Dumping
- Output Logging
- Test Description
- Model Task List

Upgrading

- Changes from v3.0 to v5.1
- Comparing With DMA/Bridge Subsystem for PCI Express

GT Locations

Debugging

- Finding Help with AMD Adaptive Computing Solutions
- Debug Tools
- Hardware Debug

Application Software Development

- Device Drivers
- Linux QDMA Software Architecture (PF/VF)
- Using the Drivers
- Reference Software Driver Flow

Additional Resources and Legal Notices

- Finding Additional Documentation
- Support Resources
- References
- Revision History
- Please Read: Important Legal Notices

Introduction

The AMD QDMA Subsystem for PCI Express (PCIe®) implements a high performance DMA for use with the PCI Express® Integrated Block with the concept of multiple queues that is different from the DMA/Bridge Subsystem for PCI Express which uses multiple AMD Card to Host (C2H) and Host to Card (H2C) channels.

This guide covers the QDMA Subsystem for PCIe® for AMD UltraScale+™ devices.

!! Important: In this guide, AMD UltraScale+™ refers to both the existing UltraScale+ and the new Spartan UltraScale+ devices.

Features

- The PCIe Integrated Block is supported in AMD UltraScale+™ devices.
- Supports 64, 128, 256, and 512-bit data path.
- Supports x1, x2, x4, x8, or x16 link widths.
- Supports Gen1, Gen2, and Gen3 link speeds. Gen4 for PCIE4C and PCIE4CE block.
- Support for both the AXI4 Memory Mapped and AXI4-Stream interfaces per queue.
- 2048 queue sets
 - 2048 H2C descriptor rings.
 - 2048 C2H descriptor rings.
 - 2048 C2H Completion (CMPT) rings.
- Supports Polling Mode (Status Descriptor Write Back) and Interrupt Mode.
- Interrupts
 - 2048 MSI-X vectors.
 - Up to 8 MSI-X per function.
 - Interrupt aggregation.
- C2H Stream interrupt moderation.
- C2H Stream Completion queue entry coalescence.
- Descriptor and DMA customization through user logic
 - Allows custom descriptor format.
 - Traffic Management.
- Supports SR-IOV with up to 4 Physical Functions (PF) and 252 Virtual Functions (VF)
 - Thin hypervisor model.
 - QID virtualization.
 - Allows only privileged/Physical functions to program contexts and registers.
 - Function level reset (FLR) support.
 - Mailbox.
- Rich programmability on a per queue basis, such as AXI4 Memory Mapped versus AXI4-Stream interfaces.

IP Facts

LogiCORE IP Facts Table	
Subsystem Specifics	
Supported Device Family ¹	AMD UltraScale+™ , AMD Spartan™ UltraScale+™
Supported User Interfaces	AXI4 Memory Map, AXI4-Stream, AXI4-Lite
Resources	Resource Use web page.
Subsystem	
Design Files	Encrypted System Verilog
Example Design	Verilog
Test Bench	Verilog
Constraints File	Xilinx Constraints File (XDC)
Simulation Model	Verilog
Supported S/W Driver	Linux, DPDK, and Windows Drivers ²
Tested Design Flows ³	
Design Entry	AMD Vivado™ Design Suite
Simulation	For supported simulators, see the <i>Vivado Design Suite User Guide: Release Notes, Installation, and Licensing</i> (UG973).
Synthesis	Vivado Synthesis
Support	
Release Notes and Known Issues	Master Answer Record: 70927
All Vivado IP Change Logs	Master Vivado IP Change Logs: 72775
Support web page	
<ol style="list-style-type: none"> 1. For a complete list of supported devices, see the Vivado IP catalog. 2. For driver details, see Xilinx DMA IP Drivers. 3. For the supported versions of the tools, see the <i>Vivado Design Suite User Guide: Release Notes, Installation, and Licensing</i> (UG973). 	

Overview

The Queue-based Direct Memory Access (QDMA) subsystem is a PCI Express® (PCIe®) based DMA engine that is optimized for both high bandwidth and high packet count data transfers. The QDMA is composed of the AMD UltraScale+™ Integrated Block for PCI Express, and an extensive DMA and bridge infrastructure that enables the ultimate in performance and flexibility.

The QDMA Subsystem for PCIe offers a wide range of setup and use options, many selectable on a per-queue basis, such as memory-mapped DMA or stream DMA, interrupt mode and polling.

The subsystem provides many options for customizing the descriptor and DMA through user logic to provide complex traffic management capabilities.

The primary mechanism to transfer data using the QDMA is for the QDMA engine to operate on instructions (descriptors) provided by the host operating system. Using the descriptors, the QDMA can move data in both the Host to Card (H2C) direction, or the Card to Host (C2H) direction. You can select on a per-queue basis whether DMA traffic goes to an AXI4 memory map (MM) interface or to an AXI4-Stream interface. In addition, the QDMA has the option to implement both an AXI4 MM Master port and an AXI4 MM Slave port, allowing PCIe traffic to bypass the DMA engine completely.

The main difference between QDMA and other DMA offerings is the concept of queues. The idea of queues is derived from the “queue set” concepts of Remote Direct Memory Access (RDMA) from high performance computing (HPC) interconnects. These queues can be individually configured by interface type, and they function in many different modes. Based on how the DMA descriptors are loaded for a single queue, each queue provides a very low overhead option for setup and continuous update functionality. By assigning queues as resources to multiple PCIe Physical Functions (PFs) and Virtual Functions (VFs), a single QDMA core and PCI Express interface can be used across a wide variety of multifunction and virtualized application spaces. The QDMA Subsystem for PCIe can be used and exercised with an AMD provided QDMA reference driver, and then built out to meet a variety of application spaces.


Related Information

[Port Descriptions](#)

QDMA Architecture

The following figure shows the block diagram of the QDMA Subsystem for PCIe.

Figure: QDMA Architecture



DMA Engines

Descriptor Engine

The Host to Card (H2C) and Card to Host (C2H) descriptors are fetched by the Descriptor Engine in one of two modes: Internal mode, and Descriptor bypass mode. The descriptor engine maintains per queue contexts where it tracks software (SW) producer index pointer (PIDX), consumer index pointer (CIDX), base address of the queue (BADDR), and queue configurations for each queue. The descriptor engine uses a round robin algorithm for fetching the descriptors. The descriptor engine has separate buffers for H2C and C2H queues, and ensures it never fetches more descriptors than available space. The descriptor engine will have only one DMA read outstanding per queue at a time and can read as many descriptors as can fit in a MRSS. The descriptor engine is responsible for reordering the out of order completions and ensures that descriptors for queues are always in order.

The descriptor bypass can be enabled on a per-queue basis and the fetched descriptors, after buffering, are sent to the respective bypass output interface instead of directly to the H2C or C2H engine. In internal mode, based on the context settings the descriptors are sent to delete per H2C memory mapped (MM), C2H MM, H2C Stream, or C2H Stream engines.

The descriptor engine is also responsible for generating the status descriptor for the completion of the DMA operations. With the exception of C2H Stream mode, all modes use this mechanism to convey completion of each DMA operation so that software can reclaim descriptors and free up any associated buffers. This is indicated by the CIDX field of the status descriptor.

Recommended: If a queue is associated with interrupt aggregation, AMD recommends that

the status descriptor be turned off, and instead the DMA status be received from the interrupt aggregation ring.

To put a limit on the number of fetched descriptors (for example, to limit the amount of buffering required to store the descriptor), it is possible to turn-on and throttle credit on a per-queue basis. In this mode, the descriptor engine fetches the descriptors up to available credit, and the total number of descriptors fetched per queue is limited to the credit provided. The user logic can return the credit through the `dsc_crdt` interface. The credit is in the granularity of the size of the descriptor.

To help a user-developed traffic manager prioritize the workload, the available descriptor to be fetched (incremental PIDX value) of the PIDX update is sent to the user logic on the `tm_dsc_sts` interface. Using this interface it is possible to implement a design that can prioritize and optimize the descriptor storage.

Related Information

[Interrupt Aggregation Ring](#)

H2C MM Engine

The H2C MM Engine moves data from the host memory to card memory through the H2C AXI-MM interface. The engine generates reads on PCIe, splitting descriptors into multiple read requests based on the MRRS and the requirement that PCIe reads do not cross 4 KB boundaries. Once completion data for a read request is received, an AXI write is generated on the H2C AXI-MM interface. For source and destination addresses that are not aligned, the hardware will shift the data and split writes on AXI-MM to prevent 4 KB boundary crossing. Each completed descriptor is checked to determine whether a writeback and/or interrupt is required.

For Internal mode, the descriptor engine delivers memory mapped descriptors straight to the H2C MM engine. The user logic can also inject the descriptor into the H2C descriptor bypass interface to move data from host to card memory. This gives the ability to do interesting things such as mixing control and DMA commands in the same queue. Control information can be sent to a control processor indicating the completion of DMA operation.

C2H MM Engine

The C2H MM Engine moves data from card memory to host memory through the C2H AXI-MM interface. The engine generates AXI reads on the C2H AXI-MM bus, splitting descriptors into multiple requests based on 4 KB boundaries. Once completion data for the read request is received on the AXI4 interface, a PCIe write is generated using the data from the AXI read as the contents of the write. For source and destination addresses that are not aligned, the hardware will shift the data and split writes on PCIe to obey Maximum Payload Size (MPS) and prevent 4 KB boundary crossings. Each completed descriptor is checked to determine whether a writeback and/or interrupt is required.

For Internal mode, the descriptor engine delivers memory mapped descriptors straight to the C2H MM engine. As with H2C MM Engine, the user logic can also inject the descriptor into the C2H descriptor bypass interface to move data from card to host memory.

For multi-function configuration support, the PCIe function number information will be provided in

the aruser bits of the AXI-MM interface bus to help virtualization of card memory by the user logic. A parity bus, separate from the data and user bus, is also provided for end-to-end parity support.

H2C Stream Engine

The H2C stream engine moves data from the host to the H2C Stream interface. For internal mode, descriptors are delivered straight to the H2C stream engine; for a queue in bypass mode, the descriptors can be reformatted and fed to the bypass input interface. The engine is responsible for breaking up DMA reads to MRRS size, guaranteeing the space for completions, and also makes sure completions are reordered to ensure H2C stream data is delivered to user logic in-order.

The engine has sufficient buffering for up to 256 descriptor reads and up to 32 KB of data. DMA fetches the data and aligns to the first byte to transfer on the AXI4 interface side. This allows every descriptor to have random offset and random length. The total length of all descriptors put together must be less than 64 KB.

For internal mode queues, each descriptor defines a single AXI4-Stream packet to be transferred to the H2C AXI-ST interface. A packet with multiple descriptors straddling is not allowed due to the lack of per queue storage. However, packets with multiple descriptors straddling can be implemented using the descriptor bypass mode. In this mode, the H2C DMA engine can be initiated when the user logic has enough descriptors to form a packet. The DMA engine is initiated by delivering the multiple descriptors straddled packet along with other H2C ST packet descriptors through the bypass interface, making sure they are not interleaved. Also, through the bypass interface, the user logic can control the generation of the status descriptor.

C2H Stream Engine

The C2H streaming engine is responsible for receiving data from the user logic and writing to the Host memory address provided by the C2H descriptor for a given Queue.

The C2H engine has two major blocks to accomplish C2H streaming DMA, Descriptor Prefetch Cache (PFCH), and the C2H-ST DMA Write Engine. The PFCH has per queue context to enhance the performance of its function and the software that is expected to program it.

PFCH cache has three main modes, on a per queue basis, called Simple Bypass Mode, Internal Cache Mode, and Cached Bypass Mode.

- In Simple Bypass Mode, the engine does not track anything for the queue, and the user logic can define its own method to receive descriptors. The user logic is then responsible for delivering the packet and associated descriptor through the simple bypass interface. The ordering of the descriptors fetched by a queue in the bypass interface and the C2H stream interface must be maintained across all queues in bypass mode.
- In Internal Cache Mode and Cached Bypass Mode, the PFCH module offers storage for up to 512 descriptors and these descriptors can be used by up to 64 different queues. In this mode, the engine controls the descriptors to be fetched by managing the C2H descriptor queue credit on demand based on received packets in the pipeline. Pre-fetch mode can be enabled on a per queue basis, and when enabled, causes the descriptors to be opportunistically pre-fetched so that descriptors are available before the packet data is available. The status can be found in prefetch context. This significantly reduces the latency by allowing packet data to be transferred to the PCIe integrated block almost immediately, instead of having to wait for the relevant descriptor to be fetched. The size of the data buffer is fixed for a queue (PFCH context) and the engine can scatter the packet across as many as seven descriptors. In cached bypass mode descriptor is bypassed to user logic for further processing, such as address translation, and sent back on the bypass interface. This mode does not assume any ordering descriptor and C2H stream packet interface, and the pre-fetch engine can match the packet and descriptors. When pre-fetch mode is enabled, do not give credits to IP. The pre-fetch engine takes care of credit management.

Completion Engine

The Completion (CMPT) Engine is used to write to the completion queues. Although the Completion Engine can be used with an AXI-MM interface and Stream DMA engines, the C2H Stream DMA engine is designed to work closely with the Completion Engine. The Completion Engine can also be used to pass immediate data to the Completion Ring. The Completion Engine can be used to write Completions of up to 64B in the Completion ring. When used with a DMA engine, the completion is used by the driver to determine how many bytes of data were transferred with every packet. This allows the driver to reclaim the descriptors.

The Completion Engine maintains the Completion Context. This context is programmed by the Driver and is maintained on a per-queue basis. The Completion Context stores information like the base address of the Completion Ring, PIDX, CIDX and a number of aspects of the Completion Engine, which can be controlled by setting the fields of the Completion Context.

The engine also can be configured on a per-queue basis to generate an interrupt or a completion status update, or both, based on the needs of the software. If the interrupts for multiple queues are aggregated into the interrupt aggregation ring, the status descriptor information is available in the interrupt aggregation ring as well.

The CMPT Engine has a cache of up to 64 entries to coalesce the multiple smaller CMPT writes into 64B writes to improve the PCIe efficiency. At any time, completions can be simultaneously coalesced for up to 64 queues. Beyond this, any additional queue that needs to write a CMPT entry causes the eviction of the least recently used queue from the cache. The depth of the cache is set to 64.

Bridge Interfaces

AXI Memory Mapped Bridge Master Interface

The AXI MM Bridge Master interface is used for high bandwidth access to AXI Memory Mapped space from the host. The interface supports up to 32 outstanding AXI reads and writes. One or more PCIe BAR of any physical function (PF) or virtual function (VF) can be mapped to the AXI-MM bridge master interface. This selection must be made prior to design compilation. The function ID, BAR ID, VF group, and VF group offset will be made available as part of `m_axib_awuser/m_axib_aruser[54:0]` user bits mapping is listed in AXI Bridge Master Ports.

Virtual function group (VFG) refers to the VF group number. It is equivalent to the PF number associated with the corresponding VF. VFG_OFFSET refers to the VF number with respect to a particular PF. Note that this is not the FIRST_VF_OFFSET of each PF.

For example, if both PF0 and PF1 have 8 VFs, FIRST_VF_OFFSET for PF0 and PF1 is 4 and 11. Below is the mapping for VFG and VFG_OFFSET.

Table: AXI-MM Interface Virtual Function Group

Function Number	PF Number	VFG	VFG_OFFSET
0	0	0	0
1	1	0	0
4	0	0	0 (Because FIRST_VF_OFFSET for PF0 is 4, the first VF of PF0 starts at FN_NUM=4 and VFG_OFFSET=0 indicates this is the first VF for PF0)
5	0	0	1 (VFG_OFFSET=1 indicates this is the second VF for PF0)
...
12	1	1	0 (VFG=1 indicates this VF is associated with PF1)
13	1	1	1

Each host initiated access can be uniquely mapped to the 64-bit AXI address space through the PCIe to AXI BAR translation.

Since all functions share the same AXI Master address space, a mechanism is needed to map requests from different functions to a distinct address space on the AXI master side. An example provided below shows how PCIe to AXI translation vector is used. Note that all VFs belonging to the same PF share the same PCIe to AXI translation vector. Therefore, the AXI address space of

each VF is concatenated together. Use VFG_OFFSET to calculate the actual starting address of AXI for a particular VF.

To summarize, AXI master write or read address is determined as:

- For PF, address = pcie2axi_vec + axib_offset.
- For VF, address = pcie2axi_vec + (VFG_OFFSET + 1)*vf_bar_size + axib_offset.

Where pcie2axi_vec is PCIe to AXI BAR translation (that can be set when the IP core is configured from the AMD Vivado™ IP catalog).

And axib_offset is the address offset in the requested target space.

AXI4-Lite Bridge Master Interface

One or more PCIe BAR of any physical function (PF) or virtual function (VF) can be mapped to the AXI4-Lite master interface. This selection must be done at the point of configuring the IP. The function ID, BAR ID (BAR hit), VF group, and VF group offset will be made available as part of `m_axil_awuser` and `m_axil_aruser[54:0]` of the AXI4-Lite interface to help the user logic identify the source of memory access.

The `m_axil_awuser/m_axil_aruser[54:0]` user bits mapping is listed in AXI4-Lite Master Ports.

Virtual function group (VFG) refers to the VF group number. It is equivalent to the PF number associated with the corresponding VF. VFG_OFFSET refer to the VF number with respect to a particular PF. Note that this is not the FIRST_VF_OFFSET of each PF.

For example, if both PF0 and PF1 has 8 VFs, and FIRST_VF_OFFSET for PF0 and PF1 is 4 and 11 and below is the mapping for VFG and VFG_OFFSET.

Table: AXI4-Lite Interface VFG

Function Number	PF Number	VFG	VFG_OFFSET
0	0	0	0
1	1	0	0
4	0	0	0 (Because FIRST_VF_OFFSET for PF0 is 4, the first VF of PF0 starts at FN_NUM=4 and VFG_OSSET=0 indicates this is the first VF for PF0)
5	0	0	1 (VFG_OSSET=1 indicates this is the second VF for PF0)
...
12	1	1	0 (VFG=1 indicates this VF is associated with PF1)

Function Number	PF Number	VFG	VFG_OFFSET
13	1	1	1

Each host initiated access can be uniquely mapped to the 64-bit AXI address space through the PCIe to AXI BAR translation.

Because all functions shares the same AXI4 master address space, a mechanism is needed to map requests from different functions to a distinct address space on the AXI master side. This below shows how PCIe to AXI translation vector is used. Note that all VFs belonging to the same PF shares the same PCIe to AXI translation vector. Therefore, the AXI address space of each VF is concatenated together. Use VFG_OFFSET to calculate the actual starting address of AXI for a particular VF.

To summarize, m_axil_awaddr is determined as:

- For PF, $m_axil_awaddr = pcie2axi_vec + axil_offset$.
- For VF, $m_axil_awaddr = pcie2axi_vec + (VFG_OFFSET + 1) * vf_bar_size + axil_offset$

Where pcie2axi_vec is PCIe to AXI BAR translation (that can be set during IP configuration.). And axib_offset is the address offset in the requested target space.

Each host initiated access can be uniquely mapped to the 64-bit AXI address space. One outstanding read and one outstanding write are supported on this interface.

Expansion ROM BAR can also be mapped to AXI4-Lite interface at the IP configuration time.

PCIe to AXI BARs

For each physical function, the PCIe configuration space consists of a set of six 32-bit memory BARs and one 32-bit Expansion ROM BAR. When SR-IOV is enabled, an additional six 32-bit BARs are enabled for each Virtual Function. These BARs provide address translation to the AXI4 memory mapped space capability, interface routing, and AXI4 request attribute configuration. Any pairs of BARs can be configured as a single 64-bit BAR. A programming example can be found in the Address Translation section (Example 3) of *AXI Bridge for PCI Express Gen3 Subsystem Product Guide* ([PG194](#)).

Request Memory Type

The memory type can be set for each PCIe BAR through attributes attr_dma_pcierbar2axibar_*_cache_pf*.

- AxCache[0] is set to 1 for modifiable, and 0 for non-modifiable.
- AxCache[1] is set to 1 for cacheable, and 0 for non-cacheable.

AXI Memory Mapped Bridge Slave Interface

The AXI-MM Bridge Slave interface is used for high bandwidth memory transfers between the

user logic and the Host. AXI to PCIe translation is supported through the AXI to PCIe BARs. The interface will split requests as necessary to obey PCIe MPS and 4 KB boundary crossing requirements. Up to 32 outstanding read and write requests are supported.

AXI4-Lite Bridge Slave CSR Interface

The AXI4-Lite slave interface is used to access the AXI Bridge and QDMA internal registers. Address bit [15] indicates if the access is for QDMA registers or AXI Bridge registers.

- When `s_axil_csr_awaddr[15]` = 1'b1, the write access is for QDMA CSR registers.
- When `s_axil_csr_awaddr[15]` = 1'b0, the write access is for Bridge registers (When accessing Bridge Registers, access from address 0x000 to 0xDFF will be redirected to PCIe core configuration space access and from address 0xE00 will be directed towards Bridge registers).
- When `s_axil_csr_araddr[15]` = 1'b1, the read access is for QDMA CSR registers.
- When `s_axil_csr_araddr[15]` = 1'b0, the read access is for Bridge registers. When accessing Bridge Registers, access from address 0x000 to 0xDFF will be redirected to PCIe core configuration space access and from address 0xE00 will be directed towards Bridge registers.

The QDMA registers are virtualized for VFs and PFs. For example, VFs and PFs can access different parts of the address space, and each has access to its own queues. To accommodate the function specific accesses, the user logic can provide function ID on `s_axil_awuser[7:0]` for write access and `s_axil_aruser[7:0]` read access, which gives the QDMA proper internal register access. One outstanding read request and one outstanding write request are supported on the AXI4-Lite slave interface.

The AXI4-Lite slave interface is also used to generate Vendor defined messages (VDM) using the Bridge registers.

Related Information

[VDM](#)

AXI to PCIe BARs

In the Bridge Slave interface, there is one BARs which can be configured as 32 bits or 64 bits. This BAR provide address translation from AXI address space to PCIe address space. The address translation is configured through BDF table programming. Refer to Slave Bridge section for BDF programming.

Related Information

[Slave Bridge](#)

Interrupt Module

The IRQ module aggregates interrupts from various sources. The interrupt sources are queue-based interrupts, user interrupts and error interrupts.

Queue-based interrupts and user interrupts are allowed on PFs and VFs, but error interrupts are

allowed only on PFs. If the SR-IOV is not enabled, each PF has the choice of MSI-X or Legacy Interrupts. With SR-IOV enabled, only MSI-X interrupts are supported across all functions. MSI-X interrupt is enabled by default. Host system (Root Complex) will enable one or all of the interrupt types supported in hardware. If MSI-X is enabled, it takes precedence. Up to eight interrupts per function are available. To allow many queues on a given function and each to have interrupts, the QDMA Subsystem for PCIe offers a novel way of aggregating interrupts from multiple queues to a single interrupt vector. In this way, all 2048 queues could in principle be mapped to a single interrupt vector. QDMA offers 256 interrupt aggregation rings that can be flexibly allocated among the 256 available functions.

PCIe Block Interface

PCIe CQ/CC

The PCIe Completer Request (CQ)/Completer Completion (CC) modules receive and process TLP requests from the remote PCIe agent. This interface to the AMD UltraScale+ Integrated Block for PCIe IP operates in an address aligned mode. The module uses the BAR information from the Integrated Block for PCIe IP to determine where the request should be forwarded. The possible destinations for these requests are:

- DMA configuration module
- AXI4 Bridge Master interface
- the AXI4-Lite Bridge Master interface

Non-posted requests are expected to receive completions from the destination, which are forwarded to the remote PCIe agent. For further details, see the *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG213](#)).

PCIe RQ/RC

The PCIe Requester Request (RQ)/Requester Completion (RC) interface generates PCIe TLPs on the RQ bus and processes PCIe Completion TLPs from the RC bus. This interface to the AMD UltraScale+ Integrated Block for PCIe® core operates in DWord aligned mode. With a 512-bit interface, straddling is enabled. While straddling is supported, all combinations of RQ straddled transactions might not be implemented. For further details, see the *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG213](#)).

PCIe Configuration

Several factors can throttle outgoing non-posted transactions. Outgoing non-posted transactions are throttled based on flow control information from the PCIe® integrated block to prevent head of line blocking of posted requests. The DMA will meter non-posted transactions based on the PCIe Receive FIFO space.

General Design of Queues

The multi-queue DMA engine of the QDMA Subsystem for PCIe uses RDMA model queue pairs to allow RNIC implementation in the user logic. Each queue set consists of Host to Card (H2C), Card to Host (C2H), and a C2H Stream Completion (CMPT). The elements of each queue are descriptors.

H2C and C2H are always written by the driver/software; hardware always reads from these queues. H2C carries the descriptors for the DMA read operations from Host. C2H carries the descriptors for the DMA write operations to the Host.

In internal mode, H2C descriptors carry address and length information and are called gather descriptors. They support 32 bits of metadata that can be passed from software to hardware along with every descriptor. The descriptor can be memory mapped (where it carries host address, card address, and length of DMA transfer) or streaming (only host address, and length of DMA transfer) based on context settings. Through descriptor bypass, an arbitrary descriptor format can be defined, where software can pass immediate data and/or additional metadata along with packet.

C2H queue memory mapped descriptors include the card address, the host address and the length. In streaming internal cached mode, descriptors carry only the host address. The buffer size of the descriptor, which is programmed by the driver, is expected to be of fixed size for the whole queue. Actual data transferred associated with each descriptor does not need to be the full length of the buffer size.


The software advertises valid descriptors for H2C and C2H queues by writing its producer index (PIDX) to the hardware. The status descriptor is the last entry of the descriptor ring, except for a C2H stream ring. The status descriptor carries the consumer index (CIDX) of the hardware so that the driver knows when to reclaim the descriptor and deallocate the buffers in the host.

For the C2H stream mode, C2H descriptors will be reclaimed based on the CMPT queue entry. Typically, this carries one entry per C2H packet, indicating one or more C2H descriptors is consumed. The CMPT queue entry carries enough information for software to claim all the descriptors consumed. Through external logic, this can be extended to carry other kinds of completions or information to the host.

CMPT entries written by the hardware to the ring can be detected by the driver using either the color bit in the descriptor or the status descriptor at the end of the CMPT ring. Each CMPT entry can carry metadata for a C2H stream packet and can also serve as a custom completion or immediate notification for the user application.

The base address of all ring buffers (H2C, C2H, and CMPT) should be aligned to a 4 KB address.

Figure: Queue Ring Architecture



The software can program 16 different ring sizes. The ring size for each queue can be selected from context programming. The last queue entry is the descriptor status, and the number of allowable entries is (queue size -1).

For example, if queue size is 8, which contains the entry index 0 to 7, the last entry (index 7) is reserved for status. This index should never be used for PIDX update, and PIDX update should never be equal to CIDX. For this case, if CIDX is 0, the maximum PIDX update would be 6.


In the example above, if traffic has already started and the CIDX is 4, the maximum PIDX update is 3.

H2C and C2H Queues

H2C/C2H queues are rings located in host memory. For both type of queues, the producer is software and consumer is the descriptor engine. The software maintains producer index (PIDX) and a copy of hardware consumer index (HW CIDX) to avoid overwriting unread descriptors. The descriptor engine also maintains consumer index (CIDX) and a copy of SW PIDX, which is to make sure the engine does not read unwritten descriptors. The last entry in the queue is dedicated for the status descriptor where the engine writes the HW CIDX and other status.

The engine maintains a total of 2048 H2C and 2048 C2H contexts in local memory. The context stores properties of the queue, such as base address (BADDR), SW PIDX, CIDX, and depth of the queue.

Figure: Simple H2C and C2H Queue



The figure above shows the H2C and C2H fetch operation.

1. For H2C, the driver writes payload into host buffer, forms the H2C descriptor with the payload buffer information and puts it into H2C queue at the PIDX location. For C2H, the driver forms the descriptor with available buffer space reserved to receive the packet write from the DMA.
2. The driver sends the posted write to PIDX register in the descriptor engine for the associated Queue ID (QID) with its current PIDX value.
3. Upon reception of the PIDX update, the engine calculates the absolute QID of the pointer update based on address offset and function ID. Using the QID, the engine will fetch the context for the absolute QID from the memory associated with the QDMA Subsystem for PCIe.
4. The engine determines the number of descriptors that are allowed to be fetched based on the context. The engine calculates the descriptor address using the base address (BADDR), CIDX, and descriptor size, and the engine issues the DMA read request.
5. After the descriptor engine receives the read completion from the host memory, the descriptor engine delivers them to the H2C Engine or C2H Engine in internal mode. In case of bypass, the descriptors are sent out to the associated descriptor bypass output interface.
6. For memory mapped or H2C stream queues programmed as internal mode, after the fetched descriptor is completely processed, the engine writes the CIDX value to the status descriptor. For queues programmed as bypass mode, user logic controls the write back through bypass in interface. The status descriptor could be moderated based on context settings. C2H stream queues always use the CMPT ring for the completions.

For C2H, the fetch operation is implicit through the CMPT ring.


Completion Queue

The Completion (CMPT) queue is a ring located in host memory. The consumer is software, and

the producer is the CMPT engine. The software maintains the consumer index (CIDX) and a copy of hardware producer index (HW PIDX) to avoid reading unwritten completions. The CMPT engine also maintains PIDX and a copy of software consumer index (SW CIDX) to make sure that the engine does not overwrite unread completions. The last entry in the queue is dedicated for the status descriptor which is where the engine writes the hardware producer index (HW PIDX) and other status.

The engine maintains a total of 2048 CMPT contexts in local memory. The context stores properties of the queue, such as base address, SW CIDX, PIDX, and depth of the queue.

Figure: Simple Completion Queue Flow



C2H stream is expected to use the CMPT queue for completions to host, but it can also be used for other types of completions or for sending messages to the driver. The message through the CMPT is guaranteed to not bypass the corresponding C2H stream packet DMA.

The simple flow of DMA CMPT queue operation with respect to the numbering above follows:

1. The CMPT engine receives the completion message through the CMPT interface, but the QID for the completion message comes from the C2H stream interface. The engine reads the QID index of CMPT context RAM.
2. The DMA writes the CMPT entry to address BASE+PIDX.
3. If all conditions are met, optionally writes PIDX to the status descriptor of the CMPT queue with color bit.
4. If interrupt mode is enabled, the CMPT engine generates the interrupt event message to the interrupt module.
5. The driver can be in polling or interrupt mode. Either way, the driver identifies the new CMPT.

- entry either by matching the color bit or by comparing the PIDX value in the status descriptor against its current software CIDX value.
6. The driver updates CIDX for that queue. This allows the hardware to reuse the descriptors again. After the software finishes processing the CMPT, that is, before it stops polling or leaving the interrupt handler, the driver issues a write to CIDX update register for the associated queue.

SR-IOV Support

The QDMA Subsystem for PCIe provides an optional feature to support Single Root I/O Virtualization (SR-IOV). The PCI-SIG® Single Root I/O Virtualization and Sharing (SR-IOV) specification (available from *PCI-SIG Specifications* (www.pcisig.com/specifications)) standardizes the method for bypassing the VMM involvement in datapath transactions and allows a single endpoint to appear as multiple separate endpoints. SR-IOV classifies the functions as:

- *Physical Functions (PF)*: Full featured PCIe® functions which include SR-IOV capabilities among others.
- *Virtual Functions (VF)*: PCIe functions featuring configuration space with Base Address Registers (BARs) but lacking the full configuration resources and controlled by the PF configuration. The main role of the VF is data transfer.

Apart from PCIe defined configuration space, QDMA Subsystem for PCI Express virtualizes data path operations, such as pointer updates for queues, and interrupts. The rest of the management and configuration functionality is deferred to the physical function driver. The Drivers that do not have sufficient privilege must communicate with the privileged Driver through the mailbox interface which is provided in part of the QDMA Subsystem for PCI Express.


Security is an important aspect of virtualization. The QDMA Subsystem for PCI Express offers the following security functionality:

- QDMA allows only privileged PF to configure the per queue context and registers. VFs inform the corresponding PFs of any queue context programming.
- Drivers are allowed to do pointer updates only for the queue allocated to them.
- The system IOMMU can be turned on to check that the DMA access is being requested by PFs or VFs. The ARID comes from queue context programmed by a privileged function.

Any PF or VF can communicate to a PF (not itself) through mailbox. Each function implements one 128B inbox and 128B outbox. These mailboxes are visible to the driver in the DMA BAR (typically BAR0) of its own function. At any given time, any function can have one outgoing mailbox and one incoming mailbox message outstanding per function.

The diagram below shows how a typical system can use QDMA with different functions and operating systems. Different Queues can be allocated to different functions, and each function can transfer DMA packets independent of each other.

Figure: QDMA in a System



X21109-062218

Limitations

The limitations of the QDMA Subsystem for PCIe are as follows:

- The DMA supports a maximum of 256 Queues on any VF function.
- Slave Bridge AXI does not support Narrow Burst transfers.

 - ☞ **Recommended:** Use AXI SmartConnect to support Narrow Burst.
 - ECC and Slave Narrow Burst support is mutually exclusive.
 - If you want an ECC feature, the recommendation is to up-size your AXI Master externally.
- SRIOV is not supported in bridge mode.
- This IP architecture assumes exclusive use of one or more complete GT quads, regardless of the designed link width. While it might be possible to share unused lanes in the GT quad with other instances of this IP, non-PCIe IPs, or custom GT-based interfaces for x2 and x1 link widths, AMD does not support evaluations or implementations of such sharing arrangements. The feasibility of sharing depends on the specific GT configuration required for other protocols, links, and lanes intended to share the GT quad. Factors affecting GT configuration include external REFCLKs, fabric design clocks and resets, GT clock management resources, connectivity rules, mode, and electrical settings.

Applications

The QDMA Subsystem for PCIe is used in a broad range of networking, computing, and data storage applications. A common usage example for the QDMA Subsystem for PCIe is to implement Data Center and Telco applications, such as Compute acceleration, Smart NIC, NVMe, RDMA-enabled NIC (RNIC), server virtualization, and NFV in the user logic. Multiple applications can be implemented to share the QDMA by assigning different queue sets and PCIe functions to each application. These Queues can then be scaled in the user logic to implement rate limiting, traffic priority, and custom work queue entry (WQE).

Licensing and Ordering

This AMD LogicORE™ IP module is provided at no additional cost with the AMD Vivado™ Design Suite under the terms of the [End User License](#).

For more information about this subsystem, visit the [QDMA Subsystem for PCIe product page](#) web page.

Product Specification

Standards

The AMD adheres to the following standards:

- AMBA AXI4-Stream Protocol Specification ([ARM IHI 0051A](#))
- PCI Express Base Specification v3.1
- PCI Local Bus Specification
- PCI-SIG® Single Root I/O Virtualization and Sharing (SR-IOV) Specification

For details, see *PCI-SIG Specifications* (<https://www.pcisig.com/specifications>).

Performance and Resource Utilization

Performance

QDMA performance and detailed analysis is available in AR [71453](#).

AMD provides two example designs for you to experiment with. Standard example design is for functional test only. To generate an example design for performance analysis, use the following Tcl command to generate a performance example design:

```
set_property CONFIG.performance_exdes {true} [get_ips qdma_0]
```

Following are the QDMA register settings recommended by AMD for better performance.

Performance numbers can vary based on systems and OS used.

Table: QDMA Performance Registers

Address	Name	Fields	Field Value	Register Value
0xB08	PFCH CFG	<ul style="list-style-type: none">• evt_pfch_fl_th[15:0]• pfch_fl_th[15:0]	<ul style="list-style-type: none">• 256• 256	0x100_0100
0xA80	PFCH_CFG_1	<ul style="list-style-type: none">• evt_qcnt_th[15:0]• pfch_qcnt[15:0]	<ul style="list-style-type: none">• 60• 60	0x3c_003c
0xA84	PFCH_CFG_2	<ul style="list-style-type: none">• fence• rsvd[1:0]• var_desc_no_drop• pfch_ll_sz_th[15:0]• var_desc_num_pfch[5:0]• num_pfch[5:0]	<ul style="list-style-type: none">• 1• 0• 0• 1024• 15• 8	0x8040_03C8

Address	Name	Fields	Field Value	Register Value
0x147C	PFCH_CFG_3	<ul style="list-style-type: none"> • rsvd[4:0] • var_desc_fl_free_cnt_th[8:0] • var_desc_lg_pkt_cam_cn_th[6:0] 	<ul style="list-style-type: none"> • 0 • 256 • 0 	0x8000
0x1484	PFCH_CFG_4	<ul style="list-style-type: none"> • glb_evt_timer_tick[14:0] • disable_glb_evt_timer • evt_timer_tick[14:0] • disable_evt_timer 	<ul style="list-style-type: none"> • 64 • 0 • 400 • 0 	0x80_0320
0x1400	CRDT_COAL_CFG_1	<ul style="list-style-type: none"> • rsvd[12:0] • dis_fence_fix • pld_fifo_th[7:0] • crdt_timer_th[9:0] 	<ul style="list-style-type: none"> • NA • 0 • 16 • 16 	0x4010
0x1404	CRDT_COAL_CFG_2	<ul style="list-style-type: none"> • rsv2[7:0] • crdt_fifo_th[7:0] • rsv1[4:0] • crdt_cnt_th[10:0] 	<ul style="list-style-type: none"> • NA • 56 • NA • 96 	0x38_0060
0x15C	GLBL_RRQ_PCIE_THROT	<ul style="list-style-type: none"> • req_throt_en • req_throt • dat_throt_en • dat_throt 	<ul style="list-style-type: none"> • 0 • 192 • 1 • 20480 	0x604_5000
0x160	GLBL_RRQ_AXIMM_THROT	<ul style="list-style-type: none"> • req_throt_en • req_throt • dat_throt_en • dat_throt 	<ul style="list-style-type: none"> • 0 • 0 • 0 • 0 	0
0x158	GLBL_RRQ_BRG_THROT	<ul style="list-style-type: none"> • req_throt_en • req_throt • dat_throt_en • dat_throt 	<ul style="list-style-type: none"> • 1 • 192 • 1 • 20480 	0x8604_5000

Address	Name	Fields	Field Value	Register Value
0xE24	H2C_REQ_THROT_PCIE	<ul style="list-style-type: none"> • req_throt_en_req • req_throt • req_throt_en_data • data_thresh 	<ul style="list-style-type: none"> • 1 • 192 • 1 • 24576 	0x8604_6000
0xE2C	H2C_REQ_THROT_AXIMM	<ul style="list-style-type: none"> • req_throt_en_req • req_throt • req_throt_en_data • data_thresh 	<ul style="list-style-type: none"> • 1 • 64 • 1 • 16384 	0x8204_4000
0x12EC	H2C_MM_DATA_THROT	<ul style="list-style-type: none"> • data_throt_en • data_throt 	<ul style="list-style-type: none"> • 1 • 20480 	0x1_5000
0x250	QDMA_GLBL_DSC_CFG	<ul style="list-style-type: none"> • c2h_uodsc_limit (Soft IP) • h2c_uodsc_limit (Soft IP) • uodsc_limit (KS-B) • Max_dsc_fetch • wb_acc_int 	<ul style="list-style-type: none"> • 5 • 8 • NA • 2 • 5 	0x50_2015
0x4C	CONFIG_BLOCK_MISC_CONTROL	<ul style="list-style-type: none"> • 10b_tag_en • rq_metering_multiplier 	<ul style="list-style-type: none"> • 0 • 256 • 9 	0x1_0009

- QDMA_C2H_INT_TIMER_TICK (0xB0C) set to 25. Corresponding to 100 ns (1 tick = 4 ns for 250 MHz user clock)
- C2H trigger mode set to user timer, with counter set to 64 and timer to match round trip latency. Global register for timer should have a value of 30 for 3 µs.
- TX/RX API burst size = 64, ring depth = 2048. The driver should update TX/RX PIDX in batches of 64.
- PCIe MPS = 256 bytes, MRRS >= 512 bytes, Extended Tag Enabled, Relaxed Ordering Enabled
- The driver will update the completion CIDX in batches of 64 to reduce number of MMIO writes before updating the C2H PIDX
- The driver should update the H2C PIDX in batches of 64, and also update for the last descriptor of the scatter gather list.
- C2H context:
 - bypass = 0 (Internal mode)
 - frcd_en = 1
 - qen = 1
 - wbk_en = 1
 - irq_en = irq_arm = int_aggr = 0
- C2H prefetch context:
 - pfch = 1
 - bypass = 0
 - valid = 1
- C2H CMPT context:
 - en_stat_desc = 1
 - en_int = 0 (Poll_mode)
 - int_aggr = 0 (Poll mode)
 - trig_mode = 4
 - counter_idx = corresponding to 64
 - timer_idx = corresponding to 3 µs
 - valid = 1
- H2C context:
 - bypass = 0 (Internal mode)
 - frcd_en = 0
 - fetch_max = 0
 - qen = 1
 - wbk_en = 1
 - wbi_chk = 1
 - wbi_intvl_en = 1
 - irq_en = 0 (Poll mode)
 - irq_arm = 0 (Poll mode)
 - int_aggr = 0 (Poll mode)

For optimal QDMA streaming performance, packet buffers of the descriptor ring should be

aligned to at least 256 bytes.

 **Recommended:** AMD recommends that you limit the total outstanding descriptor fetch to be less than 8 KB on the PCIe. For example, limit the outstanding credits across all queues to 512 for a 16B descriptor.

Performance in Descriptor Bypass Mode

When the design is configured in descriptor bypass mode, all the above setting apply. The following information provides recommendations to improve performance in bypass mode.

- When bypass in `h2c_byp_in_st_sdi` ports is set, the QDMA IP generates the status write back for every packet. AMD recommends that this port be asserted once in 32 packets or 64 packets. And if there are no more descriptors left then assert `h2c_byp_in_st_sdi` at the last descriptor. This requirement is per queue basis, and applies to AXI4 (H2C and C2H) bypass transfers and AXI4-Stream H2C transfers.
- For AXI4-Stream C2H Simple bypass mode, the `dsc_crdt_in_fence` port should be set to 1 for performance reasons. This recommendation assumes the user design already coalesced credits for each queue and sent them to the IP. In internal mode, set the fence bit in the `QDMA_C2H_PFCH_CFG_2 (0xA84)` register.

Performance Optimization Based on Available Cache/Buffer Size

Table: QDMA Soft IP

Name	Entry/Depth	Description
C2H descriptor cache depth	1024	Total number of outstanding C2H stream descriptor fetches for cache bypass and internal. This cache depth is not relevant in simple bypass mode, in simple bypass mode you can have longer descriptor cache.
Prefetch cache depth	64	C2H prefetch tags available. If you have more than 64 active queues for packets < 512 B, performance can reduce depending on the data pattern. If you see performance degradation, you can implement simple bypass mode, where you can maintain the descriptor flow.

Name	Entry/Depth	Description
C2H payload FIFO depth	512	Units of 64 B. The amount of C2H data that C2H engine can buffer. This amount of buffer can sustain the host read latency up to 2 us (512 *4 ns). If latency is more than 2 us, there could be performance degradation.
Common reorder buffer depth	512	Units of 64 B for Soft IP. Shared buffer space that can be flexibly allocated between the read engines. Throttle CSRs can be used to limit the amount of outstanding read data used by each engine in this common buffer space.

Resources Utilization

For QDMA Resource Utilization, see [Resource Use web page](#).

Minimum Device Requirements

Gen3x16 capability requires a minimum of a -2 speed grade.

Table: Minimum Device Requirements

Capability Link Speed	Capability Link Width	Supported Speed Grades
AMD UltraScale+™ Devices with PCIE4 Block		
Gen1/Gen2	x1, x2, x4, x8, x16	-1, -1L, -1LV, -2, -2L, -2LV, -3
Gen3	x1, x2, x4	-1, -1L, -1LV, -2, -2L, -2LV, -3
	x8	-1, -2, -2L, -3
	x16	-2, -2L, -3
AMD UltraScale+™ Devices with PCIE4C and PCIE4CE Block		
Gen1/Gen2	x1, x2, x4, x8, x16	-1, -2, -2L, -2LV, -3
Gen3	x1, x2, x4	-1, -2, -2L, -2LV, -3
	x8	-1, -2, -2L, -3
	x16	-2, -2L, -3
Gen4	x1, x2, x4, x8	-2, -2L, -3

Capability Link Speed	Capability Link Width	Supported Speed Grades
AMD Spartan™ UltraScale+™ Devices (PCIE4CE)		
Gen1/Gen2	x1, x2, x4, x8	-1, -1L, -1LV, -2
Gen3	x1, x2, x4, x8	-1, -1L ¹ , -1LV ¹ , -2
Gen4	x1, x2, x4, x8	-2

1. Gen3 x8 supports normal latency (NL) and low latency (LL) mode, where core clock is 250 MHz and 500 MHz respectively. -1L and -1LV speed grades only support NL mode.

 **Note:** Subject to the documented product features and minimum device requirements, this IP supports AMD UltraScale+™ devices with one or more PCIE4/PCIE4C/PCIE4CE integrated blocks for PCIe. Based on the available programmable logic resources, the following are NOT supported even if this IP is supported by the device architecture:

- Zynq UltraScale+ MPSoC devices ZU5 and smaller
- Kintex UltraScale+ FPGA devices KU3 and smaller
- Artix UltraScale+ FPGA devices AU25P and smaller

Contact [Support](#) for information about implementing this IP in devices containing at least one integrated block for PCIe, but are not supported based on the available programmable logic resources.

QDMA Operations

Descriptor Engine

The descriptor engine is responsible for managing the consumer side of the Host to Card (H2C) and Card to Host (C2H) descriptor ring buffers for each queue. The context for each queue determines how the descriptor engine will process each queue individually. When descriptors are available and other conditions are met, the descriptor engine will issue read requests to PCIe to fetch the descriptors. Received descriptors are offloaded to either the descriptor bypass out interface (bypass mode) or delivered directly to a DMA engine (internal mode). When a H2C Stream or Memory Mapped DMA engine completes a descriptor, status can be written back to the status descriptor, an interrupt, and/or a marker response can be generated to inform software and user logic of the current DMA progress. The descriptor engine also provides a Traffic Manager Interface which notifies user logic of certain status for each queue. This allows the user logic to make informed decisions if customization and optimization of DMA behavior is desired.

Descriptor Context

The Descriptor Engine stores per queue configuration, status and control information in

descriptor context that can be stored in block RAM or UltraRAM, and the context is indexed by H2C or C2H QID. Prior to enabling the queue, the hardware and credit context must first be cleared. After this is done, the software context can be programmed and the qen bit can be set to enable the queue. After the queue is enabled, the software context should only be updated through the direct mapped address space to update the Producer Index and Interrupt Arm® bit, unless the queue is being disabled. The hardware context and credit context contain only status. It is only necessary to interact with the hardware and credit contexts as part of queue initialization to clear them to all zeros. Once the queue is enabled, context is dynamically updated by hardware. Any modification of the context through the indirect bus when the queue is enabled can result in unexpected behavior. Reading the context when the queue is enabled is not recommended as it can result in reduced performance.

Related Information

[QDMA_DMAP_SEL_H2C_DSC_PIDX\[2048\] \(0x18004\)](#)

[QDMA_DMAP_SEL_C2H_DSC_PIDX\[2048\] \(0x18008\)](#)

Software Descriptor Context Structure (0x0 C2H and 0x1 H2C)

The descriptor context is used by the descriptor engine.

Table: Software Descriptor Context Structure Definition

Bit	Bit Width	Field Name	Description
[255:140]	116	reserved	Reserved. Set to 0s.
[139]	1	int_aggr	If set, interrupts will be aggregated in interrupt ring.
[138:128]	[10:0]	vec	MSI-X vector used for interrupts for direct interrupt or interrupt aggregation entry for aggregated interrupts.
[127:64]	64	dsc_base	4K aligned base address of descriptor ring.
[63]	1	is_mm	This field determines if the queue is Memory Mapped or not. If this field is set, the descriptors will be delivered to associated H2C or C2H MM engine. 1: Memory Mapped 0: Stream
[62]	1	mrkr_dis	If set, disables the marker response in internal mode. Not applicable for C2H ST.

Bit	Bit Width	Field Name	Description
[61]	1	irq_req	Interrupt due to error waiting to be sent (waiting for irq_arm). This bit should be cleared when the queue context is initialized. Not applicable for C2H ST.
[60]	1	err_wb_sent	A writeback/interrupt was sent for an error. Once this bit is set no more writebacks or interrupts will be sent for the queue. This bit should be cleared when the queue context is initialized. Not applicable for C2H ST.
[59:58]	2	err	Error status. Bit[1] dma – An error occurred during DMA operation. Check engine status registers. Bit[0] dsc – An error occurred during descriptor fetch or update. Check descriptor engine status registers. This field should be set to 0 when the queue context is initialized.
[57]	1	irq_no_last	No interrupt was sent and the producer index (PIDX) or consumer index (CIDX) was idle in internal mode. When the irq_arm bit is set, the interrupt will be sent. This bit will clear automatically when the interrupt is sent or if the PIDX of the queue is updated. This bit should be initialized to 0 when the queue context is initialized. Not applicable for C2H ST.
[56:54]	3	port_id	Port_id The port id that will be sent on user interfaces for events associated with this queue.
[53]	1	irq_en	Interrupt enable. An interrupt to the host will be sent on host status updates. Set to 0 for C2H ST.

Bit	Bit Width	Field Name	Description
[52]	1	wbk_en	Writeback enable. A memory write to the status descriptor will be sent on host status updates.
[51]	1	mm_chn	Set to 0 and cannot be modified.
[50]	1	bypass	If set, the queue will operate under Bypass mode, otherwise it will be in Internal mode.
[49:48]	2	dsc_sz	Descriptor fetch size. 0: 8B, 1: 16B; 2: 32B; 3: 64B. If bypass mode is not enabled, 32B is required for Memory Mapped DMA, 16B is required for H2C Stream DMA, and 8B is required for C2H Stream DMA. If the queue is configured for bypass mode, any descriptor size can be selected. The descriptors will be delivered on the bypass output interface. It is up to the user logic to process the descriptors before they are fed back into the descriptor bypass input.
[47:44]	4	rng_sz	Descriptor ring size index. This index selects one of 16 register (offset 0x204:0x240) which has different ring sizes.
[43:41]	3	reserved	Reserved
[40:37]	4	fetch_max	Maximum number of descriptor fetches outstanding for this queue. The max outstanding is fetch_max + 1. Higher value can increase the single queue performance,
[36]	1	at	Address type of base address. 0: untranslated 1: translated This will be the address type (AT) used on PCIe for descriptor fetches and status descriptor writebacks.

Bit	Bit Width	Field Name	Description
[35]	1	wbi_intvl_en	Write back/Interrupt interval. Enables periodic status updates based on the number of descriptors processed. Applicable to Internal mode. Not Applicable to C2H ST. The writeback interval is determined by register QDMA_GLBL_DSC_CFG (0x250) bits[2:0].
[34]	1	wbi_chk	Writeback/Interrupt after pending check. Enable status updates when the queue has completed all available descriptors. Applicable to Internal mode.
[33]	1	fcrd_en	Enable fetch credit. The number of descriptors fetched will be qualified by the number of credits given to this queue. Set to 1 for C2H ST.
[32]	1	qen	Indicates that the queue is enabled.
[31:25]	7	reserved	Reserved
[24:17]	8	fnc_id	Function ID
[16]	1	irq_arm	Interrupt arm. When this bit is set, the queue is allowed to generate an interrupt.
[15:0]	16	pidx	Producer index.

Hardware Descriptor Context Structure (0x2 C2H and 0x3 H2C)

Table: Hardware Descriptor Structure Definition

Bit	Bit Width	Field Name	Description
[47]	1	reserved	Reserved
[46:43]	4	fetch_pnd	Descriptor fetch pending
[42]	1	evt_pnd	Event pending
[41]	1	idl_stp_b	Queue invalid and no descriptors pending. This bit is set when the queue is enabled. The bit is cleared when the queue has

Bit	Bit Width	Field Name	Description
			been disabled (software context qen bit) and no more descriptors are pending. [0] Queue is disabled (software context qen bit) and no more descriptors pending. [1] Queue is enabled.
[40]	1	dsc_pnd	Descriptors pending. Descriptors are defined to be pending if the last CIDX completed does not match the current PIDX.
[39:32]	8		Reserved
[31:16]	16	crd_use	Credits consumed. Applicable if fetch credits are enabled in the software context.
[15:0]	16	cidx	Consumer index of last fetched descriptor.

Credit Descriptor Context Structure


Table: Credit Descriptor Context Structure Definition

Bit	Bit Width	Field Name	Description
[31:16]	16	reserved	Reserved
[15:0]	16	credit	Fetch credits received. Applicable if fetch credits are enabled in the software context.

The credit descriptor context is for internal DMA use only and it can be read from the indirect bus for debug. This context stores credits for each queue received through the Descriptor Credit Interface with the CREDIT_ADD operation. If the credit operation has the dsc_crdt_in_fence bit set to 1, credits are added only as the read request for the descriptor is generated.

Descriptor Fetch

Figure: Descriptor Fetch Flow



1. The descriptor engine is informed of the availability of descriptors through an update to a queue's descriptor PIDX. This portion of the context is direct mapped to the QDMA_DMAP_SEL_H2C_DSC_PIDX and QDMA_DMAP_SEL_C2H_DSC_PIDX address space.
2. On a PIDX update, the descriptor engine evaluates the number of descriptors available based on the last fetched consumer index (CIDX). The availability of new descriptors is communicated to the user logic through the Traffic Manager Status Interface.
3. If fetch crediting is enabled, the user logic is required to provide a credit for each descriptor that should be fetched.
4. If descriptors are available and either fetch credits are disabled or are non-zero, the descriptor engine will generate a descriptor fetch to PCIe. The number of fetched descriptors is further qualified by the PCIe Max Read Request Size (MRRS) and descriptor fetch credits, if enabled. A descriptor fetch can also be stalled due to insufficient completion space. In each direction, C2H and H2C are allocated 256 entries for descriptor fetch completions. Each entry is the width of the datapath. If sufficient space is available, the fetch is allowed to proceed. A given queue can only have one descriptor fetch pending on PCIe at any time.
5. The host receives the read request and provides the descriptor read completion to the descriptor engine.
6. Descriptors are stored in a buffer until they can be offloaded. If the queue is configured in bypass mode, the descriptors are sent to the Descriptor Bypass Output port. Otherwise they are delivered directly to a DMA engine. Once delivered, the descriptor fetch completion buffer space is deallocated.

Note: Available descriptors are always <ring size> - 2. At any time, the software should not update the PIDX to more than <ring size> - 2.

For example, if queue size is 8, which contains the entry index 0 to 7, the last entry (index 7) is reserved for status. This index should never be used for the PIDX update, and the PIDX update

should never be equal to CIDX. For this case, if CIDX is 0, the maximum PIDX update would be 6.

Internal Mode

A queue can be configured to operate in Descriptor Bypass mode or Internal mode by setting the software context bypass field. In internal mode, the queue requires no external user logic to handle descriptors. Descriptors that are fetched by the descriptor engine are delivered directly to the appropriate DMA engine and processed. Internal mode allows credit fetching and status updates to the user logic for run time customization of the descriptor fetch behavior.

Internal Mode Writeback and Interrupts (AXI MM and H2C ST)

Status writebacks and/or interrupts are generated automatically by hardware based on the queue context. When `wbi_intvl_en` is set, writebacks/interrupts are sent based on the interval selected in the register QDMA_GLBL_DSC_CFG (0x250) bits[2:0]. Due to the slow nature of interrupts, in interval mode, interrupts might be late or skip intervals. If the `wbi_chk` context bit is set, a writeback/interrupt is sent when the descriptor engine has detected that the last descriptor at the current PIDX has completed. It is recommended the `wbi_chk` bit be set for all internal mode operation, including when interval mode is enabled. An interrupt is not generated until the `irq_arm` bit is set by the software. After an interrupt is sent, the `irq_arm` bit is cleared by the hardware. Should an interrupt be needed when the `irq_arm` bit is not set, the interrupt is held in a pending state until the `irq_arm` bit is set.

Descriptor completion is defined to be when the descriptor data transfer has completed and its write data is acknowledged on AXI (H2C bresp for AXI MM, Valid/Ready of ST), or it is accepted by the PCIe Controller's transaction layer for transmission (C2H MM).

Descriptor Bypass Mode

Descriptor Bypass mode also supports crediting and status updates to user logic. In addition, Descriptor Bypass mode allows the user logic to customize processing of descriptors and status updates. Descriptors fetched by the descriptor engine are delivered to user logic through the descriptor bypass out interface. This allows user logic to pre-process or store the descriptors, if desired. On the descriptor bypass out interface, the descriptors can be a custom format (adhering to the descriptor size). To perform DMA operations, the user logic drives descriptors (must be QDMA format) into the descriptor bypass input interface.

Descriptor Bypass Mode Writeback/Interrupts

In bypass mode, the user logic has explicit control over status updates to the host, and marker responses back to user logic. Along with each descriptor submitted to the Descriptor Bypass Input Port for a Memory Mapped Engine (H2C and C2H) or H2C Stream DMA engine, there is a CIDX, and `sdi` field. The CIDX is used to identify which descriptor has completed in any status update (host writeback, marker response, or coalesced interrupt) generated at the completion of the descriptor. If the `sdi` field of the descriptor is input, then a writeback to the host is generated if the context `wbk_en` bit is set. An interrupt can also be sent if the `sdi` bit is set if the context

`irq_en` and `irq_arm` bits are set.

If interrupts are enabled, the user logic must monitor the traffic manager output for the `irq_arm`. After the `irq_arm` bit is observed for the queue, a descriptor with the `sdi` bit is sent to the DMA. Once a descriptor with the `sdi` bit is sent, another `irq_arm` assertion must be observed before another descriptor with the `sdi` bit can be sent. If you set the `sdi` bit when the `arm` bit is not properly observed, an interrupt might or might not be sent, and software might hang indefinitely waiting for an interrupt. When interrupts are not enabled, setting the `sdi` bit has no restriction. However, excessive writeback events can severely reduce the descriptor engine performance and consume write bandwidth to the host.

Descriptor completion is defined to be when the descriptor data transfer has completed and its write data has been acknowledged on AXI4 (H2C bresp for AXI MM, Valid/Ready of ST), or been accepted by the PCIe Controller's transaction layer for transmission (C2H MM).

Marker Response

Marker responses can be generated for any descriptor by setting the `mrkr_req` bit. Marker responses are generated after the descriptor is completed. Similar to host writebacks, excessive marker response requests can reduce descriptor engine performance. Along with `mrkr_req` signals, `sdi` can also be set. In this case, the marker response is sent on queue status ports and writeback is sent to the host. The marker responses are sent on queue status ports that can be identified by the queue id.

Descriptor completion is defined as when the descriptor data transfer has completed and its write data is acknowledged on AXI (H2C bresp for AXI4, Valid/Ready of ST), or is accepted by the PCIe Controller's transaction layer for transmission (C2H MM).

Traffic Manager Output Interface

The traffic manager interface provides details of a queue's status to user logic, allowing user logic to manage descriptor fetching and execution. In normal operation, for an enabled queue, each time the `irq_arm` bit is asserted or PIDX of a queue is updated, the descriptor engine asserts `tm_dsc_sts_valid`. The `tm_dsc_sts_avl` signal indicates the number of new descriptors available since the last update. Through this mechanism, user logic can track the amount of work available for each queue. This can be used for prioritizing fetches through the descriptor engine's fetch crediting mechanism or other user optimizations. On the valid cycle, the `tm_dsc_sts_irq_arm` indicates that the `irq_arm` bit was zero and was set. In bypass mode, this is essentially a credit for an interrupt for this queue. When a queue is invalidated by software or due to error, the `tm_dsc_sts_qinv` signal will be set. If this bit is observed, the descriptor engine will have halted new descriptor fetches for that queue. In this case, the contents on `tm_dsc_sts_avl` indicate the number of available fetch credits held by the descriptor engine. This information can be used to help user logic reconcile the number of credits given to the descriptor engine, and the number of descriptors it should expect to receive. Even after `tm_dsc_sts_qinv` is asserted, valid descriptors already in the fetch pipeline will continue to be delivered to the DMA engine (internal mode) or delivered to the descriptor bypass output port (bypass mode).

Other fields of the `tm_dsc_sts` interface identify the queue id, DMA direction (H2C or C2H), internal or bypass mode, stream or memory mapped mode, queue enable status, queue error status, and port ID.

While the `tm_dsc_sts` interface is a valid/ready interface, it should not be back-pressured for optimal performance. Since multiple events trigger a `tm_dsc_sts` cycle, if internal buffering is filled, descriptor fetching will be halted to prevent generation of new events.

Related Information

[QDMA Traffic Manager Credit Output Ports](#)

Descriptor Credit Input Interface

The credit interface is relevant when a queue's `fcrd_en` context bit is set. It allows the user logic to prioritize and meter descriptors fetched for each queue. You can specify the DMA direction, qid, and credit value. For a typical use case, the descriptor engine uses credit inputs to fetch descriptors. Internally, credits received and consumed are tracked for each queue. If credits are added when the queue is not enabled, the credits will be returned through the Traffic Manager Output Interface with `tm_dsc_sts_qinv` asserted, and the credits in `tm_dsc_sts_avl` are not valid. Monitor `tm_dsc_sts` interface to keep an account for each queue on how many credits are consumed.

Related Information

[QDMA Descriptor Credit Input Ports](#)

Errors

Errors can potentially occur during both descriptor fetch and descriptor execution. In both cases, once an error is detected for a queue it will invalidate the queue, log an error bit in the context, stop fetching new descriptors for the queue which encountered the error, and can also log errors in status registers. If enabled for writeback, interrupts, or marker response, the DMA will generate a status update to these interfaces. Once this is done, no additional writeback, interrupts, or marker responses (internal mode) will be sent for the queue until the queue context is cleared. As a result of the queue invalidation due to an error, a Traffic Manager Output cycle will also be generated to indicate the error and queue invalidation. After the queue is invalidated, if there is an error you can determine the cause by reading the error registers and context for that queue. You must clear and remove that queue, and then add the queue back later when needed.

Although additional descriptor fetches will be halted, fetches already in the pipeline will continue to be processed and descriptors will be delivered to a DMA engine or Descriptor Bypass Out interface as usual. If the descriptor fetch itself encounters an error, the descriptor will be marked with an error bit. If the error bit is set, the contents of the descriptor should be considered invalid. It is possible that subsequent descriptor fetches for the same queue do not encounter an error and will not have the error bit set.

Memory Mapped DMA

In memory mapped DMA operations, both the source and destination of the DMA are memory

mapped space. In an H2C transfer, the source address belongs to PCIe address space while the destination address belongs to AXI MM address space. In a C2H transfer, the source address belongs to AXI MM address space while the destination address belongs to PCIe address space. PCIe-to-PCIe, and AXI MM-to-AXI MM DMAs are not supported. Aside from the direction of the DMA, transfer H2C and C2H DMA behave similarly and share the same descriptor format.

Operation

The memory mapped DMA engines (H2C and C2H) are enabled by setting the run bit in the Memory Mapped Engine Control Register. When the run bit is deasserted, descriptors can be dropped. Any descriptors that have already started the source buffer fetch will continue to be processed. Reassertion of the run bit will result in resetting internal engine state and should only be done when the engine is quiesced. Descriptors are received from either the descriptor engine directly or the Descriptor Bypass Input interface. Any queue that is in internal mode should not be given descriptors through the Descriptor Bypass Input interface. Any descriptor sent to an MM engine that is not running will be dropped. For configurations where a mix of Internal Mode queues and Bypass Mode queues are enabled, round robin arbitration is performed to establish order.

The DMA Memory Mapped engine first generates the read request to the source interface, splitting the descriptor at alignment boundaries specific to the interface. Both PCIe and AXI read interfaces can be configured to split at different alignments. Completion space for read data is preallocated when the read is issued. Likewise for the write requests, the DMA engine will split at appropriate alignments. On the AXI interface each engine will use a single AXI ID. The DMA engine will reorder the read completion/write data to the order in which the reads were issued. Once sufficient read completion data is received the write request will be issued to the destination interface in the same order that the read data was requested. Before the request is retired, the destination interfaces must accept all the write data and provide a completion response. For PCIe the write completion is issued when the write request has been accepted by the transaction layer and will be sent on the link next. For the AXI Memory Mapped interface, the bresp is the completion criteria. Once the completion criteria has been met, the host writeback, interrupt and/or marker response is generated for the descriptor as appropriate.

The DMA Memory Mapped engines also support the no_dma field of the Descriptor Bypass Input, and zero-length DMA. Both cases are treated identically in the engine. The descriptors propagate through the DMA engine as all other descriptors, so descriptor ordering within a queue is still observed. However no DMA read or write requests are generated. The status update (writeback, interrupt, and/or marker response) for zero-length/no_dma descriptors is processed when all previous descriptors have completed their status update checks.

Related Information

[Descriptor Bypass Mode Writeback/Interrupts](#)

Errors

There are two primary error categories for the DMA Memory Mapped Engine. The first is an error bit that is set with an incoming descriptor. In this case, the DMA operation of the descriptor is not

processed but the descriptor proceeds through the engine to status update phase with an error indication. This should result in a writeback, interrupt, and/or marker response depending on context and configuration. It also results in the queue being invalidated. The second category of errors for the DMA Memory Mapped Engine are errors encountered during the execution of the DMA itself. This can include PCIe read completions errors, and AXI bresp errors (H2C), or AXI bresp errors and PCIe write errors due to bus master enable or function level reset (FLR), as well as RAM ECC errors. The first enabled error is logged in the DMA engine. Refer to the Memory Mapped Engine error logs. If an error occurs on the read, the DMA write is aborted if possible. If the error was detected when pulling write data from RAM, it is not possible to abort the request. Instead invalid data parity is generated to ensure the destination is aware of the problem. After the descriptor which encountered the error has gone through the DMA engine, it proceeds to generate status updates with an error indication. As with descriptor errors, it results in the queue being invalidated.

AXI Memory Mapped Descriptor for H2C and C2H (32B)

Table: AXI Memory Mapped Descriptor Structure for H2C and C2H

Bit	Bit Width	Field Name	Description
[255:192]	64	reserved	Reserved
[191:128]	64	dst_addr	Destination Address
[127:80]	36	reserved	Reserved
[79:64]	16	lengthInByte	Read length in byte (64K-1)
[63:0]	64	src_addr	Source Address

Internal mode memory mapped DMA must configure the descriptor queue to be 32B and follow the above descriptor format. In bypass mode, the descriptor format is defined by the user logic, which must drive the H2C or C2H MM bypass input port.

AXI Memory Mapped Writeback Status Structure for H2C and C2H

The MM writeback status register is located after the last entry of the (H2C or C2H) descriptor.

Table: AXI Memory Mapped Writeback Status Structure for H2C and C2H

Bit	Bit Width	Field Name	Description
[63:48]	16	reserved	Reserved
[47:32]	16	pidx	Producer Index at time of writeback
[31:16]	16	cidx	Consumer Index
[15:2]	14	reserved	Reserved

Bit	Bit Width	Field Name	Description
[1:0]	2	err	Error bit 1: Descriptor fetch error bit 0: DMA error

Stream Mode DMA

H2C Stream Engine

The H2C Stream Engine is responsible for transferring streaming data from the host and delivering it to the user logic. The H2C Stream Engine operates on H2C stream descriptors. Each descriptor specifies the start address and the length of the data to be transferred to the user logic. The H2C Stream Engine parses the descriptor and issues read requests to the host over PCIe, splitting the read requests at the MRRS boundary. There can be up to 256 requests outstanding in the H2C Stream Engine to hide the host read latency. The H2C Stream Engine implements a re-ordering buffer of 32 KB to re-order the TLPs as they come back. Data is issued to the user logic in order of the requests sent to PCIe.

If the status descriptor is enabled in the associated H2C context, the engine could additionally send a status write back to host once it is done issuing data to the user logic.

Internal and Bypass Modes

Each queue in QDMA Subsystem for PCIe can be programmed in either of the two H2C Stream modes: internal and bypass. This is done by specifying the mode in the queue context. The H2C Stream Engine knows whether the descriptor being processed is for a queue in internal or bypass mode.

The following figures show the internal mode and bypass mode flows.

Figure: H2C Internal Mode Flow





Figure: H2C Bypass Mode Flow

For a queue in the Internal mode, after the descriptor is fetched from the host it is fed straight to the H2C Stream Engine for processing. In this case, a packet of data cannot span over multiple descriptors. Thus for a queue in internal mode, each descriptor generates exactly one AXI4-Stream packet on the QDMA H2C AXI4-Stream output. If the packet is present in host memory in non-contiguous space, then it has to be defined by more than one descriptor and this requires that the queue be programmed in bypass mode.

In the Bypass mode, after the descriptors are fetched from the host they are sent straight to the user logic using the QDMA bypass output port. The QDMA does not parse these descriptors at all. The user logic can store these descriptors and then send the required information from these descriptors back to QDMA using the QDMA H2C Stream descriptor bypass-in interface. Using this information, the QDMA constructs descriptors which are then fed to the H2C Stream Engine for processing.

When `fcrd_en` is enabled in the software context, DMA will wait for the user application to provide credits, Credit return in the figure above. When `fcrd_en` is not set, the DMA uses a pointer update, fetches descriptors and sends the descriptor out. The user application should not send in credits. Credit return in the above figure does not apply in this case.

The following are the advantages of using the bypass mode:

- The user logic can have a custom descriptor format. This is possible because QDMA Subsystem for PCIe does not parse descriptors for queues in bypass mode. The user logic parses these descriptors and provides the information required by the QDMA on the H2C Stream bypass-in interface.
- Immediate data can be passed from the software to the user logic without DMA operation.
- The user logic can do traffic management by sending the descriptors to the QDMA when it is ready to sink all the data. Descriptors can be cached in local RAM.
- Perform address translation.

There are some requirements imposed on the user logic when using the bypass mode. Because the bypass mode allows a packet to span multiple descriptors, the user logic needs to indicate to QDMA which descriptor marks the Start-Of-Packet (SOP) and which marks the End-Of-Packet (EOP). At the QDMA H2C Stream bypass-in interface, among other pieces of information, the user logic needs to provide: Address, Length, SOP, and EOP. It is required that once the user logic feeds SOP descriptor information into QDMA, it must eventually feed EOP descriptor information also. Descriptors for these multi-descriptor packets must be fed in sequentially. Other descriptors not belonging to the packet must not be interleaved within the multi-descriptor packet. The user logic must accumulate the descriptors up to the EOP descriptor, before feeding them back to QDMA. Not doing so can result in a hang. The QDMA will generate a TLAST at the QDMA H2C AXI4-Stream data output once it issues the last beat for the EOP descriptor. This is guaranteed because the user is required to submit the descriptors for a given packet sequentially.

The H2C stream interface is shared by all the queues, and has the potential for a head of line blocking issue if the user logic does not reserve the space to sink the packet. Quality of service can be severely affected if the packet sizes are large. The Stream engine is designed to saturate PCIe for packet sizes as low as 128B, so AMD recommends that you restrict the packet size to be host page size or maximum transfer unit as required by the user application.

A performance control provided in the H2C Stream Engine is the ability to stall requests from being issued to the PCIe RQ/RC if a certain amount of data is outstanding on the PCIe side as seen by the H2C Stream Engine. To use this feature, the SW must program a threshold value in the H2C_REQ_THROT (0xE24) register. After the H2C Stream Engine has more data outstanding to be delivered to the user logic than this threshold, it stops sending further read requests to the PCIe RQ/RC. This feature is disabled by default and can be enabled with the H2C_REQ_THROT (0xE24) register. This feature helps improve the C2H Stream performance, because the H2C Stream Engine can make requests at a much faster rate than the C2H Stream Engine. This can potentially use up the PCIe side resources for H2C traffic which results in C2H traffic suffering. The H2C_REQ_THROT (0xE24) register also allows the SW to separately enable and program the threshold of the maximum number of read requests that can be outstanding in the H2C Stream engine. Thus, this register can be used to individually enable and program the thresholds for the outstanding requests and data in the H2C Stream engine.

H2C Stream Descriptor (16B)

Table: H2C Descriptor Structure

Bit	Bit Width	Field Name	Description
[127:96]	32	addr_h	Address High. Higher 32 bits of the source address in Host
[95:64]	32	addr_l	Address Low. Lower 32 bits of the source address in Host
[63:48]	16	reserved	Reserved
[47:32]	16	len	Packet Length. Length of the data to be fetched for this descriptor. This is also the packet length since in internal mode, a packet cannot span multiple descriptors. The maximum length of the packet can be 64K-1 bytes.
[31:0]	32	metadata	Metadata. QDMA passes this field on the H2C-ST TUSER along with the data on every beat. For a queue in internal mode, it can be used to pass messages from SW to user logic along with the data.

This H2C descriptor format is only applicable for internal mode. For bypass mode, the user logic can define its own format as needed by the user application.

Descriptor Metadata

Similar to bypass mode, the internal mode also provides a mechanism to pass information directly from the software to the user logic. In addition to address and length, the H2C Stream descriptor also has a 32b metadata field. This field is not used by the QDMA Subsystem for PCIe for the DMA operation. Instead, it is passed on to the user logic on the H2C AXI4-Stream tuser on every beat of the packet. Passing metadata on the tuser is not supported for a queue in bypass mode and consequently there is no input to provide the metadata on the QDMA H2C Stream bypass-in interface.

Zero Length Descriptor

The length field in a descriptor can be zero. In this case, the H2C Stream Engine will issue a zero byte read request on PCIe. After the QDMA receives the completion for the request, the H2C Stream Engine will send out one beat of data with `tlast` on the QDMA H2C AXI4-Stream interface. The zero byte packet will be indicated on the interface by setting the `zero_b_dma` bit in the tuser. The user logic must set both the SOP and EOP for a zero byte descriptor. If not

done, an error will be flagged by the H2C Stream Engine.

H2C Stream Status Descriptor Writeback

When feeding the descriptor information on the bypass input interface, the user logic can request the QDMA Subsystem for PCIe to send a status write back to the host when it is done fetching the data from the host. The user logic can also request that a status be issued to it when the DMA is done. These behaviors can be controlled using the `sdi` and `mrkr_req` inputs in the bypass input interface.

The H2C writeback status register is located after the last entry of the H2C descriptor list.

Note: The format of the H2C-ST status descriptor written to the descriptor ring is different from that written into the interrupt coalesce entry.

Table: AXI4-Stream H2C Writeback Status Descriptor Structure

Bit	Bit Width	Field Name	Description
[63:32]	32	reserved	Reserved
[47:32]	16	pidx	Producer Index
[31:16]	16	cidx	Consumer Index
[15:2]	14	reserved	Reserved (Producer Index)
[1:0]	2	error	Error 0x0 : No Error 0x1 : Descriptor or data error was encountered on this queue 0x2 and 0x3 : Reserved

Related Information

[QDMA Descriptor Bypass Input Ports](#)

H2C Stream Data Aligner

The H2C engine has a data aligner that aligns the data to zero Bytes (0B) boundary before issuing it to the user logic. This allows the start address of a descriptor to be arbitrarily aligned and still receive the data on the H2C AXI4-Stream data bus without any holes at the beginning of the data. The user logic can send a batch of descriptors from SOP to EOP with arbitrary address and length alignments for each descriptor. The aligner will align and pack the data from the different descriptors and will issue a continuous stream of data on the H2C AXI4-Stream data bus. The `tlast` on that interface will be asserted when the last beat for the EOP descriptor is being issued.

[Handling Descriptors With Errors](#)

If an error is encountered while fetching a descriptor, the QDMA Descriptor Engine flags the descriptor with error. For a queue in internal mode, the H2C Stream Engine handles the error descriptor by not performing any PCIe or DMA activity. Instead, it waits for the error descriptor to pass through the pipeline and forces a writeback after it is done. For a queue in bypass mode, it is the responsibility of the user logic to not issue a batch of descriptors with an error descriptor. Instead, it must send just one descriptor with error input asserted on the H2C Stream bypass-in interface and set the SOP, EOP, no_dma signal, and sdi or mrkr - req signal to make the H2C Stream Engine send a writeback to Host.

Handling Errors in Data From PCIe

If the H2C Stream Engine encounters an error coming from PCIe on the data, it keeps the error sticky across the full packet. The error is indicated to the user on the err bit on the H2C Stream Data Output. Once the H2C Stream sends out the last beat of a packet that saw a PCIe data error, it also sends a Writeback to the Software to inform it about the error.

C2H Stream Engine

The C2H Stream Engine DMA writes the stream packets to the host memory into the descriptor provided by the host driver through the C2H descriptor queue.

The Prefetch Engine is responsible for calculating the number of descriptors needed for the DMA that is writing the packet. The buffer size is fixed per queue basis. For internal and cached bypass mode, the prefetch module can fetch up to 512 descriptors for a maximum of 64 different queues at any given time.

The Prefetch Engine also offers low latency feature pfch_en = 1, where the engine can prefetch up to qdma_c2h_pfch_cfg.num_pfch descriptors upon receiving the packet, so that subsequent packets can avoid the PCIe latency.

The QDMA requires software to post full ring size so the C2H stream engine can fetch the needed number of descriptors for all received packets. If there are not enough descriptors in the descriptor ring, the QDMA will stall the packet transfer. For performance reasons, the software is required to post the PIDX as soon as possible to ensure there are always enough descriptors in the ring.

C2H stream packet data length is limited to 31 * C2H buffer size. C2H buffer size can be programmed from 0xAB0 to 0xAEC address, for details refer to [qdma_v5_1_pf_registers.csv](#) file.

In older versions (such as 2018.3), C2H stream packet data length was limited to 7 * C2H buffer size.

C2H Stream Descriptor (8B)

Table: AXI4-Stream C2H Descriptor Structure

Bit	Bit Width	Field Name	Description
[63:0]	64	addr	Destination Address

C2H Prefetch Engine

The prefetch engine interacts between the descriptor fetch engine and C2H DMA write engine to pair up the descriptor and its payload.

Table: C2H Prefetch Context Structure

Bit	Bit Width	Field Name	Description
[45]	1	valid	Context is valid
[44:29]	16	sw_crdt	Software credit This field is written by the hardware for internal use. The software must initialize it to 0 and then treat it as read-only.
[28]	1	pfch	Queue is in prefetch This field is written by the hardware for internal use. The software must initialize it to 0 and then treat it as read-only.
[27]	1	pfch_en	Enable prefetch
[26]	1	err	Error detected on this queue During the descriptor per-fetch process, if there are any errors detected, they are logged here. This is per queue basis.
[25:8]	18	reserved	Reserved
[7:5]	3	port_id	Port ID
[4:1]	4	buf_size_idx	Buffer size index
[0]	1	bypass	C2H bypass mode, set this bit for simple bypass mode.

C2H Stream Modes

The C2H descriptors can be from the descriptor fetch engine or C2H bypass input interfaces. The descriptors from the descriptor fetch engine are always in cache mode. The prefetch engine keeps the order of the descriptors to pair with the C2H data packets from the user. The descriptors from the C2H bypass input interfaces have one interface for both simple mode and cache mode (note that both simple bypass and cache bypass use the same interface). For simple mode, the user application keeps the order of the descriptors to pair with the C2H data packets. For cache mode, the prefetch engine keeps the order of the descriptors to pair with the C2H data packet from the user.

The prefetch context has a bypass bit. When it is 1 ' b1, the user application sends the credits for

the descriptors. When it is 1'b0, the prefetch engine handles the credits for the descriptors. The descriptor context has a bypass bit. When it is 1'b1, the descriptor fetch engine sends out the descriptors on the C2H bypass output interface. The user application can convert it and later loop it back to the QDMA Subsystem for PCIe on the C2H bypass input interface. When the bypass context bit is 1'b0, the descriptor fetch engine sends the descriptors to the prefetch engine directly.

There is a 1K descriptor entry buffer to take in descriptors from bypass input ports. 1K deep buffer is shared with all the Qs.

On a per queue basis, three modes are supported.

- Cache Internal Mode
- Cache Bypass Mode
- Simple Bypass Mode

The selection between Simple Bypass Mode and Cache Bypass Mode is done by setting the bypass bits in Software Descriptor Context and C2H Prefetch Context as shown in the following table.

Note: If you already have the descriptor cached on the device, there is no need to fetch one from the host and you should follow the simple bypass mode for the C2H Stream application. In simple bypass mode, do not provide credits to fetch the descriptor, and instead, you need to send in the descriptor on the descriptor bypass interface.

Note: AXI4-Stream C2H *Simple Bypass mode* and *Cache Bypass mode* both use same bypass in ports (c2h_byp_in_st_csh_* ports).

Table: C2H Stream Modes

	c2h_byp_in port	desc_byp software descriptor bypass	C2H prefetch context
Simple bypass mode	c2h_byp_in_st_csh_*	1	1
Cache bypass mode	c2h_byp_in_st_csh_*	1	0
Cache internal mode	N/A	0	0

Simple Bypass Mode

For simple bypass mode, the descriptor fetch engine sends the descriptors out on the C2H bypass out interface. The user application converts the descriptor and loops it back to the QDMA on the simple mode C2H bypass input interface. The user application sends the credits for the descriptors, and it also keeps the order of the descriptors.

For simple bypass transfer to work, a prefetch tag is needed and it can be fetched from the QDMA IP.

The user application must request a prefetch tag before sending any traffic for a simple bypass queue through the C2H ST engine. Invalid queues or non-bypass queues should not request any tags using this method, as it might reduce performance by freezing tags that never get used. The prefetch tag needs to be reserved upfront before any traffic can start. One prefetch tag per target host is required. In most applications, one prefetch tag for a host is needed. In Simple Bypass mode, the tag is not tied to any descriptor ring. For the queues that share the same prefetch tag, the data and descriptors need to come in the same order. For Simple Bypass, the data and descriptors are both controlled by the user, so they need to guarantee the order is maintained.

For example when the data stream has packets in the order of Q0, Q1, and Q2, on descriptor input, you can not send Q1, Q2, Q0 etc. The order needs to be maintained.

The user application writes to the MDMA_C2H_PFCH_BYP_QID (0x1408) register with the qid for a simple bypass queue, then reads from MDMA_C2H_PFCH_BYP_TAG (0x140C) register to retrieve the corresponding prefetch tag. This tag must be driven with all bypass_in descriptors for as long as the current qid is valid. If a current qid is invalidated, a new prefetch tag must be requested with a valid qid.

Prefetched tag must be assigned to input port c2h_byp_in_st_csh_pfch_tag[6:0] for all transfers. The prefetch tag points to the CAM that stores the active queues in the prefetch engine. Also the qid that was used to prefetch tag needs to be used as the qid for all simple bypass packets. Assign the qid to dma<n>_s_axis_c2h_ctrl_qid.

The steps to fetch the prefetch tag are as follows:

1. Software instruction:

- a. Initialize a queue (qid).
- b. Write to MDMA_C2H_PFCH_BYP_QID 0x1408 with valid qid.
- c. Read MDMA_C2H_PFCH_BYP_TAG 0x140C to obtain the prefetch tag.
- d. The prefetch tag and the qid that was used to fetch the tag should be used for all simple bypass packets. This information needs to be communicated to the user side.


2. User side:

- a. Assign the qid used to fetch the tag to dma<n>_s_axis_c2h_ctrl_qid.
- b. Assign the actual qid of the packet transfer to
dma<n>_s_axis_c2h_cmpt_ctrl_qid.
- c. Assign the prefetch tag value to c2h_byp_in_st_csh_pfch_tag.
- d. Assign the actual qid of the packet transfer to c2h_byp_in_st_csh_qid.

Note: The c2h_byp_in_st_csh_pfch_tag[6:0] port can have the same prefetch_tag for as long as the original qid is valid.

Simple bypass flow shown below does not include fetch of the "prefetch_tag".

Figure: C2H Simple Bypass Mode Flow



Note: No sequence is required between descriptor bypass in, data payload and completion packets.

If you already have descriptors, there is no need to update the pointers or provide credits. Instead, send the descriptors in the descriptor bypass interface, and send the data and Completion (CMPT) packets.

When simple bypass mode is selected, the queue that is used to fetch the prefetch ID, acts like management Q and it controls the buffer sizes.

The buffer size that is set for this management Q is used for all the queues, irrespective of the buffer size set for other queues. In the simple bypass mode, you provide descriptors and data packets, so the buffer size should be set properly to accommodate all packet sizes in all queues. Users need to make sure to send adequate descriptors for each data packets. Number of descriptors = packet size/Management Q buffer size, and this applies for all queues.

Cache Bypass Mode

For Cache Bypass mode, the descriptor fetch engine sends the descriptors out on the C2H bypass output interface. The user application converts the descriptor and loops it back to the QDMA on the cache mode C2H bypass input interface. The prefetch engine sends the credits for the descriptors, and it keeps the order of the descriptors.


For Cache Internal mode, the descriptor fetch engine sends the descriptors to the prefetch engine. The prefetch engine sends out the credits for the descriptors and keeps the order of the descriptors. In this case, the descriptors do not go out on the C2H bypass output and do not come back on the C2H bypass input interfaces. In Cache Internal mode prefetch tag is

maintained by the IP internally.

In Cache Bypass or Cache Internal mode, prefetch mode can be turned on which prefetches the descriptor and reduces transfer latency significantly. When prefetch mode is enabled, the user application can not send credits as input in QDMA Descriptor Credit input ports. Credits for all queues are maintained by prefetch engine.

In cache bypass mode, prefetch tag is maintained by the IP internally. Signal `c2h_byp_out_pfch_tag[6:0]` should be looped back as an input `c2h_byp_in_st_csh_pfch_tag[6:0]`. The prefetch tag points to the cam that stores the active queues in the prefetch engine.

Figure: C2H Cache Bypass Mode Flow



Note: No sequence is required between payload and completion packets.

Related Information

[QDMA Descriptor Bypass Input Ports](#)

[QDMA Descriptor Bypass Output Ports](#)

[QDMA Descriptor Credit Input Ports](#)

[Queue Status Ports](#)

C2H Stream Packet Type

The following are some of the different C2H stream packets.

Regular Packet

The regular C2H packet has both the data packet and Completion (CMPT) packet. They are a one-to-one match.

The regular C2H data packet can be multiple beats.

- `dma<n>_s_axis_c2h_ctrl_qid` = C2H descriptor queue ID.
- `dma<n>_s_axis_c2h_ctrl_len` = length of the packet.
- `dma<n>_s_axis_c2h_mty` = empty byte should be set in last beat.
- `dma<n>_s_axis_c2h_ctrl_has_cmpt` = 1'b1. This data packet has a corresponding CMPT packet.

The regular C2H CMPT packet is one beat.

- `dma<n>_s_axis_c2h_cmpt_ctrl_qid` = Completion queue ID of the packet. This can be different from the C2H descriptor QID.
- `dma<n>_s_axis_c2h_cmpt_ctrl_cmpt_type` = HAS_PLD. This completion packet has a corresponding data packet.
- `dma<n>_s_axis_c2h_cmpt_ctrl_wait_pld_pkt_id` = This completion packet has to wait for the data packet with this ID to be sent before the CMPT packet can be sent.

When the user application sends the data packet, it must count the packet ID for each packet.

The first data packet has a packet ID of 1, and it increments for each data packets.

For the regular C2H packet, the data packet and the completion packet is a one-to-one match.

Therefore, the number of data packets with `dma<n>_s_axis_c2h_ctrl_has_cmpt` as 1'b1 should be equal to the number of CMPT packet with

`dma<n>_s_axis_c2h_cmpt_ctrl_cmpt_type` as HAS_PLD.


The QDMA Subsystem for PCIe has a shallow completion input FIFO of depth 2. For better performance, add FIFO for completion input as shown in the diagram below. Depth and width of the FIFO depends on the use case. Width is dependent on the largest CMPT size for the application, and depth is dependent on performance needs. For best performance for 64 Byte CMPT, a depth of 512 is recommended.

When the user application sends the data payload, it counts every packet. The first packet starts with a `pkt_pld_id` of 1. The second packet has a `pkt_pld_id` of 2, and so on. It is a 16-bits counter once the count reaches 16'hffff it wraps around to 0 and count forward.

The user application defines the CMPT type.

- If the `dma<n>_s_axis_c2h_cmpt_ctrl_cmpt_type` is `HAS_PLD`, the CMPT has a corresponding data payload. The user application must place `pkt_pld_id` of that packet in the `dma<n>_s_axis_c2h_cmpt_ctrl_wait_pld_pkt_id` field. The DMA only sends out this CMPT after it sends out the corresponding data payload packet.
- If the `dma<n>_s_axis_c2h_cmpt_ctrl_cmpt_type` is `NO_PLD_NO_WAIT`, the CMPT does not have any data payload, and it does not need to wait for payload. Then the DMA sends out this CMPT.
- If the `dma<n>_s_axis_c2h_cmpt_ctrl_cmpt_type` is `NO_PLD_BUT_WAIT`, the CMPT does not have a corresponding data payload packet. The CMPT must wait for a particular data payload packet before the CMPT is sent out. Therefore, the user application must place the `pld_pkt_id` of that particular data payload into the `dma<n>_s_axis_c2h_cmpt_ctrl_wait_pld_pkt_id` field. The DMA does not send out the CMPT until the data payload with that `pld_pkt_id` is sent out.

Figure: CMPT Input FIFO



X22048-120718

Immediate Data Packet

The user application can have a packet that only writes to the Completion Ring without having a corresponding data packet transfer to the host. This type of packet is called immediate data packet. For the immediate data packet, the QDMA does not send the data payload, but it writes to the CMPT Queue. The immediate packet does not consume a descriptor.

For the immediate data packet, the user application only sends the CMPT packet to the DMA, and it does not send the data packet.

The following is the setting of the immediate completion packet. There is no corresponding data packet.

In some applications, the immediate completion packet does not need to wait for any data packet. But in some applications, it might still need to wait for the data payload packet. When the completion type is `NO_PLD_NO_WAIT`, the completion packet can be sent out without waiting for any data packet. When the completion type is `NO_PLD_BUT_WAIT`, the completion packet must specify the data packet ID that it needs to wait for.

- `dma<n>_s_axis_c2h_cmpt_user_cmpt_type = NO_PLD_NO_WAIT or NO_PLD_BUT_WAIT.`
- `dma<n>_s_axis_c2h_cmpt_ctrl_wait_pld_pkt_id = Do not increment packet count.`

Marker Packet

The C2H Stream Engine of the QDMA provides a way for the user application to insert a marker into the QDMA along with a C2H packet. This marker then propagates through the C2H Engine pipeline and comes out on the Queue status port interface. The marker is inserted by setting the marker bit in the C2H Stream packet. The marker response is indicated by the QDMA to the user application by setting the `qsts_out_op[7:0] = 0x0 (CMPT Marker response)` bits on the Queue status ports. For a marker packet, QDMA does not send out a payload packet but still writes to the Completion Ring. Not all marker responses are generated because of a corresponding marker request. The QDMA sometimes generates marker responses when it encounters exceptional events. See the following section for details about when QDMA internally generates marker responses.

The primary purpose of giving the user application the ability of sending in a marker into QDMA is to determine when all the traffic prior to the marker has been flushed. This can be used in the shut down sequence in the user application. Although not a requirement, the marker can be sent by the user application with the `user_trig` bit set when sending in the marker into QDMA. This allows the QDMA to generate an interrupt and truly ensures that all traffic prior to the marker is flushed out. The QDMA Completion Engine takes the following actions when it receives a marker from the user application:

The DMA also has an option not to send completion information to completion ring during a Marker packet. Port `dma<n>_s_axis_c2h_cmpt_ctrl_no_wrb_marker` can be set during a marker packet. This option disables any write to completion ring when that Marker packet is generated. When this signal is set for a Marker packet, the DMA has no data or completion transfers, but responds with marker response on `qsts_out_op[7:0]`.

- Sends the Completion that came along with the marker to the C2H Stream Completion Ring.
- Sends lower 24bits of completion data to the Queue status data port `qsts_out_data[26:3]`.
- Generates Status Descriptor if enabled (if `user_trig` was set when marker was inserted).
- Generates an Interrupt if enabled and not outstanding.
- Sends the marker response. If an Interrupt was not sent due to it being enabled but outstanding, the `retry_marker_req` bit in the marker response is set to inform the user that an Interrupt could not be sent for this marker request. See the Queue status ports interface description for details of these fields.

The marker packet has both the data packet and CMPT packet. They are one-to-one match. The following is the setting of the data packet with marker:

- 1 beat of data
- `dma<n>_s_axis_c2h_ctrl_marker = 1'b1`
- `dma<n>_s_axis_c2h_ctrl_len` = data width (for example, 64 if data width is 512 bits)
- `dma<n>_s_axis_c2h_mty = 0`
- `dma<n>_s_axis_c2h_ctrl_has_cmpt = 1'b1`

The following is the setting of the CMPT packet with marker:

- 1 beat of CMPT packet
- `dma<n>_s_axis_c2h_cmpt_ctrl_marker = 1'b1`
- `dma<n>_s_axis_c2h_cmpt_ctrl_cmpt_type = HAS_PLD`
- `dma<n>_s_axis_c2h_cmpt_ctrl_wait_pld_pkt_id` = This completion packet has to wait for the data payload packet with this ID to be sent before the CMPT packet is sent.
- `dma<n>_s_axis_c2h_cmpt_ctrl_no_wrb_marker = 1'b0`. If this is set to 1'b1 then no CMPT packet.

The immediate data packet and the marker packet do not consume the descriptor; instead, they write to the C2H Completion Ring. The software needs to size the C2H Completion Ring large enough to accommodate the outstanding immediate packets and the marker packets. When a marker request is received by the DMA and the completion ring is full, the marker response is sent out. However, because the completion ring is full, the completion entry is dropped and the queue is invalidated.

Zero Length Packet

The length of the data packet can be zero. On the input, the user needs to send one beat of data. The zero length packet consumes the descriptor. The QDMA sends out 1DW payload data. The following is the setting of the zero length packet:

- 1 beat of data
- `dma<n>_s_axis_c2h_ctrl_len = 0`
- `dma<n>_s_axis_c2h_mty = 0`

 **Note:** Zero Byte packets are not supported in Internal mode and Cache bypass mode. The QDMA might hang if zero byte packets are dropped due to not available descriptor. Zero Byte Packets are supported in Simple bypass mode.

Disable Completion Packet

The user application can disable the completion for a specific packet. The QDMA provides direct memory access (DMA) to the payload, but does not write to the C2H Completion Ring. The user application only sends the data packet to the DMA, and does not send the CMPT packet.

The following is the setting of the disable completion packet:

- `dma<n>_s_axis_c2h_ctrl_has_cmpt = 1'b0`

Related Information

[QDMA Descriptor Bypass Output Ports](#)

Handling Descriptors With Errors

If an error is encountered while fetching a descriptor (in pre-fetch or regular mode), the QDMA Descriptor Engine flags the descriptor with error. For a queue in internal mode, the C2H Stream Engine handles the error descriptor by not performing any PCIe or DMA activity. Instead, it waits for the error descriptor to pass through the pipeline and forces a writeback after it is done. For a queue in bypass mode, it is the responsibility of the user logic to not issue a batch of descriptors with an error descriptor. Instead, it must send just one descriptor with error input asserted on the C2H Stream bypass-in interface and set the SOP, EOP, no_dma signal, and sdi or mrkr_req signal to make the C2H Stream Engine send a writeback to Host.

Completion Engine

The Completion Engine writes the C2H AXI4-Stream Completion (CMPT) in the CMPT queue. The user application sends a CMPT packet and other information, such as, but not limited to, CMPT QID, and CMPT_TYPE to the QDMA Subsystem for PCIe. The QDMA uses this information to process the CMPT packet. The QDMA can be instructed to write the CMPT packet unchanged in the CMPT queue. Alternatively, the user application can instruct the QDMA to insert certain fields, like error and color, in the CMPT packet before writing it into the CMPT queue. Additionally, using the CMPT interface signals, the user application instructs the QDMA to order the writing of the CMPT packet in a specific way, relative to traffic on the C2H data input. Although not a requirement, a CMPT is typically used with a C2H queue. In such a case, the CMPT is used to inform the SW that a certain number of C2H descriptors are used up by the DMA of C2H data. This allows the SW to reclaim the C2H descriptors. A CMPT can also be used without a corresponding C2H DMA operation, in which case, it is known as Immediate Data. The user-defined portion of the CMPT packet typically needs to specify the length of the data packet transferred and whether or not descriptors were consumed as a result of the data packet transfer. Immediate and marker type packets do not consume any descriptors. The exact contents of the user-defined data are up to the user to determine.

☞ **Note:** Maximum buffer size register 0xB50 bits[31:26] is programmed to 0 (default value). This value might result in an overflow depending on the simulator or the synthesis tool used. To avoid overflow, set 0xB50 bits[31:26] to maximum value of 63.

☞ **Note:** The calculation of the completion ring size must account for completion data, immediate data, and marker packets. Therefore, you must assign the completion ring size accurately.

Completion Context Structure

The completion context is used by the Completion Engine.

Table: Completion Context Structure Definition

Bit	Bit Width	Field Name	Description
[256:183]	17		Reserved. Initialize to 0.
[182:180]	3	port_id	<p>Port ID. The Completion Engine checks the port_id of events received at its input to the port_id configured here. If the check fails, the input is dropped, and an error is logged in the C2H_ERR_STAT register. The following are checked for port_id:</p> <ol style="list-style-type: none"> 1. All events on the s_axis_c2h_cmpt interface. These include CMPTs, Immediate data, markers and VirtIO control messages. 2. CMPT CIDX pointer updates (checked only when the update is coming from the AXI side).
[179]	1		Reserved. Initialize to 0's
[178:175]	4	baddr4_low	Since the minimum alignment supported is 64B in this case, this field must be 0
[174:147]	28		Reserved. Initialize to 0's
[146]	1	dir_c2h	<p>DMA direction is C2H. The CMPT engine can be used to manage the completion/used ring of a C2H as well as an H2C queue.</p> <p>0x0: DMA direction is H2C 0x1: DMA direction is C2H</p>
[145]	1		Reserved. Initialize to 0
[144]	1	dis_int_on_vf	Disable interrupt with VF
[143]	1	int_aggr	Interrupt Aggregation Set to configure the QID in interrupt aggregation mode
[142:132]	11	vec	Interrupt Vector Number
131	1	at	Address Translation This bit is used to determine whether the queue addresses are translated or

Bit	Bit Width	Field Name	Description
			untranslated. This information is sent to the PCIe on CMPT and Status writes. 0: Address is untranslated 1: Address is translated
130	1	ovf_chk_dis	Completion Ring Overflow Check Disable If set, then the CMPT Engine does not check whether writing a completion entry in the completion ring overflows the ring or not. The result is that QDMA invariably sends out Completions without first checking if it is going to overflow the Completion Ring and not take any actions that it normally takes when it encounters a Completion Ring overflow scenario. It is up to the software and user logic to negotiate and ensure that they do not cause a Completion Ring overflow
[129]	1	full_upd	Full Update If reset, the all fields other than the CIDX of a Completion-CIDX-update are ignored. Only the CIDX field is copied from the update to the context. If set, the Completion CIDX update can update the following fields in this context: <ul style="list-style-type: none"> • timer_ix • counter_ix • trig_mode • en_int • en_stat_desc
[128]	1	timer_running	If set, it indicates that a timer is running on this queue. This timer is for the purpose of CMPT interrupt moderation. Ideally, the software must ensure that there is no running timer on this QID before shutting the queue down. This is a field used internally by the hardware. The software must initialize it to 0 and then treat it as read-only.

Bit	Bit Width	Field Name	Description
[127]	1	user_trig_pend	If set, it indicates that a user logic initiated interrupt is pending to be generated. The user logic can request an interrupt through the s_axis_c2h_cmpt_ctrl_user_trig signal. This bit is set when the user logic requests an interrupt while another one is already pending on this QID. When the next Completion CIDX update is received by QDMA, this pending bit might or might not generate an interrupt depending on whether or not there are entries in the Completion ring waiting to be read. This is a field used internally by the hardware. The software must initialize it to 0 and then treat it as read-only.
[126:125]	2	err	Indicates that the Completion Context is in error. This is a field written by the hardware. The software must initialize it to 0 and then treat it as read-only. The following errors are indicated here: 0: No error. 1: A bad CIDX update from software was detected. 2: A descriptor error was detected. 3: A Completion packet was sent by the user logic when the Completion Ring was already full.
[124]	1	valid	Context is valid.
[123:108]	16	cidx	Current value of the hardware copy of the Completion Ring Consumer Index.
[107:92]	16	pidx	Completion Ring Producer Index. This is a field written by the hardware. The software must initialize it to 0 and then treat it as read-only.
[91:90]	2	desc_size	Completion Entry Size: 0: 8B 1: 16B 2: 32B

Bit	Bit Width	Field Name	Description
			3: 64B
[89:32]	58	baddr	64B aligned base address of Completion ring – bit [63:6].
[31:28]	4	qsize_idx	Completion ring size index. This index selects one of 16 register (offset 0x204 :0x240) which has different ring sizes.
[27]	1	color	Color bit to be used on Completion.
[26:25]	2	int_st	<p>Interrupt State:</p> <p>0: ISR</p> <p>1: TRIG</p> <p>This is a field used internally by the hardware. The software must initialize it to 0 and then treat it as read-only. When out of reset, the hardware initializes into ISR state, and is not sensitive to trigger events. If the software needs interrupts or status writes, it must send an initial Completion CIDX update. This makes the hardware move into TRIG state and as a result it becomes sensitive to any trigger conditions.</p>
[24:21]	4	timer_idx	Index to timer register for TIMER based trigger modes.
[20:17]	4	counter_idx	Index to counter register for COUNT based trigger modes.
[16:13]	4		Reserved. Initialize to 0
[12:5]	8	fnc_id	Function ID
[4:2]	3	trig_mode	<p>Interrupt and Completion Status Write Trigger Mode:</p> <p>0x0: Disabled</p> <p>0x1: Every</p> <p>0x2: reserved</p> <p>0x3: User</p> <p>0x4: User_Timer</p> <p>0x5: reserved</p>
[1]	1	en_int	Enable Completion interrupts.

Bit	Bit Width	Field Name	Description
[0]	1	en_stat_desc	Enable Completion Status writes.

Completion Status Structure

The Completion Status is located at the last location of Completion ring, that is, Completion Ring Base Address + (Size of the completion length (8,16,32) * (Completion Ring Size – 1)).

In order to make the QDMA Subsystem for PCIe write Completion Status to the Completion ring, Completion Status must be enabled in the Completion context. In addition to affecting Interrupts, the trigger mode defined in the Completion context also moderates the writing of Completion Statuses. Subject to Interrupt/Status moderation, a Completion Status can be written when either of the following happens:

1. A CMPT packet is written to the Completion ring.
2. A CMPT-CIDX update from the SW is received, and indicates that more Completion entries are waiting to be read.
3. The timer associated with the respective CMPT QID expires and is programmed in a timer-based trigger mode.

Table: AXI4-Stream Completion Status Structure

Bit	Bit Width	Field Name	Description
[63:37]	27		Reserved
[36:35]	2	error	Error. 0x0: No error 0x1 Bad CIDX update received 0x2: Descriptor error 0x3: CMPT ring overflow error
[34:33]	2	int_state	Interrupt State. 0: ISR 1: TRIG
[32]	1	color	Color status bit
[31:16]	16	cidx	Consumer Index (RO)
[15:0]	16	pidx	Producer Index

Completion Entry Structure

The size of a Completion (CMPT) Ring entry is 512-bits. This includes user defined data, an optional error bit, and an optional color bit. The user defined data has four size options: 8B, 16B, 32B and 64B. The bit locations of the optional error and color bits in the CMPT entry are

configurable individually. This is done by specifying the locations of these fields using the AMD Vivado™ IDE IP customization options while compiling the QDMA Subsystem for PCIe. There are seven color bit location options and eight error bit location options. The location is specified as an offset from the LSB bit of the Completion entry.

When the user application drives a Completion packet into the QDMA Subsystem for PCIe, it provides a `dma<n>_s_axis_cmpt_ctrl_col_idx[2:0]` value and a `dma<n>_s_axis_cmpt_ctrl_err_idx[2:0]` value at the interface. These indices are used by the QDMA Subsystem for PCIe to use the correct locations of the color and error bits. For example, if `dma<n>_s_axis_cmpt_ctrl_col_idx[2:0] = 0` and `dma<n>_s_axis_cmpt_ctrl_err_idx[2:0] = 1`, then the QDMA Subsystem for PCIe uses the C2H Stream Completion Color bits position option 0 for color location, and C2H Stream Completion Error bits position option 1 for error location. An index of seven for color or error signals implies that the DMA will not update the corresponding color or error bits when Completion entry is updated (those fields are ignored). The C2H Stream Completions bits options are set in the PCIe DMA Tab in the AMD Vivado™ IDE.

The error and color bit location values that are used at compile time are available for the software to read from the MMIO registers. There are seven registers for this purpose, `QDMA_C2H_CMPT_FORMAT` (0xBC4) to `QDMA_GLBL_ERR_MASK` (0x24C). Each of these registers holds one color and one error bit location.

- C2H Stream Completions bits option 0 for color bit location and option 0 for error bit location are available through the `QDMA_C2H_CMPT_FORMAT_0` register.
- C2H Stream Completions bits option 1 for color bit location and option 1 for error bit location are available through the `QDMA_C2H_CMPT_FORMAT_1` register.
- And so on.

Based on the CMPT data size selection (8, 16, 32 or 64 Bytes), the data in `s_axis_c2h_cmpt_tdata[511:0]` signal is registered in the completion entry as shown in the following table.

Table: Completion Entry Structure

Name	Size (Bits)	Index
User-defined bits for 64 Bytes settings	510-512	Depending on whether there are color and error bits present.
User-defined bits for 32 Bytes settings	254-256	Depending on whether there are color and error bits present.
User-defined bits for 16 Bytes settings	126-128	Depending on whether there are color and error bits present.
User-defined bits for 8 Bytes settings	62-64	Depending on whether there are color and error bits present.

Name	Size (Bits)	Index
Err		<p>The Error bit location is defined by registers QDMA_C2H_CMPT_FORMAT_0 (0xBC4) to QDMA_C2H_CMPT_FORMAT_6 (0xBDC). These register show color bit position that is user defined during IP generation. You can index into this register based on input CMPT ports</p> <p><code>dma<n>_s_axis_c2h_cmpt_ctrl_err_idx</code></p> <p>You can choose not to include err bit (index value 7). In such a case, user-defined data takes up that space</p>
Color		<p>The Color bit location is defined by registers QDMA_C2H_CMPT_FORMAT_0 (0xBC4) to QDMA_C2H_CMPT_FORMAT_6 (0xBDC). These register show color bit position that is user defined during IP generation. You can index into this register based on input CMPT ports</p> <p><code>dma<n>_s_axis_c2h_cmpt_ctrl_col_idx</code></p> <p>If you do not include a color bit (index value 7), the user-defined data takes up that space.</p>

Related Information

[PCIe DMA Tab](#)

[QDMA_CSR \(0x0000\)](#)

Completion Input Packet

The user application sends the CMPT packet to the QDMA.

The CMPT packet and data packet do not require a one-to-one match. For example, the immediate data packet only has the CMPT packet, and does not have the data packet. The disable completion packet only has the data packet and does not have the CMPT packet.

Each CMPT packet has a CMPT ID. It is the ID for the associated CMPT queue. Each CMPT queue

has a CMPT Context. The driver sets up the mapping of the C2H descriptor queue to the CMPT queue. There also can be a CMPT queue that is not associated to a C2H queue. The following is the CMPT packet from the user application.

Table: CMPT Input Packet

Name	Size	Index
Data (s_axis_c2h_cmpt_tdata[511:0])	512 bits	[511:0]

The CMPT packet has four options (8, 16, 32 or 64 Bytes). It has one pump of data with 512 bits.

Completion Status/Interrupt Moderation

The QDMA Subsystem for PCIe provides a means to moderate the Completion interrupts and Completion Status writes on a per queue basis. The software can select one out of five modes for each queue. The selected mode for a queue is stored in the QDMA Subsystem for PCIe in the Completion ring context for that queue. After a mode has been selected for a queue, the driver can always select another mode when it sends the completion ring CIDX update to the QDMA. The Completion interrupt moderation is handled by the Completion engine. The Completion engine stores the Completion ring contexts of all the queues. It is possible to individually enable or disable the sending of interrupts and Completion Statuses for every queue and this information is present in the Completion ring context. It is worth mentioning that the modes being described here moderate not only interrupts but also Completion Status writes. Also, since interrupts and Completion Status writes can be individually enabled/disabled for each queue, these modes work only if the interrupt/Completion Status is enabled in the Completion context for that queue.

The QDMA Subsystem for PCIe keeps only one interrupt outstanding per queue. This policy is enforced by QDMA even if all other conditions to send an interrupt are met for the mode. The way the QDMA Subsystem for PCIe considers an interrupt serviced is by receiving a CIDX update for that queue from the driver.

The basic policy followed in all the interrupt moderation modes is that when there is no interrupt outstanding for a queue, the QDMA Subsystem for PCIe keeps monitoring the trigger conditions to be met for that mode. Once the conditions are met, an interrupt is sent out. While the QDMA subsystem is waiting for the interrupt to be served, it remains sensitive to interrupt conditions being met and remembers them. When the CIDX update is received, the QDMA subsystem evaluates whether the conditions are still being met. If they are still being met, another interrupt is sent out. If they are not met, no interrupt is sent out and the QDMA resumes monitoring for the conditions to be met again.

The interrupt moderation modes that the QDMA subsystem provides are not necessarily precise. Thus, if the user application sends two CMPT packets with an indication to send an interrupt, it is not necessary that two interrupts are generated. The main reason for this behavior is that when the driver is interrupted to read the Completion ring, and it is under no obligation to read exactly up to the Completion for which the interrupt was generated. Thus, the driver might not read up to

the interrupting Completion, or it might even read beyond the interrupting Completion descriptor if there are valid descriptors to be read there. This behavior requires the QDMA Subsystem for PCIe to re-evaluate the trigger conditions every time it receives the CIDX update from the driver.

The detailed description of each mode is given below:

TRIGGER_EVERY

This mode is the most aggressive in terms of interruption frequency. The idea behind this mode is to send an interrupt whenever the completion engine determines that an unread completion descriptor is present in the Completion ring.

TRIGGER_USER

The QDMA Subsystem for PCIe provides a way to send a CMPT packet to the subsystem with an indication to send out an interrupt when the subsystem is done sending the packet to the host. This allows the user application to perform interrupt moderation when the TRIGGER_USER mode is set.

TRIGGER_USER_COUNT

This mode allows the QDMA Subsystem for PCIe is sensitive to either of two triggers. One of these triggers is sent by the user along with the CMPT packet. The other trigger is the presence of more than a programmed threshold of unread Completion entries in the Completion Ring, as seen by the hardware. This threshold is driver programmable on a per-queue basis. The QDMA evaluates whether or not to send an interrupt when either of these triggers is detected. As explained in the preceding sections, other conditions must be satisfied in addition to the triggers for an interrupt to be sent.

TRIGGER_USER_TIMER

In this mode, the QDMA Subsystem for PCIe is sensitive to either of two triggers. One of these triggers is sent by the user along with the CMPT packet. The other trigger is the expiration of the timer that is associated with the CMPT queue. The period of the timer is driver programmable on a per-queue basis. The QDMA evaluates whether or not to send an interrupt when either of these triggers is detected. As explained in the preceding sections, other conditions must be satisfied in addition to the triggers for an interrupt to be sent. For more information, see [Completion Timer](#).

TRIGGER_USER_TIMER_COUNT

This mode allows the QDMA Subsystem for PCIe is sensitive to any of three triggers. The first trigger is sent by the user along with the CMPT packet. The second trigger is the expiration of the timer that is associated with the CMPT queue. The period of the timer is driver programmable on a per-queue basis. The third trigger is the presence of more than a programmed threshold of unread Completion entries in the Completion Ring, as seen by the hardware. This threshold is driver programmable on a per-queue basis. The QDMA evaluates whether or not to send an interrupt when any of these triggers is detected. As explained in the preceding sections, other conditions must be satisfied in addition to the triggers for an interrupt to be sent.

TRIGGER_DIS

In this mode, the QDMA Subsystem for PCIe does not send Completion interrupts in spite of them being enabled for a given queue. The only way that the driver can read the completion ring in this case is when it regularly polls the ring. The driver requires to make use of the color bit feature provided in the Completion ring when this mode is set as this mode also disables the sending of any Completion Status descriptors to the Completion ring.

When a queue is programmed in TRIGGER_USER_TIMER_COUNT mode, the software can choose to not read all the Completion entries available in the Completion ring as indicated by an interrupt (or a Completion Status write). In such a case, the software can give a Completion CIDX update for the partial read. This works because the QDMA restarts the timer upon reception of the CIDX update and once the timer expires, another interrupt is generated. This process repeats until all the completion entries are read.

However, in the TRIGGER_EVERY, TRIGGER_USER and TRIGGER_USER_COUNT modes, an interrupt is sent, if at all, as a result of a Completion packet being received by the QDMA from the user logic. For every request by the user logic to send an interrupt, the QDMA sends one and only one interrupt. Thus in this case, if the software does not read all the Completion entries available to be read and the user logic does not send any more Completions requesting interrupts, the QDMA does not generate any more interrupts. This results in the residual Completions sitting in the Completion ring indefinitely. To avoid this from happening, when in TRIGGER_EVERY, TRIGGER_USER and TRIGGER_USER_COUNT mode, the software must read all the Completion entries in the Completion ring as indicated by an interrupt (or a Completion Status write).

The following are the flowcharts of different modes. These flowcharts are from the point of view of the Completion Engine. The Completion packets come in from the user logic and are written to the Completion Ring. The software (SW) update refers to the Completion Ring CIDX update sent from software to hardware.

Figure: Flowchart for EVERY Mode





Figure: Flowchart for USER Mode



X20641-040518

Figure: Flowchart for USER_COUNT Mode






Figure: Flowchart for USER_TIMER Mode



X20637-040519

Figure: Flowchart for USER_TIMER_COUNT Mode



Completion Timer

The Completion Timer engine supports the timer trigger mode in the Completion context. It supports 2048 queues, and each queue has its own timer. When the timer expires, a timer expire signal is sent to the Completion module. If multiple timers expire at the same time, they are sent out in a round robin manner.

Reference Timer

The reference timer is based on the timer tick. The register QDMA_C2H_INT (0xB0C) defines the value of a timer tick. The 16 registers QDMA_C2H_TIMER_CNT (0xA00-0xA3c) has the timer counts based on the timer tick. The `timer_idx` in the Completion context is the index to the 16 QDMA_C2H_TIMER_CNT registers. Each queue can choose its own `timer_idx`.

Handling Exception Events

C2H Completion On Invalid Queue

When QDMA receives a Completion on a queue which has an invalid context as indicated by the Valid bit in the C2H CMPT Context, the Completion is silently dropped.

C2H Completion On A Full Ring

The maximum number of Completion entries in the Completion Ring is 2 less than the total number of entries in the Completion Ring. The C2H Completion Context has PIDX and CIDX in it. This allows the QDMA to calculate the number of Completions in the Completion Ring. When the QDMA receives a Completion on a queue that is full, QDMA takes the following actions:

- Invalidates the C2H Completion Context for that queue.
- Marks the C2H Completion Context with error.
- Drops the Completion.
- If enabled, sends a Status Descriptor marked with error.
- If enabled and not outstanding, sends an Interrupt.
- Sends a Marker Response with error.
- Logs the error in the C2H Error Status Register.

C2H Completion With Descriptor Error

When the QDMA C2H Engine encounters a Descriptor Error, the following actions are taken in the context of the C2H Completion Engine:

- Invalidates the C2H Completion Context for that queue.
- Marks the C2H Completion Context with error.
- Sends the Completion out to the Completion Ring. It is marked with an error.
- If enabled and not outstanding, sends a Status Descriptor marked with error.
- If enabled and not outstanding, sends an Interrupt. Note that the Completion Engine can only send an interrupt and/or status descriptor if not outstanding. One implication of this is that if the interrupt happens to be outstanding when the descriptor error is encountered, a queue interrupt will not be sent to the software. Despite that, the error is logged and an error interrupt is still sent, if not masked by the software
- Sends a Marker Response with error.

C2H Completion With Invalid CIDX

The C2H Completion Engine has logic to detect that the CIDX value in the CIDX update points to an empty location in the Completion Ring. When it detects such error, the C2H Completion Engine:

- Invalidates the Completion Context.
- Marks the Completion Context with error.
- Logs an error in the C2H error status register.

Port ID Mismatch

The CMPT context specifies the port over which CMPTs are expected for that CMPT queue. If the `port_id` in the incoming CMPT is not the same as the `port_id` in the CMPT context, the CMPT

Engine treats the incoming CMPT as a mis-directed CMPT and drops it. It also logs an error. Note that the CMPT queue is not invalidated when a `port_id` mismatch occurs.

AXI Bridge Operations

Bridge

The Bridge core is an interface between the AXI4 and the PCI Express integrated block. It contains the memory mapped AXI4 to AXI4-Stream Bridge, and the AXI4-Stream Enhanced Interface Block for PCIe. The memory mapped AXI4 to AXI4-Stream Bridge contains a register block and two functional half bridges, referred to as the Slave Bridge and Master Bridge.

- The slave bridge connects to the AXI4 Interconnect as a slave device to handle any issued AXI4 master read or write requests.
- The master bridge connects to the AXI4 Interconnect as a master to process the PCIe generated read or write TLPs.
- The register block contains registers used in the Bridge core for dynamically mapping the AXI4 memory mapped (MM) address range provided using the AXIBAR parameters to an address for PCIe range.

The core uses a set of interrupts to detect and flag error conditions.

Related Information

[Bridge Register Space](#)

Slave Bridge

The slave bridge provides termination of memory-mapped AXI4 transactions from an AXI4 master device (such as a processor). The slave bridge provides a way to translate addresses that are mapped within the AXI4 memory mapped address domain to the domain addresses for PCIe. Write transactions to the Slave Bridge are converted into one or more MemWr TLPs, depending on the configured Max Payload Size setting, which are passed to the integrated block for PCIe. When a remote AXI master initiates a read transaction to the slave bridge, the read address and qualifiers are captured and a MemRd request TLP is passed to the core and a completion timeout timer is started. Completions received through the core are correlated with pending read requests and read data is returned to the AXI4 master. The slave bridge can support up to 32 AXI4 write requests, and 32 AXI4 read requests.

 **Note:** If slave reads and writes are valid, IP prioritizes reads over writes. You are recommended to have proper arbitration (leave some gaps between reads so writes can pass through).

BDF Table

Address translations for AXI address is done based on BDF table programming (0x2420 to 0x2434). These BDF table entries can be programmed through the AXI Slave CSR interface, `s_axil_csr_*`. There are three regions that you can use for slave data transfers. Each region can be further divided into many windows for a different address translation. These regions and

number of windows should be configured in the IP wizard configuration. Each entry in the BDF table programming represents one window. If you need 2 windows then 2 entries need to be programmed and so on.

There are some restrictions in programming BDF table.

1. All PCIe slave bridge data transfers must be quiesced before programming the BDF table.
2. There are six registers for each BDF table entry. All six registers must be programmed to make a valid entry. Even if some registers have 0s, you need to program 0s in those registers.
3. All the six registers need to be programmed in an order for an entry to be valid. Order is listed below.
 - a. 0x2420
 - b. 0x2424
 - c. 0x2428
 - d. 0x242C
 - e. 0x2430
 - f. 0x2434

BDF table entry start address = $0x2420 + (0x20 * i)$, where i = table entry number.

Protection

Specifying protection levels for different windows within a BAR is facilitated by AXI4 prot field via Trustzones. Any access from PMC will have $a^{\star}prot[1]=0$ and hence will get full access. For the BDF space the protection domain ID itself is stored in the BDF table. When a request comes in with $a^{\star}rprot[1]=0$, it will be allowed full access. Requests with $a^{\star}prot[1]=1$ will only be allowed to access BDF entries that have lower protection level.

The following table describes this behavior:

Table: AXI BAR Protection Levels

Access Type	BDF Table Value (prot[2:0])	Value in $a^{\star}prot[2:0]$ (AXI Interface)	Action
Secure access	3'hXXX	3'hX0X (bit 1=0)	Allow
Non-secure access to secure entry	3'hX0X	3'hX1X (bit 1=1)	Do not allow
Non-secure access to less secure entry	3'hX1X	3'hX1X (bit 1=1)	Allow if bits [2] and [0] match between $a^{\star}prot$ and BDF entry

Address Translation

Address translation can be turned off by selecting the No Address Translation option during IP configuration. When this option is selected, one full 64-bit BAR space is given for slave data

transfer. You must set up any address translation if needed. If No Address Translation is not selected DMA will do address translation.

One 64 bit BAR space is divided into 8 (called window) and 8 window space is available for address translation. Address translations can be done for address bits more than the window size. Window size = BAR size/Number of windows. (number of windows are fixed to 8). Address translation for Slave Bridge transfer are described in the following examples:

Slave Address Translation Examples

Example 1: BAR Size of 64 KB, with 1 Window Size 4 KB

Window 0: 4 KB with address translation of 0x7 for bits [63:13].

1. Selections in Vivado IP configuration in the AXI BARs tab are as follows:

- AXI BAR size 64K: 0xFFFF bits [15:0]
- Set Aperture Base Address: 0x0000_0000_0000_0000
- Set Aperture High Address: 0x0000_0000_0000_FFFF

2. The BDF table programming:

- Program 1 entries for 1 window
- Window Size = AXI BAR size/8 = 64K / 8 = 0x1FFF = 8 KB (13 bits). Each window max size is 8 KB.
- In this example for window size of 4K, 0x1 is programmed at 0x2430 bits [25:0].
- Address translation for bits [63:13] are programmed at 0x2420 and 0x2424.
- In this example, address translation for bits [63:13] are set to 0x7.

Table: BDF Table Programming

Program Value	Registers
0x0000_E000	Address translation value Low
0x0	Address translation value High
0x0	PASID/ Reserved
0x0	[11:0]: Function Number
0xC0000001	[31:30] Read/Write Access permission [29] : R0 access Error [28:26] Protection ID [25:0] Window Size ([25:0]*4K = actual size of the window)
0x0	reserved

For this example Slave address 0x0000_0000_0000_0100 will be address translated to 0x0000_0000_0000_E100.

Example 2: BAR Size of 64 KB, with 1 Window 8 KB

Window 0:8 KB with address translation of 0x6 ('b110) for bits [63:13].

1. Selections in Vivado IP configuration in the AXI BARs tab are as follows:

- AXI BAR size 64K: 0xFFFF bits [15:0]
- Set Aperture Base Address: 0x0000_0000_0000_0000
- Set Aperture High Address: 0x0000_0000_0000_FFFF

2. The BDF table programming:

- Program 1 entries for 1 window.
- Window Size = AXI BAR size/8 = 64K / 8 = 0x1FFF = 8 KB (13 bits). Each window max size is 8 KB.
- In this example for window size of 8K, 0x2 is programmed at 0x2430 bits [25:0].
- Address translation for bits [63:13] are programmed at 0x2420 and 0x2424.
- In this example, address translation for bits [63:13] are set to 0x6 ('b110).

Table: BDF Table Programming

Offset	Program Value	Registers
0x2420	0x0000_C000	Address translation value Low
0x2424	0x0	Address translation value High
0x2428	0x0	PASID/ Reserved
0x242C	0x0	[11:0]: Function Number
0x2430	0xC0000002	[31:30] Read/Write Access permission [29] : R0 access Error [28:26] Protection ID [25:0] Window Size ([25:0]*4K = actual size of the window)
0x2434	0x0	reserved

For this example, the Slave address 0x0000_0000_0000_0100 will be address translated to 0x0000_0000_0000_C100.

Example 3: BAR Size of 32 GB, and 4 Windows of Various Sizes

- Window 0: 4 KB with address translation of 0x7 for bits [63:32].
- Window 1: 4 GB with address translation of 0x0 for bits [63:32].
- Window 2: 64 KB with address translation of 0xB BBBB for bits [63:32].
- Window 3: 1 GB with address translation of 0x11111 for bits [63:32].

1. Selections in AMD Vivado™ IP configuration in the AXI BARs tab are as follows:
 - AXI BAR size 32G: 0x7_FFFF_FFFF bits [34:0].
 - Set Aperture Base Address: 0x0000_0000_0000_0000.
 - Set Aperture High Address: 0x0000_0007_FFFF_FFFF.
2. The BDF table programming:
 - Window Size = AXI BAR size/8 = 32 GB / 8 = 0xFFFF_FFFF = 4 GB (32 bits). Each window max size is 4 GB.
 - Program 4 entries for 4 windows:
 - BDF entry 0 table starts at 0x2420.
 - BDF entry 1 table starts at 0x2440.
 - BDF entry 2 table starts at 0x2460.
 - BDF entry 3 table starts at 0x2480.
 - Window 0 size 4 KB.
 - Program 0x1 to 0x2430 bits [25:0].
 - Address translation for bits [34:32] are programmed at 0x2420 and 0x2424.
 - Program 0x0000_0000 to 0x2420.
 - Program 0x0000_0007 to 0x2424
 - Window 1 size 4 GB.
 - Program 0x10_0000 to 0x2450 bits [25:0].
 - Address translation for bits [63:32] are programmed at 0x2440 and 0x2444.
 - Program 0x0000_0000 to 0x2440.
 - Program 0x0000_0000 to 0x2444
 - Window 2 size 64 KB.
 - Program 0x10 to 0x2470 bits [25:0].
 - Address translation for bits [63:32] are programmed at 0x2460 and 0x2464.
 - Program 0x0000_0000 to 0x2460
 - Program 0x0000_BBBB to 0x2464
 - Window 3 size 1 GB.
 - Program 0x4_0000 to 0x2490 bits [25:0].
 - Address translation for bits [63:32] are programmed at 0x2480 and 0x2484.
 - Program 0x0000_0000 to 0x2480.
 - Program 0x0001_1111 to 0x2484

Table: BDF Table Programming Entry 0

Offset	Program Value	Registers
0x2420	0x0000_0000	Address translation value Low
0x2424	0x7	Address translation value High
0x2428	0x0	PASID/ Reserved
0x242C	0x0	[11:0]: Function Number

Offset	Program Value	Registers
0x2430	0xC0000001	[31:30] Read/Write Access permission [29] : R0 access Error [28:26] Protection ID [25:0] Window Size ([25:0]*4K = actual size of the window)
0x2434	0x0	reserved

Table: BDF Table Programming Entry 1

Offset	Program Value	Registers
0x2440	0x0000	Address translation value Low
0x2444	0x0	Address translation value High
0x2448	0x0	PASID/ Reserved
0x244C	0x0	[11:0]: Function Number
0x2450	0xC010_0000	[31:30] Read/Write Access permission [29] : R0 access Error [28:26] Protection ID [25:0] Window Size ([25:0]*4K = actual size of the window)
0x2454	0x0	reserved

Table: BDF Table Programming Entry 2

Offset	Program Value	Registers
0x2460	0x_0000	Address translation value Low
0x2464	0xB BBBB	Address translation value High
0x2468	0x0	PASID/ Reserved
0x246C	0x0	[11:0]: Function Number
0x2470	0xC000_00010	[31:30] Read/Write Access permission [29] : R0 access Error [28:26] Protection ID [25:0] Window Size ([25:0]*4K = actual size of the window)
0x2474	0x0	reserved

Table: BDF Table Programming Entry 3

Offset	Program Value	Registers
0x2480	0x0000_0000	Address translation value Low
0x2484	0x1_1111	Address translation value High
0x2488	0x0	PASID/ Reserved
0x248C	0x0	[11:0]: Function Number
0x2490	0xC004_0000	[31:30] Read/Write Access permission [29] : R0 access Error [28:26] Protection ID [25:0] Window Size ([25:0]*4K = actual size of the window)
0x2494	0x0	reserved

For the above example:

- The slave address 0x0000_0000_0000_0100 translated to 0x0000_0007_0000_0100.
- The slave address 0x0000_0001_0000_0100 translated to 0x0000_0000_0000_0100.
- The slave address 0x0000_0002_0000_0100 translated to 0x0000_BBBB_0000_0100.
- The slave address 0x0000_0003_0000_0100 translated to 0x0001_1111_0000_0100.

The slave bridge does not support narrow burst AXI transfers. To avoid narrow burst transfers, connect the AXI smart-connect module which will convert narrow burst to full burst AXI transfers.

Master Bridge

The master bridge processes both PCIe MemWr and MemRd request TLPs received from the integrated block for PCI Express and provides a means to translate addresses that are mapped within the address for PCIe domain to the memory mapped AXI4 address domain. Each PCIe MemWr request TLP header is used to create an address and qualifiers for the memory mapped AXI4 bus and the associated write data is passed to the addressed memory mapped AXI4 Slave. The Master Bridge can support up to 32 active PCIe MemWr request TLPs. PCIe MemWr request TLPs support is as follows:

- 4 for 64-bit AXI4 data width
- 8 for 128-bit AXI4 data width
- 16 for 256-bit AXI4 data width
- 32 for 512-bit AXI4 data width

Each PCIe MemRd request TLP header is used to create an address and qualifiers for the memory mapped AXI4 bus. Read data is collected from the addressed memory mapped AXI4 bridge slave and used to generate completion TLPs which are then passed to the integrated block for PCI

Express. The Master Bridge in AXI Bridge mode can support up to 32 active PCIe MemRd request TLPs with pending completions for improved AXI4 pipe-lining performance.

AXI Transactions for PCIe

The following tables are the translation tables for AXI4-Stream and memory-mapped transactions.

Table: AXI4 Memory-Mapped Transactions to AXI4-Stream PCIe TLPs

AXI4 Memory-Mapped Transaction	AXI4-Stream PCIe TLPs
INCR Burst Read of AXIBAR	MemRd 32 (3DW)
INCR Burst Write to AXIBAR	MemWr 32 (3DW)
INCR Burst Read of AXIBAR	MemRd 64 (4DW)
INCR Burst Write to AXIBAR	MemWr 64 (4DW)

Table: AXI4-Stream PCIe TLPs to AXI4 Memory Mapped Transactions

AXI4-Stream PCIe TLPs	AXI4 Memory-Mapped Transaction
MemRd 32 (3DW) of PCIEBAR	INCR Burst Read
MemWr 32 (3DW) to PCIEBAR	INCR Burst Write
MemRd 64 (4DW) of PCIEBAR	INCR Burst Read
MemWr 64 (4DW) to PCIEBAR	INCR Burst Write

For PCIe® requests with lengths greater than 1 Dword, the size of the data burst on the Master AXI interface is always equal the width of the AXI data bus even when the request received from the PCIe link is shorter than the AXI bus width.

Slave axi write strobe (wstrb) signal can be used to facilitate data alignment to an address boundary. Write strobe signal can be 0 in the beginning of a valid data cycle and it appropriately calculates an offset to the given address. However, the valid data identified by the write strobe signal must be continuous from the first byte enable to the last byte enable.

Transaction Ordering for PCIe

The QDMA Subsystem for PCIe subsystem conforms to PCIe® transaction ordering rules. See the [PCI-SIG Specifications](#) for the complete rule set. The following behaviors are implemented in the QDMA Subsystem for PCIe subsystem to enforce the PCIe transaction ordering rules on the highly-parallel AXI bus of the bridge.

- The bresp to the remote (requesting) AXI4 master device for a write to a remote PCIe device is not issued until the MemWr TLP transmission is guaranteed to be sent on the PCIe link before any subsequent TX-transfers.
- If Relaxed Ordering bit is not set within the TLP header, then a remote PCIe device read to a remote AXI slave is not permitted to pass any previous remote PCIe device writes to a remote AXI slave received by the QDMA Subsystem for PCIe subsystem. The AXI read address phase is held until the previous AXI write transactions have completed and bresp has been received for the AXI write transactions. If the Relaxed Ordering attribute bit is set within the TLP header, then the remote PCIe device read is permitted to pass.
- Read completion data received from a remote PCIe device are not permitted to pass any remote PCIe device writes to a remote AXI slave received by the QDMA Subsystem for PCIe subsystem prior to the read completion data. The bresp for the AXI write(s) must be received before the completion data is presented on the AXI read data channel.

 **Note:** The transaction ordering rules for PCIe might have an impact on data throughput in heavy bidirectional traffic.

Malformed TLP

The integrated block for PCI Express® detects a malformed TLP. For the IP configured as an Endpoint core, a malformed TLP results in a fatal error message being sent upstream if error reporting is enabled in the Device Control register.

Abnormal Conditions

This section describes how the Slave side and Master side (see the following tables) of the QDMA Subsystem for PCIe subsystem handle abnormal conditions.

Slave Bridge Abnormal Conditions

Slave bridge abnormal conditions are classified as: Illegal Burst Type and Completion TLP Errors. The following sections describe the manner in which the Bridge handles these errors.

Illegal Burst Type

The slave bridge monitors AXI read and write burst type inputs to ensure that only the INCR (incrementing burst) type is requested. Any other value on these inputs is treated as an error condition and the Slave Illegal Burst (SIB) interrupt is asserted. In the case of a read request, the Bridge asserts SLVERR for all data beats and arbitrary data is placed on the s_axi_rdata bus. In the case of a write request, the Bridge asserts SLVERR for the write response and all write data is discarded.

Completion TLP Errors

Any request to the bus for PCIe (except for a posted Memory write) requires a completion TLP to

complete the associated AXI request. The Slave side of the Bridge checks the received completion TLPs for errors and checks for completion TLPs that are never returned (Completion Timeout). Each of the completion TLP error types are discussed in the subsequent sections.

Unexpected Completion

When the slave bridge receives a completion TLP, it matches the header RequesterID and Tag to the outstanding RequesterID and Tag. A match failure indicates the TLP is an Unexpected Completion which results in the completion TLP being discarded and a Slave Unexpected Completion (SUC) interrupt strobe being asserted. Normal operation then continues.

Unsupported Request

A device for PCIe might not be capable of satisfying a specific read request. For example, if the read request targets an unsupported address for PCIe, the completer returns a completion TLP with a completion status of 0b001 - Unsupported Request. The completer that returns a completion TLP with a completion status of Reserved must be treated as an unsupported request status, according to the PCI Express Base Specification v3.0. When the slave bridge receives an unsupported request response, the Slave Unsupported Request (SUR) interrupt is asserted and the DECERR response is asserted with arbitrary data on the AXI4 memory mapped bus.

Completion Timeout

A Completion Timeout occurs when a completion (Cpl) or completion with data (CplD) TLP is not returned after an AXI to PCIe memory read request, or after a PCIe Configuration Read/Write request. For PCIe Configuration Read/Write request, completions must complete within the C_COMP_TIMEOUT parameter selected value from the time the request is issued. For PCIe Memory Read request, completions must complete within the value set in the Device Control 2 register in the PCIe Configuration Space register. When a completion timeout occurs, an OKAY response is asserted with all 1s data on the memory mapped AXI4 bus.

Poison Bit Received on Completion Packet

An Error Poison occurs when the completion TLP EP bit is set, indicating that there is poisoned data in the payload. When the slave bridge detects the poisoned packet, the Slave Error Poison (SEP) interrupt is asserted and the SLVERR response is asserted with arbitrary data on the memory mapped AXI4 bus.

Completer Abort

A Completer Abort occurs when the completion TLP completion status is 0b100 - Completer Abort. This indicates that the completer has encountered a state in which it was unable to complete the transaction. When the slave bridge receives a completer abort response, the Slave Completer Abort (SCA) interrupt is asserted and the SLVERR response is asserted with arbitrary data on the memory mapped AXI4 bus.

Table: Slave Bridge Response to Abnormal Conditions

Transfer Type	Abnormal Condition	Bridge Response
Read	Illegal burst type	SIB interrupt is asserted. SLVERR response given with arbitrary read data.

Transfer Type	Abnormal Condition	Bridge Response
Write	Illegal burst type	SIB interrupt is asserted. Write data is discarded. SLVERR response given.
Read	Unexpected completion	SUC interrupt is asserted. Completion is discarded.
Read	Unsupported Request status returned	SUR interrupt is asserted. DECERR response given with arbitrary read data.
Read	Completion timeout	SCT interrupt is asserted. SLVERR response given with arbitrary read data.
Read	Poison bit in completion	Completion data is discarded. SEP interrupt is asserted. SLVERR response given with arbitrary read data.
Read	Completer Abort (CA) status returned	SCA interrupt is asserted. SLVERR response given with arbitrary read data.

Master Bridge Abnormal Conditions

The following sections describe the manner in which the master bridge handles abnormal conditions.

AXI DECERR Response

When the master bridge receives a DECERR response from the AXI bus, the request is discarded and the Master DECERR (MDE) interrupt is asserted. If the request was non-posted, a completion packet with the Completion Status = Unsupported Request (UR) is returned on the bus for PCIe.

AXI SLVERR Response

When the master bridge receives a SLVERR response from the addressed AXI slave, the request is discarded and the Master SLVERR (MSE) interrupt is asserted. If the request was non-posted, a completion packet with the Completion Status = Completer Abort (CA) is returned on the bus for PCIe.

Max Payload Size for PCIe, Max Read Request Size or 4K Page Violated

When the master bridge receives a SLVERR response from the addressed AXI slave, the request is discarded and the Master SLVERR (MSE) interrupt is asserted. If the request was non-posted, a completion packet with the Completion Status = Completer Abort (CA) is returned on the bus for PCIe.

Completion Packets

When the MAX_READ_REQUEST_SIZE is greater than the MAX_PAYLOAD_SIZE, a read request for PCIe can ask for more data than the master bridge can insert into a single completion packet. When this situation occurs, multiple completion packets are generated up to MAX_PAYLOAD_SIZE, with the Read Completion Boundary (RCB) observed.

Poison Bit

When the poison bit is set in a transaction layer packet (TLP) header, the payload following the header is corrupted. When the master bridge receives a memory request TLP with the poison bit set, it discards the TLP and asserts the Master Error Poison (MEP) interrupt strobe.

Zero Length Requests

When the master bridge receives a read request with the Length = 0x1, FirstBE = 0x00, and LastBE = 0x00, it responds by sending a completion with Status = Successful Completion.

When the master bridge receives a write request with the Length = 0x1, FirstBE = 0x00, and LastBE = 0x00 there is no effect.

Table: Master Bridge Response to Abnormal Conditions

Transfer Type	Abnormal Condition	Bridge Response
Read	DECERR response	MDE interrupt strobe asserted. Completion returned with Unsupported Request status.
Write	DECERR response	MDE interrupt strobe asserted.
Read	SLVERR response	MSE interrupt strobe asserted. Completion returned with Completer Abort status.
Write	SLVERR response	MSE interrupt strobe asserted.
Write	Poison bit set in request	MEP interrupt strobe asserted. Data is discarded.
Read	DECERR response	MDE interrupt strobe asserted. Completion returned with

Transfer Type	Abnormal Condition	Bridge Response
		Unsupported Request status.
Write	DECERR response	MDE interrupt strobe asserted.

Link Down Behavior

Interrupts

The QDMA Subsystem for PCIe supports up to 2K total MSI-X vectors. A single MSI-X vector can be used to support multiple queues. Each function can support up to 8 vectors (8 *256 function = 2K vectors).

The QDMA supports Interrupt Aggregation. Each vector has an associated Interrupt Aggregation Ring. The QID and status of queues requiring service are written into the Interrupt Aggregation Ring. When a PCIe® MSI-X interrupt is received by the Host, the software reads the Interrupt Aggregation Ring to determine which queue needs service. Mapping of queues to vectors is programmable vector number provided in the queue context. It supports MSI-X interrupt modes for SR-IOV and non-SR-IOV.

Asynchronous and Queue Based Interrupts

The QDMA supports both asynchronous interrupts and queue-based interrupts.

The asynchronous interrupts are used for capturing events that are not synchronous to any DMA operations, namely, errors, status, and debug conditions.

Interrupts are broadcast to all PFs, and maintain status for each PF in a queue based scheme.


The queue based interrupts include the interrupts from the H2C MM, H2C stream, C2H MM, and C2H stream.

Interrupt Engine

The Interrupt Engine handles the queue based interrupts and the error interrupt.

The following figure shows the Interrupt Engine block diagram.

Figure: Interrupt Engine Block Diagram



X20891-111020

The Interrupt Engine gets the interrupts from H2C MM, H2C stream, C2H MM, C2H stream, or error interrupt.

It handles the interrupts in two ways: direct interrupt or indirect interrupt. The interrupt sources has the information that shows if it is direct interrupt or indirect interrupt. It also has the information of the vector. If it is direct interrupt, the vector is the interrupt vector that is used to generate the PCIe MSI-X message (the interrupt vector index of the MSIX table). If it is indirect interrupt, the vector is the ring index of the Interrupt Aggregation Ring. The interrupt source gets the information of interrupt type and vector from the Descriptor Software Context, the Completion Context, or the error interrupt register.

Direct Interrupt

For direct interrupt, the Interrupt Engine gets the interrupt vector from the source, and it then sends out the PCIe MSI-X message directly.

Interrupt Aggregation Ring


For the indirect interrupt, it does interrupt aggregation. The following are some restrictions for the interrupt aggregation.

- Each Interrupt Aggregation Ring can only be associated with one function. But multiple rings can be associated with the same function.
- The interrupt engine supports up to three interrupts from the same source, until the software services the interrupts.
- Interrupt aggregation ring size needs to be $> 3 * \text{number of Qs}$.

The Interrupt Engine processes the indirect interrupt with the following steps.

- Interrupt source provides the index to which interrupt ring it belongs too.
- Reads interrupt context for that queue.
- Writes to the Interrupt Aggregation Ring.
- Sends out the PCIe MSI-X message.

This following figure shows the indirect interrupt block diagram.

Figure: Indirect Interrupt

The Interrupt Context includes the information of the Interrupt Aggregation Ring. It has 256 entries to support up to 256 Interrupt Aggregation Rings.

Color bit is added so software does not read more entries than what it should read. When the software allocates the memory space for the Interrupt Aggregation Ring, the `coal_color` starts with `1'b0`. The software needs to initialize the color bit of the Interrupt Context to be `1'b1`. When the hardware completes the entire ring and flips to first entry in the next pass, it also flips the color value to 0 and starts writing 0 in color bit space. The software does the same after it completes the last entry with a color value 1, and goes to the first entry in the second pass and expects a color value 0. If the software does not see a color value 0, which indicates an old entry, it waits for new entry with a color value 0.

The software reads the Interrupt Aggregation Ring to get the Qid, and the `int_type` (H2C or C2H). From the Qid, the software can identify whether the queue is stream or MM.

The `stat_desc` in the Interrupt Aggregation Ring is the status descriptor from the Interrupt source. When the status descriptor is disabled, the software can get the status descriptor information from the Interrupt Aggregation Ring.

There can be two cases:

- The interrupt source is C2H stream. Then it is the status descriptor of the C2H Completion Ring. The software can read the `pidx` of the C2H Completion Ring.
- The interrupt source is others (H2C stream, H2C MM, C2H MM). Then it is the status descriptor of that source. The software can read the `cidx`.

Finally, the Interrupt Engine sends out the PCIe MSI-X message using the interrupt vector from the Interrupt Context. When there is an interrupt from any source, the interrupt engine updates PIDX and check for `int_st` of that interrupt context. If `int_st` is 0 (WAITING_TRIGGER) then the interrupt engine will send a interrupt. If `int_st` is 1 (ISR_RUNNING), the interrupt engine will not send interrupt. If the interrupt engine sends interrupt it will update `int_sts` to 1 and once software updated CIDX and CIDX matches PIDX `int_sts` will be cleared. The process is explained below.

When the PCIe MSI-X interrupt is received by the Host, the software reads the Interrupt Aggregation Ring to determine which queue needs service. After the software reads the ring, it will do a dynamic pointer update for the software CIDX to indicate the cumulative pointer that the

software reads to. The software does the dynamic pointer update using the register QDMA_DMAP_SEL_INT_CIDX[2048] (0x18000). If the software CIDX is equal to the PIDX, this will trigger a write to the Interrupt Context to clear int_ston the interrupt state of that queue. This is to indicate the QDMA that the software already reads all of the entries in the Interrupt Aggregation Ring. If the software CIDX is not equal to the PIDX, the interrupt engine will send out another PCIe MSI-X message. Therefore, the software can read the Interrupt Aggregation Ring again. After that, the software can do a pointer update of the interrupt source ring. For example, if it is C2H stream interrupt, the software will update pointer of the interrupt source ring, which is the C2H Completion Ring.

These are the steps for the software:

1. After the software gets the PCIe MSI-X message, it reads the Interrupt Aggregation Ring entries.
2. The software uses the coal_color bit to identify the written entries. Each entry has Qid and Int_type (H2C or C2H). From the Qid and Int_type, the software can check if it is stream or MM. This points to a corresponding source ring. For example, if it is C2H stream, the source ring is the C2H Completion Ring. The software can then read the source ring to get information, and do a dynamic pointer update of the source ring after that.
3. After the software finishes reading of all written entries in the Interrupt Aggregation Ring, it does one dynamic pointer update of the software cidx using the register QDMA_DMAP_SEL_INT_CIDX[2048] (0x18000). This communicates to the hardware of the Interrupt Aggregation Ring pointer used by the software.
If the software cidx is not equal to the pidx, the hardware will send out another PCIe MSI-X message, so that the software can read the Interrupt Aggregation Ring again.

When the software does the dynamic pointer update for the Interrupt Aggregation Ring using the register QDMA_DMAP_SEL_INT_CIDX[2048] (0x18000), it sends the ring index of the Interrupt Aggregation Ring.

The following diagram shows the indirect interrupt flow. The Interrupt module gets the interrupt requests. It first writes to the Interrupt Aggregation Ring. Then it waits for the write completions. After that, it sends out the PCIe MSI-X message. The interrupt requests can keep on coming, and the Interrupt module keeps on processing them. In the meantime, the software reads the Interrupt Aggregation Ring, and it does the dynamic pointer update. If the software CIDX is not equal to the PIDX, it will send out another PCIe MSI-X message.

Interrupt Context Structure

The following is the Interrupt Context Structure (0x8).

Table: Interrupt Context Structure (0x8)

Signal	Bit	Owner	Description
rsvd	[255:126]	Driver	Reserved. Initialize to 0s
func	[125:114]	Driver	Function number

Signal	Bit	Owner	Description
rsvd	[113:83]	Driver	Reserved. Initialize to 0s
at	[82]	Driver	1'b0: un-translated address 1'b1: translated address
pidx	[81:70]	DMA	Producer Index, updated by DMA IP.
page_size	[69:67]	Driver	Interrupt Aggregation Ring size: 0: 4 KB 1: 8 KB 2: 12 KB 3: 16 KB 4: 20 KB 5: 24 KB 6: 28 KB 7: 32 KB
baddr_4k	[66:15]	Drive	Base address of Interrupt Aggregation Ring – bit [63:12]
color	[14]	DMA	Color bit
int_st	[13]	DMA	Interrupt State: 0: WAIT_TRIGGER 1: ISR_RUNNING
Rsvd	[12]	NA	Reserved
vec	[11:1]	Driver	Interrupt vector index in msix table
valid	[0]	Driver	Valid

The software needs to size the Interrupt Aggregation Ring appropriately. Each source can send up to three messages to the ring. Therefore, the size of the ring needs satisfy the following formula.

Number of entry \geq 3 x number of queues

The Interrupt Context is programmed by the context access. The QDMA_IND_CTXT_CMD.Qid has the ring index, which is from the interrupt source. The operation of MDMA_CTXT_CMD_CLR can clear all of the bits in the Interrupt Context. The MDMA_CTXT_CMD_INV can clear the valid bit.

- Context access through QDMA_TRQ_SEL_IND:
 - QDMA_IND_CTXT_CMD.Qid = Ring index
 - QDMA_IND_CTXT_CMD.Sel = MDMA_CTXT_SEL_INT_COAL (0x8)
 - QDMA_IND_CTXT_CMD.cmd.Op =
 - MDMA_CTXT_CMD_WR
 - MDMA_CTXT_CMD_RD
 - MDMA_CTXT_CMD_CLR
 - MDMA_CTXT_CMD_INV

After the interrupt engine looks up the Interrupt Context, the interrupt engine writes to the Interrupt Aggregation Ring. The interrupt engine also updates the Interrupt Context with the new PIDX, color, and the interrupt state.

Interrupt Aggregation Entry

This is the Interrupt Aggregation Ring entry structure. It has 8B data.

Table: Interrupt Aggregation Ring Entry Structure

Signal	Bit	Owner	Description
Coal_color	[63]	DMA	The color bit of the Interrupt Aggregation Ring. This bit inverts every time pidx wraps around on the Interrupt Aggregation Ring.
Qid	[62:39]	DMA	This is from Interrupt source. Queue ID.
Int_type	[38:38]	DMA	0: H2C 1: C2H
Rsvd	[37:37]	DMA	Reserved
Stat_desc	[36:0]	DMA	This is the status descriptor of the Interrupt source.


The following is the information in the stat_desc.

Table: stat_desc Information

Signal	Bit	Owner	Description
Error	[36:35]	DMA	This is from interrupt source: c2h_err[1:0], or h2c_err[1:0].
Int_st	[34:33]	DMA	This is from Interrupt source. Interrupt state.

Signal	Bit	Owner	Description
			0: WRB_INT_ISR 1: WRB_INT_TRIGGER 2: WRB_INT_ARMED
Color	[32:32]	DMA	This is from Interrupt source. This bit inverts every time pidx wraps around and this field gets copied to color field of descriptor.
Cidx	[31:16]	DMA	This is from Interrupt source. Cumulative consumed pointer.
Pidx	[15:0]	DMA	This is from Interrupt source. Cumulative pointer of total interrupt Aggregation Ring entry written.

Interrupt Flow

Figure: Interrupt Flow

X20890-052418

Error Interrupt

There are Leaf Error Aggregators in different places. They log the errors and propagate the errors to the Central Error Aggregator. Each Leaf Error Aggregator has an error status register and an error mask register. The error mask is enable mask. Irrespective of the enable mask value, the error status register always logs the errors. Only when the error mask is enabled, the Leaf Error Aggregator will propagate the error to the Central Error Aggregator.

The Central Error Aggregator aggregates all of the errors together. When any error occurs, it can generate an Error Interrupt if the `err_int_arm` bit is set in the error interrupt register `QDMA_GLBL_ERR_INT` (0B04). The `err_int_arm` bit is set by the software and cleared by the hardware when the Error Interrupt is taken by the Interrupt Engine. The Error Interrupt is for all of the errors including the H2C errors and C2H errors. The Software must set this `err_int_arm` bit to generate interrupt again.

The Error Interrupt supports the direct interrupt only. Register `QDMA_GLBL_ERR_INT` bit[23], `en_coal` must always be programmed to 0 (direct interrupt).


The Error Interrupt gets the vector from the error interrupt register `QDMA_GLBL_ERR_INT`. For the direct interrupt, the vector is the interrupt vector index of the MSI-X table.

Here are the processes of the Error Interrupt.

1. Reads the Error Interrupt register `QDMA_C2H_GLBL_INT` (0B04) to get function and vector numbers.
2. Sends out the PCIe MSI-X message.

The following figure shows the error interrupt register block diagram.

Figure: Error Interrupt Handling



Legacy Interrupt

The QDMA Subsystem for PCIe supports the legacy interrupt for physical function, and it is expected that the single queue will be associated with interrupt.

To enable the legacy interrupt, the software needs to set the `en_lgcy_intr` bit in the register `QDMA_GLBL_GLBL_INTERRUPT_CFG` (0x2C4). When `en_lgcy_intr` is set, the QDMA will not send out MSI-X interrupt.

When the legacy interrupt wire INTA is asserted, the QDMA hardware sets the `lgcy_intr_pending` bit in the `QDMA_GLBL_GLBL_INTERRUPT_CFG` (0x2C4) register. When the software receives the legacy interrupt, it needs to clear the `lgcy_intr_pending` bit. The hardware will keep the legacy interrupt wire asserted until the software clears the


`lgcy_intr_pending` bit.

 **Note:** INTB, INTC, and INTD legacy interrupts are not supported.

User Interrupt

You can generate interrupt to host system using the user interrupt interface. You need to provide `usr_irq_in_fnc`, `usr_irq_in_vec`, and `usr_irq_in_vld` interrupts and they should be held active until `usr_irq_out_ack` is returned. `usr_irq_in_fnc` is a function number associated with an interrupt. If it is for MSI-X interrupt `usr_irq_in_vec` should be provided. If the interrupt is for legacy interrupt, vector is not needed.

Figure: Interrupt



Queue Management

Function Map Table

The Function Map Table is used to allocate queues to each function. The index into the RAM is the function number. Each entry contains the base number of the physical QID and the number of queues allocated to the function. It provides a function based, queue access protection mechanism by translating and checking accesses to logical queues (through QDMA_TRQ_SEL_QUEUE_PF and QDMA_TRQ_SEL_QUEUE_VF address space) to their physical queues. Direct register accesses to queue space beyond what is allocated to the function in the table is canceled and an error is logged.

Function map can be accessed through the indirect context register space QDMA_IND_CTXT_CMD registers, with QDMA_IND_CTXT_CMD.sel = 0xC. When accessed through indirect context register space, the context structure is defined by the Function Map Context Structure table. Along with FMAP table programming in the IP, you must program the FMAP table in the Mailbox IP. This is needed for function level reset (FLR) procedure.

1. Program Function Map Context structure in QDMA_IND_CTXT_DATA (0x804 - 0x820) registers as listed in the following table.
2. Program QDMA_IND_CTXT_CMD registers
 - a. [19:7] : function number
 - b. [6:5] : 2'h1 (write a context data)
 - c. [4:1] : 4'hC (FMAP)

For more information on QDMA_IND_CTXT_CMD (0x844), refer `qdma_v5_1_pf_registers.csv` available in the [Register Reference File](#).

Because these spaces exist only in the PF address map, only a physical function can modify this table.

Table: Function Map Context Structure (0xC)

Bits	Bit Width	Field Name	Description
[255:44]			Reserved. Set to 0.
[43:32]	12	Qid_max	Maximum number of queues this function has.
[31:11]			Reserved. Set to 0.
[10:0]	11	Qid_base	The base queue ID for the function.

Context Programming

- Program all mask registers to 1. They are QDMA_IND_CTXT_MASK_0 (0x824) to QDMA_IND_CTXT_MASK_7 (0x840) .
- Program context values on to the following registers: QDMA_IND_CTXT_DATA_0 (0x804) to QDMA_IND_CTXT_DATA_7 (0x820)
- A host profile table context needs to be programmed before any context settings QDMA_CTXT_SEL_C_HOST_PROFILE. Select 0xA in QDMA_IND_CTXT_CMD (0x844), and write all data field to 0s and program context. All other values are reserved.
- Refer to software descriptor context structure, C2H prefetch context structure and C2H CPMT context structure to program the context data registers.
- Program any context to corresponding queue in the following context command register: QDMA_IND_CTXT_CMD (0x844).
 - Qid is given in bits [19:7].
 - Opcode bits [6:5] selects what operations must be done.
 - 0 = QDMA_CTXT_CLR : All content of context is zeroed out. Qin is sent out on tm_dsc_sts
 - 1 = QDMA_CTXT_WR : Write context
 - 2 = QDMA_CTXT_RD : Read context
 - 3 = QDMA_CTXT_INV : Qin is set to zero and other context values are intact. Qin is sent out on tm_dsc_sts and unused credits are sent out.
 - The context that is accessed is given in bits [4:1].
 - 4'h0 = QDMA_CTXT_SEL_DEC_SW_C2H; C2H Descriptor SW Context
 - 4'h1 = QDMA_CTXT_SEL_DEC_SW_H2C; H2C descriptor SW context
 - 4'h2 = QDMA_CTXT_SEL_DEC_HW_C2H; C2H Descriptor HW Context
 - 4'h3 = QDMA_CTXT_SEL_DEC_HW_H2C; H2C Descriptor HW Context
 - 4'h4 = QDMA_CTXT_SEL_DEC_CR_C2H; C2H Descriptor HW Context
 - 4'h5 = QDMA_CTXT_SEL_DEC_CR_H2C; H2C Descriptor HW Context
 - 4'h6 = QDMA_CTXT_SEL_WRB; CMPT / used ring Context
 - 4'h7 = QDMA_CTXT_SEL_PFTCH; C2H PFCH Context
 - 4'h8 = QDMA_CTXT_SEL_INT_COAL; Interrupt Aggregation Context
 - 4'h9 = Reserved
 - 4'hA = QDMA_CTXT_SEL_C_HOST_PROFILE; Host Profile Table (Only QDMA_CTXT_CMD_WR and QDMA_CTXT_CMD_RD supported)
 - 4'hB = QDMA_CTXT_SEL_TIMER; Timer Context (Only QDMA_CTXT_CMD_INV supported)
 - 4'hC = QDMA_CTXT_SEL_FMAP FMAP table write (Only QDMA_CTXT_CMD_WR and QDMA_CTXT_CMD_RD supported)
 - 4'hD = QDMA_CTXT_SEL_FNC_STS (Per function BME enable/Disable)
 - Context programming write/read does not occur when bit [0] is set. For more information on register 0x844, refer qdma_v5_1_pf_registers.csv available in the [Register Reference File](#).

Related Information

[QDMA_CSR \(0x0000\)](#)

Queue Setup

- Clear Descriptor Software Context.
- Clear Descriptor Hardware Context.
- Clear Descriptor Credit Context.
- Set-up Descriptor Software Context.
- Clear Prefetch Context.
- Clear Completion Context.
- Set-up Completion Context.
 - If interrupts/status writes are desired (enabled in the Completion Context), an initial Completion CIDX update is required to send the hardware into a state where it is sensitive to trigger conditions. This initial CIDX update is required, because when out of reset, the hardware initializes into an unarmed state.
- Set-up Prefetch Context.

Queue Teardown

Queue Tear-down (C2H Stream):

- Send marker packet to drain the pipeline.
- Wait for marker completion.
- Invalidate/clear prefetch context.
- Invalidate/clear completion context.
- Invalidate/clear descriptor software context.
- Invalidate timer context (clear cmd is not supported).

Queue Tear-down (H2C Stream & MM):

- Invalidate/clear descriptor software context.

Virtualization

QDMA implements SR-IOV passthrough virtualization where the adapter exposes a separate virtual function (VF) for use by a virtual machine (VM). A physical function (PF) can be optionally made privileged with full access to QDMA registers and resources, but only VFs implement per queue pointer update registers and interrupts. VF drivers must communicate with the driver attached to the PF through the mailbox for configuration, resource allocation, and exception handling. The QDMA implements function level reset (FLR) to enable operating system on VM to reset the device without interfering with the rest of the platform.

Table: Privileged Access

Type	Notes
Queue context/other control registers	Registers for Context access only controlled by PFs (All 4 PFs).

Type	Notes
Status and statistics registers	Mainly PF only registers. VFs need to coordinate with a PF driver for error handling. VFs need to communicate through the mailbox with driver attached to PF.
Data path registers	Both PFs and VFs must be able to write the registers involved in data path without needing to go through a hypervisor. Pointer update for H2C/C2H Descriptor Fetch can be done directly by VF or PF for the queues associated with the function using its own BAR space. Any pointer updates to queue that do not belong to the function will be dropped with error logged.
Other protection recommendations	Turn on IOMMU to protect bad memory accesses from VMs.
PF driver and VF driver communication	The VF driver needs to communicate with the PF driver to request operations that have global effect. This communication channel needs this ability to pass messages and generate interrupts. This communication channel utilizes a set of hardware mailboxes for each VF.

Mailbox

In a virtualized environment, the driver attached to a PF has enough privilege to program and access QDMA registers. For all the lesser privileged functions, certain PFs and all VFs must communicate with privileged drivers using the mailbox mechanism. The communication API must be defined by the driver. The QDMA IP does not define it.


Each function (both PF and VF) has an inbox and an outbox that can fit a message size of 128B. A VF accesses its own mailbox, and a PF accesses its own mailbox and all the functions (PF or VF) associated with that PF.

☞ **Note:** pcie_qdma_mailbox IP supports up to 4PFs and 240VFs. You can build a mailbox system in the PL to support more number of PFs and VFs (CPM5 limit is 16PFs and 2KVF). Adding pcie_qdma_mailbox IP increases PL utilization.

The QDMA mailbox allows the following access:

- From a VF to the associated PF.
- From a PF to any VF belonging to its own virtual function group (VFG).
- From a PF (typically a driver that does not have access to QDMA registers) to another PF.

Figure: Mailbox




VF To PF Messaging

A VF is allowed to post one message to a target PF mailbox until the target function (PF) accepts it. Before posting the message the source function should make sure its `o_msg_status` is cleared, then the VF can write the message to its Outgoing Message Registers. After finishing message writing, the VF driver sends `msg_send` command through write 0x1 at the control/status register (CSR) address 0x5004. The mailbox hardware then informs the PF driver by asserting `i_msg_status` field.

The function driver should enable the periodic polling of the `i_msg_status` to check the availability of incoming messages. At a PF side, `i_msg_status = 0x1` indicates one or more message is pending for the PF driver to pick up. The `cur_src_fn` in the Mailbox Status Register gives the function ID of the first pending message. The PF driver should then set the Mailbox Target Function Register to the source function ID of the first pending message. Then access to a PF's Incoming Message Registers is indirectly, which means the mailbox hardware will always return the corresponding message bytes sent by the Target function. Upon finishing the message reading, the PF driver should also send `msg_rcv` command through write 0x2 at the CSR address. The hardware will deassert the `o_msg_status` at the source function side. The following figure illustrates the messaging flow from a VF to a PF at both the source and destination sides.

Figure: VF to PF Messaging Flow



VF (#n) to PF Message Flow
Status polling can be changed to interrupt driven

X21105-062118

PF To VF Messaging

The messaging flow from a PF to the VFs that belong to its VFG is slightly different than the VF to PF flow because:

A PF can send messages to multiple destination functions, therefore, it might receive multiple acknowledgments at the moment when checking the status. As illustrated in the following figure, a PF driver must set Mailbox Target Function Register to the destination function ID before doing any message operation; for example, checking the incoming message status, write message, or send the command. At the VF side (receiving side), whenever a VF driver get the `i_msg_status = 0x1`, the VF driver should read its Incoming Message Registers to pick up the message.

Depending on the application, the VF driver can send the `msg_rcv` immediately after reading the message or after the corresponding message being processed.


To avoid one-by-one polling of the status of outgoing messages, the mailbox hardware provides a set of Acknowledge Status Registers (ASR) for each PF. Upon the mailbox receiving the `msg_rcv` command from a VF, it deasserts the `o_msg_status` field of the source PF and it also sets the corresponding bit in the Acknowledge Status Registers. For a given VF with function ID `<N>`, acknowledge status is at:

- Acknowledge Status Register address: `<N> / 32 + 0x22420 Register Address`
- Acknowledge Status bit location: `<N> / 32`

The mailbox hardware asserts the `ack_status` filed in the Status Register (0x22400) when there

is any bit was asserted in the Acknowledge Status Register (ASR). The PF driver can poll the `ack_status` before actually reading out the Acknowledge status registers. The PF driver might detect multiple completions through one register access. After being processed, the PF driver should also write the value back to the same register address to clear the status.

Figure: PF to VF Messaging Flow



Mailbox Interrupts

The mailbox module supports interrupt as the alternative event notification mechanism. Each mailbox has an Interrupt Control Register (at the offset 0x22410 for a PF, or at the offset 0x5010 for a VF). Set 1 to this register to enable the interrupt. Once the interrupt is enabled, the mailbox will send the interrupt to the QDMA given there is any pending event for the mailbox to process, namely, any incoming message pending or any acknowledgment for the outgoing messages. Configure the interrupt vector through the Function Interrupt Vector Register (0x22408 for a PF, or 0x5008 for a VF) according to the driver configuration.

Enabling the interrupt does not change the event logging mechanism, which means the user must check the pending events through reading the Function Status Registers. The first step to respond to an interrupt request is disabling the interrupt. It is possible that the actual number of the pending events is more than the number of the events at the moment when the mailbox is sent the interrupt.

Recommended: AMD recommends that the user application interrupt handler process all the pending events that present in the status register. Upon finishing the interrupt response, the user application re-enables the interrupt.

The mailbox will check its event status at the time the interrupt control change from disabled to enabled. If there is any new events that arrived the mailbox between reading the interrupt status and the re-enabling the interrupt, the mailbox will generate a new interrupt request immediately.

Related Information

[QDMA_PF_MAILBOX \(0x22400\)](#)

[QDMA_VF_MAILBOX \(0x5000\)](#)

Function Level Reset

The function level reset (FLR) mechanism enables the software to quiesce and reset Endpoint hardware with function-level granularity. When a VF is reset, only the resources associated with this VF are reset. When a PF is reset, all resources of the PF, including that of its associated VFs, are reset. Because FLR is a privileged operation, it must be performed by the PF driver running in the management system.

Use Mode

- Hypervisor requests for FLR when a function is attached and detached (that is; power on and off).
- You can request FLR as follows:

```
echo 1 > /sys/bus/pci/devices/$BDF/reset
```

where \$BDF is the bus device function number of the targeted function.

FLR Process

A complete FLR process involves of three major steps.

1. Pre-FLR: Pre-FLR resets all QDMA context structure, mailbox, and user logic of the target function.
 - Each function has a register called FLR Control Status register, which keeps track of the pre-FLR status of the function. The offset is calculated as FLR Control Status register offset = MB_base + 0x100, which is located at offset 0x100 from the mailbox memory space of the function. Note that PF and VF have different MB_base. The definition of FLR Control Status register is shown in the following table.
 - The software writes 1 to Bit[0] flr_status of the target function to initiate pre-FLR. The hardware clears Bit[0] flr_status when pre-FLR completes. The software keeps polling on Bit[0] flr_status, and only proceeds to the next step when it returns 0.

Table: FLR Control Status Register

Offset	Field	R/W Type	Width	Default	Description
0x100	pre_flr_st	RW	32	0	[31:1] reserved. [0]: 1 Initiates pre-FLR. [0]: 0 pre-FLR done.

Offset	Field	R/W Type	Width	Default	Description
					bit[0] is set by the driver and cleared by the hardware.

2. Quiesce: The software must ensure all pending transaction is completed. This can be done by polling the Transaction Pending bit in the Device Status register (in PCIe Configuration Space), until it is cleared or times out after a certain period of time.
3. PCIe-FLR: PCIe-FLR resets all resources of the target function in the PCIe controller.

Note: Initiate Function Level Reset bit (bit 15 of PCIe Device Control Register) of the target function should be set to 1 to trigger FLR process in PCIe.

OS Support

If the PF driver is loaded and alive (i.e., use mode 1), all three steps aforementioned are performed by the driver. However, for an AMD UltraScale+ device, if you want to perform FLR before loading the PF driver (as defined in Use Mode above), an OS kernel patch is provided to allow OS to perform the correct FLR sequence through functions defined in //.../source/drivers/pci/quick.c.

Port ID

Port ID is the categorization of some queues on the FPGA side. When the DMA is shared by more than one user application, the port ID provides indirection to QID so that all the interfaces can be further demuxed with lower cost. However, when used by a single application, the port ID can be ignored and drive the port id inputs to 0s.

Host Profile

Host profile must be programmed to represent root port host. Host profile can be programmed through context programming. Select QDMA_CTXT_SEL_C_HOST_PROFILE (4'hA) in QDMA_IND_CTXT_CMD. Host profile context structure is given in the following table:

Table: Host Profile Context Structure

Bit	Bit Width	Field Name	Description
[255:188]	68	Reserved	Reserved
[187:186]	2		H2C AXI4-MM write awprot
[185:182]	4		H2C AXI4-MM write awcache
[181:178]	4		H2C AXIMM steering

Bit	Bit Width	Field Name	Description
[177:104]	74	Reserved	Reserved
[103:102]	2		C2H AXI4-MM read arprot
[101:98]	4		C2H AXI4-MM read awcache
[97:94]	4		C2H AXIMM steering
[0:93]	94	Reserved	Reserved

 **Note:** H2C AXI4-MM steering bit and C2H AXI4-MM steering bit should set to 0s. If not, DMA AXI4-MM transfers do not work. For most cases, host profile context structure is all 0s, and host profile must still be programmed to represent a host.

System Management

Resets

The QDMA Subsystem for PCIe supports all the PCIe defined resets, such as link down, reset, hot reset, and function level reset (FLR) (supports only Quiesce mode).

Soft Reset

Reset the QDMA logic through the `soft_reset_n` port. This port needs to be held in reset for a minimum of 100 clock cycles (`axi_aclk` cycles).

This does not reset PCIe hard block. It resets only the DMA portion of logic. This reset can be asserted if there is a DMA hang or some error condition.

Soft Reset Use Cases

The uses cases that prompt the use of `soft_reset` include:

- DMA hangs and user is not getting proper values.
- DMA transfer have errors, but the PCIe links are good.
- DMA records some asynchronous error

After `soft_reset`, you must reinitialize the queues and program all queue context.

VDM

Vendor Defined Messages (VDMs) are an expansion of the existing messaging capabilities with PCIe. PCI Express Specification defines additional requirements for Vendor Defined

Messages, header formats and routing information. For details, see *PCI-SIG Specifications* (<https://www.pcisig.com/specifications>).

QDMA allows the transmission and reception of VDMs. To enable this feature, select Enable Bridge Slave Mode in the Vivado Customize IP dialog box. This enables the `st_rx_msg` interface. RX Vendor Defined Messages are stored in shallow FIFO before they are transmitted to the output port. When there are many back-to-back VDM messages, FIFO will overflow and these message will be dropped. So it is better to repeat VDM messages at regular intervals.

Throughput for VDMs depend on several factors: PCIe speed, data width, message length, and the internal VDM pipeline.

Internal VDM pipelines must be replaced with the Internal RX VDM FIFO interface for network on chip (NoC) access, which has a shallow buffer of 64B.

 **Note:** New VDM messages will be dropped if more than 64B of VDM are received before the FIFO is serviced through NoC.

Internal RX VDM FIFO interface cannot handle back-to-back messages. Pipeline throughput can only handle one in every four accesses, which is about 25% efficiency from the host access.

!! Important: Do not use back-to-back VDM access.

RX Vendor Defined Messages:

1. When QDMA receives a VDM, the incoming messages will be received on the `st_rx_msg` port.
2. The incoming data stream will be captured on the `st_rx_msg_data` port (per-DW).
3. The user application needs to drive the `st_rx_msg_rdy` to signal if it can accept the incoming VDMs.
4. Once `st_rx_msg_rdy` is High, the incoming VDM is forwarded to the user application.
5. The user application needs to store this incoming VDMs and track of how many packets were received.

TX Vendor Defined Messages:

1. To enable transmission of VDM from QDMA, program the TX Message registers in the Bridge through the Slave interface.
2. Bridge has TX Message Control, Header L (bytes 8-11), Header H (bytes 12-15) and TX Message Data registers as shown in the PCIe TX Message Data FIFO Register (`TX_MSG_DFIFO`).
3. Issue a Write to offset 0xE64 through Slave interface for the TX Message Header L register.
4. Program offset 0xE68 for the required VDM TX Header H register.
5. Program up to 16DW of Payload for the VDM message starting from DW0 – DW15 by sending Writes to offset 0xE6C one by one.
6. Program the `msg_routing`, `msg_code`, data length, requester function field and `msg_execute` field in the `TX_MSG_CTRL` register in offset 0xE60 to send the VDM TX packet.
7. The TX Message Control register also indicates the completion status of the message in bit 23. User needs to read this bit to confirm the successful transmission of the VDM packet.
8. All the fields in the registers are RW except bit 23 (`msg_fail`) in TX Control register which

- is cleared by writing a 1.
9. VDM TX packet will be sent on the AXI-ST RQ transmit interface.

Related Information

[VDM Ports](#)

[Bridge Register Space](#)

Config Extend

PCIe extended interface can be selected for more configuration space. When the Configuration Extend Interface is selected, you are responsible for adding logic to extend the interface to make it work properly.

Expansion ROM

If selected, the Expansion ROM is activated and can be a value from 2 KB to 4 GB. According to the PCI Local Bus Specification (*PCI-SIG Specifications* (<https://www.pcisig.com/specifications>))), the maximum size for the Expansion ROM BAR should be no larger than 16 MB. Selecting an address space larger than 16 MB can result in a non-compliant core.

Errors

Bridge Errors

Slave Bridge Abnormal Conditions

Slave bridge abnormal conditions are classified as: Illegal Burst Type and Completion TLP Errors. The following sections describe the manner in which the Bridge handles these errors.

Illegal Burst Type

The slave bridge monitors AXI read and write burst type inputs to ensure that only the INCR (incrementing burst) type is requested. Any other value on these inputs is treated as an error condition and the Slave Illegal Burst (SIB) interrupt is asserted. In the case of a read request, the Bridge asserts SLVERR for all data beats and arbitrary data is placed on the s_axi_rdata bus. In the case of a write request, the Bridge asserts SLVERR for the write response and all write data is discarded.

Completion TLP Errors

Any request to the bus for PCIe (except for a posted Memory write) requires a completion TLP to complete the associated AXI request. The Slave side of the Bridge checks the received completion TLPs for errors and checks for completion TLPs that are never returned (Completion Timeout). Each of the completion TLP error types are discussed in the subsequent sections.

Unexpected Completion

When the slave bridge receives a completion TLP, it matches the header RequesterID and Tag to the outstanding RequesterID and Tag. A match failure indicates the TLP is an Unexpected Completion which results in the completion TLP being discarded and a Slave Unexpected Completion (SUC) interrupt strobe being asserted. Normal operation then continues.

Unsupported Request

A device for PCIe might not be capable of satisfying a specific read request. For example, if the read request targets an unsupported address for PCIe, the completer returns a completion TLP with a completion status of 0b001 - Unsupported Request. The completer that returns a completion TLP with a completion status of Reserved must be treated as an unsupported request status, according to the PCI Express Base Specification v3.0. When the slave bridge receives an unsupported request response, the Slave Unsupported Request (SUR) interrupt is asserted and the DECERR response is asserted with arbitrary data on the AXI4 memory mapped bus.

Completion Timeout

A Completion Timeout occurs when a completion (Cpl) or completion with data (CplD) TLP is not returned after an AXI to PCIe memory read request, or after a PCIe Configuration Read/Write request. For PCIe Configuration Read/Write request, completions must complete within the C_COMP_TIMEOUT parameter selected value from the time the request is issued. For PCIe Memory Read request, completions must complete within the value set in the Device Control 2 register in the PCIe Configuration Space register. When a completion timeout occurs, an OKAY response is asserted with all 1s data on the memory mapped AXI4 bus.

Poison Bit Received on Completion Packet

An Error Poison occurs when the completion TLP EP bit is set, indicating that there is poisoned data in the payload. When the slave bridge detects the poisoned packet, the Slave Error Poison (SEP) interrupt is asserted and the SLVERR response is asserted with arbitrary data on the memory mapped AXI4 bus.

Completer Abort

A Completer Abort occurs when the completion TLP completion status is 0b100 - Completer Abort. This indicates that the completer has encountered a state in which it was unable to complete the transaction. When the slave bridge receives a completer abort response, the Slave Completer Abort (SCA) interrupt is asserted and the SLVERR response is asserted with arbitrary data on the memory mapped AXI4 bus.

Table: Slave Bridge Response to Abnormal Conditions

Transfer Type	Abnormal Condition	Bridge Response
---------------	--------------------	-----------------

Transfer Type	Abnormal Condition	Bridge Response
Read	Illegal burst type	SIB interrupt is asserted. SLVERR response given with arbitrary read data.
Write	Illegal burst type	SIB interrupt is asserted. Write data is discarded. SLVERR response given.
Read	Unexpected completion	SUC interrupt is asserted. Completion is discarded.
Read	Unsupported Request status returned	SUR interrupt is asserted. DECERR response given with arbitrary read data.
Read	Completion timeout	SCT interrupt is asserted. SLVERR response given with arbitrary read data.
Read	Poison bit in completion	Completion data is discarded. SEP interrupt is asserted. SLVERR response given with arbitrary read data.
Read	Completer Abort (CA) status returned	SCA interrupt is asserted. SLVERR response given with arbitrary read data.

Master Bridge Abnormal Conditions

The following sections describe the manner in which the master bridge handles abnormal conditions.

AXI DECERR Response

When the master bridge receives a DECERR response from the AXI bus, the request is discarded and the Master DECERR (MDE) interrupt is asserted. If the request was non-posted, a completion packet with the Completion Status = Unsupported Request (UR) is returned on the bus for PCIe.

AXI SLVERR Response

When the master bridge receives a SLVERR response from the addressed AXI slave, the request is discarded and the Master SLVERR (MSE) interrupt is asserted. If the request was non-posted, a completion packet with the Completion Status = Completer Abort (CA) is returned on the bus for

PCIe.

Max Payload Size for PCIe, Max Read Request Size

When the master bridge receives a SLVERR response from the addressed AXI slave, the request is discarded and the Master SLVERR (MSE) interrupt is asserted. If the request was non-posted, a completion packet with the Completion Status = Completer Abort (CA) is returned on the bus for PCIe.

Completion Packets

When the MAX_READ_REQUEST_SIZE is greater than the MAX_PAYLOAD_SIZE, a read request for PCIe can ask for more data than the master bridge can insert into a single completion packet. When this situation occurs, multiple completion packets are generated up to MAX_PAYLOAD_SIZE, with the Read Completion Boundary (RCB) observed.

Poison Bit

When the poison bit is set in a transaction layer packet (TLP) header, the payload following the header is corrupted. When the master bridge receives a memory request TLP with the poison bit set, it discards the TLP and asserts the Master Error Poison (MEP) interrupt strobe.

Zero Length Requests

When the master bridge receives a read request with the Length = 0x1, FirstBE = 0x00, and LastBE = 0x00, it responds by sending a completion with Status = Successful Completion.

When the master bridge receives a write request with the Length = 0x1, FirstBE = 0x00, and LastBE = 0x00 there is no effect.

Table: Master Bridge Response to Abnormal Conditions

Transfer Type	Abnormal Condition	Bridge Response
Read	DECERR response	MDE interrupt strobe asserted. Completion returned with Unsupported Request status.
Write	DECERR response	MDE interrupt strobe asserted.
Read	SLVERR response	MSE interrupt strobe asserted. Completion returned with Completer Abort status.
Write	SLVERR response	MSE interrupt strobe asserted.

Transfer Type	Abnormal Condition	Bridge Response
Write	Poison bit set in request	MEP interrupt strobe asserted. Data is discarded.
Read	DECERR response	MDE interrupt strobe asserted. Completion returned with Unsupported Request status.
Write	DECERR response	MDE interrupt strobe asserted.

Linkdown Errors

If the PCIe link goes down during DMA operations, transactions might be lost and the DMA might not be able to complete. In such cases, the AXI4 interfaces continue to operate. Outstanding read requests on the C2H Bridge AXI4 MM interface receive correct completions or completions with a slave error response. The DMA logs a link down error in the status register. It is the responsibility of the driver to have a timeout and handle recovery of a link down situation.

Data Path Errors

Data protection is supported on the primary data paths. CRC error can occur on C2H streaming, H2C streaming. Parity error can occur on Memory Mapped, Bridge Master and Bridge Slave interfaces. Error on Write payload can occur on C2H streaming, Memory Mapped and Bridge Slave. Double bit error on write payload and read completions for Bridge Slave interface causes parity error. Parity errors on requests to the PCIe are dropped by the core, and a fatal error is logged by the PCIe. Parity errors are not recoverable and can result in unexpected behavior. Any DMA during and after the parity error should be considered invalid. If there is a parity error and transfer hangs or stops, the DMA will log the error. You must investigate and fix the parity issues. Once the issues are fixed, clear that queue and reopen the queue to start a new transfer.

DMA Errors

All DMA errors are logged in their respective error status register. Each block has error status and error mask register so error can be passed on to higher level and eventually to QDMA_GLBL_ERR_STAT register.

Errors can be fatal error based on register settings. If there is an fatal error DMA will stop the transfer and will send interrupt if enabled. After debug and analysis, you must invalidate and restart the queue to start the DMA transfer.

Error Aggregator


There are Leaf Error Aggregators in different places. They log the errors and propagate them to the central place. The Central Error Aggregator aggregates the errors from all of the Leaf Error Aggregators.

The QDMA_GLBL_ERR_STAT register is the error status register of the Central Error Aggregator. The bit fields indicate the locations of Leaf Error Aggregators. Then, look for the error status register of the individual Leaf Error Aggregator to find the exact error.

The register QDMA_GLBL_ERR_MASK is the error mask register of the Central Error Aggregator. It has the mask bits for the corresponding errors. When the mask bit is set to 1'b1, it will enable the corresponding error to be propagated to the next level to generate an Interrupt. The detail information of the error generated interrupt is described in the interrupt section. Error interrupt is controlled by the register QDMA_GLBL_ERR_INT (0xB04).

Each Leaf Error Aggregator has an error status register and an error mask register. The error status register logs the error. The hardware sets the bit when the error happens, and the software can write 1'b1 to clear the bit if needed. The error mask register has the mask bits for the corresponding errors. When the mask bit is set to 1'b1, it will enable the propagation of the corresponding error to the Central Error Aggregator. The error mask register does not affect the error logging to the error status register.

Figure: Error Aggregator



The error status registers and the error mask registers of the Leaf Error Aggregators are as follows.

C2H Streaming Error

- QDMA_C2H_ERR_STAT (0xAF0): This is the error status register of the C2H streaming errors.
- QDMA_C2H_ERR_MASK (0xAF4): This is the error mask register. The software can set the bit to enable the corresponding C2H streaming error to be propagated to the Central Error Aggregator.
- QDMA_C2H_FIRST_ERR_QID (0xB30): This is the Qid of the first C2H streaming error.

C2H MM Error

- QDMA_C2H MM Status (0x1040)
- C2H MM Error Code Enable Mask (0x1054)
- C2H MM Error Code (0x1058)
- C2H MM Error Info (0x105C)

QDMA H2C0 MM Error

- H2C0 MM Status (0x1240)
- H2C MM Error Code Enable Mask (0x1254)
- H2C MM Error Code (0x1258)
- H2C MM Error Info (0x125C)

TRQ Error

- QDMA_GLBL_TRQ_ERR_STS (0x264): This is the error status register of the Trq errors.
- QDMA_GLBL_TRQ_ERR_MSK (0x268): This is the error mask register.
- QDMA_GLBL_TRQ_ERR_LOG_A (0x26C): This is the error logging register. It shows the select, function and the address of the access when the error happens.

Descriptor Error

- QDMA_GLBL_DSC_ERR_STS (0x254)
- QDMA_GLBL_DSC_ERR_MSK (0x258): This is the error logging register. It has the QID, DMA direction, and the consumer index of the error.
- QDMA_GLBL_DSC_ERR_LOG0 (0x25C)
- QDMA_GLBL_TRQ_ERR_STS (0x264): This is the error status register of the TRQ errors.

RAM Double Bit Error

- QDMA_RAM_DB_E_STS_A (0xFC)
- QDMA_RAM_DB_E_MSK_A (0xF8)

RAM Single Error

- QDMA_RAM_SBE_STS_A (0xF4)
- QDMA_RAM_SBE_MSK_A (0xF0)

Related Information

[Register Space](#)

[C2H Streaming Fatal Error Handling](#)

- QDMA_C2H_FATAL_ERR_STAT (0xAF8): The error status register of the C2H streaming fatal errors.
- QDMA_C2H_FATAL_ERR_MASK (0xAFC): The error mask register. The SW can set the bit to enable the corresponding C2H fatal error to be sent to the C2H fatal error handling logic.
- QDMA_C2H_FATAL_ERR_ENABLE (0xB00): This register enables two C2H streaming fatal error handling processes:

bit[0]

Stop the data transfer by disabling the write request from the C2H DMA write engine by setting enable_wrq_dis bit [0] to 1.

bit[1]

Invert the write payload parity on the data transfer by setting enable_wpl_par_inv bit [1] to 1.

Related Information

[QDMA_CSR \(0x0000\)](#)

Port Descriptions

The QDMA Subsystem for PCIe connects directly to the PCIe Integrated Block. The datapath interfaces to the PCIe Integrated Block IP are 128, 256 or 512-bits wide, and runs at up to 250 MHz depending on the configuration of the IP. The datapath width applies to all data interfaces. Ports associated with this core are described below.

Table: Parameters

Parameter Name	Description
PL_LINK_CAP_MAX_LINK_WIDTH	Phy lane width
C_M_AXI_ADDR_WIDTH	AXI4 Master interface Address width
C_M_AXI_ID_WIDTH	AXI4 Master interface id width
C_M_AXI_DATA_WIDTH	AXI4 Master interface data width 128 or 256 or 512 bits
C_S_AXI_ID_WIDTH	AXI4 Bridge Slave interface id width
C_S_AXI_ADDR_WIDTH	AXI4 Bridge Slave interface Address width
C_S_AXI_DATA_WIDTH	AXI4 Bridge Slave interface data width 128 or 256 or 512 bits
C_S_AXI_ID_WIDTH	AXI4 Bridge Slave interface id width
AXI_DATA_WIDTH	AXI4 DMA transfer data width. Example 128 or 256 or 512 bits

Related Information

[QDMA Architecture](#)

QDMA Global Ports

Table: QDMA Global Port Descriptions

Port Name	I/O	Description
sys_clk	I	Should be driven by the ODIV2 port of reference clock IBUFDS_GTE4. See the <i>UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide</i> (PG213).
sys_clk_gt	I	PCIe reference clock. Should be driven from the port of reference clock IBUFDS_GTE4. See the <i>UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide</i> (PG213).
sys_rst_n	I	Reset from the PCIe edge connector reset signal.
pci_exp_txp [PL_LINK_CAP_MAX_LINK_WIDTH-1:0]	O	PCIe TX serial interface.
pci_exp_txn [PL_LINK_CAP_MAX_LINK_WIDTH-1:0]	O	PCIe TX serial interface.
pci_exp_rxp [PL_LINK_CAP_MAX_LINK_WIDTH-1:0]	I	PCIe RX serial interface.
pci_exp_rxn [PL_LINK_CAP_MAX_LINK_WIDTH-1:0]	I	PCIe RX serial interface.
user_lnku	O	Output active-High identifies that the PCI Express core is linked up with a host device.
axi_aclk	O	User clock out. PCIe derived clock output for all interface signals output from and input to QDMA. Use this clock to drive inputs and gate outputs from QDMA.
axi_aresetn	O	User reset out. AXI reset signal synchronous with the clock provided on the axi_aclk output. This reset should drive all corresponding AXI

Port Name	I/O	Description
		Interconnect aresetn signals.
soft_reset_n	I	Soft reset (active-Low). Use this port to assert reset and reset the DMA logic. This will reset only the DMA logic. User should assert and de-assert this port.
phy_ready	O	Phy ready out status.
csr_prog_done	O	This port is enabled only when the AXI-Lite CSR Slave Interface option is selected in the Basic tab in the IP customization GUI. This port indicates whether access to AXI Lite CSR interface is available. 1'b0: The AXI Lite CSR Slave interface is not accessible. 1'b1: The AXI Lite CSR Slave interface is accessible.

All AXI interfaces are clocked out and in by the `axi_aclk` signal. You are responsible for using `axi_aclk` to drive all signals into the DMA.

AXI Bridge Master Ports

Table: AXI4 Memory Mapped Master Bridge Read Address Interface Port Descriptions

Signal Name	I/O	Description
<code>m_axib_araddr</code> [C_M_AXI_ADDR_WIDTH-1:0]	O	This signal is the address for a memory mapped read to the user logic from the host.
<code>m_axib_arid</code> [C_M_AXI_ID_WIDTH-1:0]	O	Master read address ID.
<code>m_axib_arlen[7:0]</code>	O	Master read address length.
<code>m_axib_arsize[2:0]</code>	O	Master read address size.
<code>m_axib_arprot[2:0]</code>	O	Master read protection type.
<code>m_axib_arvalid</code>	O	The assertion of this signal means there is a valid read request to the address on <code>m_axib_araddr</code> .
<code>m_axib_arready</code>	I	Master read address ready.
<code>m_axib_arlock</code>	O	Master read lock type.
<code>m_axib_arcache[3:0]</code>	O	Master read memory type.

Signal Name	I/O	Description
m_axib_arburst[1:0]	O	Master read address burst type.
m_axib_aruser[28:0]	O	Master read user bits. m_axib_aruser[10:0] = reserved m_axib_aruser[11] = bridge traffic m_axib_aruser[15:12] = bar id m_axib_aruser[18:16] = reserved m_axib_aruser[30:19] = function number m_axib_aruser[31] = reserved m_axib_aruser[39:32] = bus number m_axib_aruser[42:40] = vf group m_axib_aruser[54:43] = vfg offset

Table: AXI4 Memory Mapped Master Bridge Read Interface Port Descriptions

Signal Name	I/O	Description
m_axib_rdata [C_M_AXI_DATA_WIDTH-1:0]	I	Master read data.
m_axib_ruser [C_M_AXI_DATA_WIDTH/8-1:0]	I	m_axib_ruser[C_M_DATA_WIDTH/8-1:0] = read data odd parity, per byte.
m_axib_rid [C_M_AXI_ID_WIDTH-1:0]	I	Master read ID.
m_axib_rresp[1:0]	I	Master read response.
m_axib_rlast	I	Master read last.
m_axib_rvalid	I	Master read valid.
m_axib_rready	O	Master read ready.

Table: AXI4 Memory Mapped Master Bridge Write Address Interface Port Descriptions

Signal Name	I/O	Description
m_axib_awaddr [C_M_AXI_ADDR_WIDTH-1:0]	O	This signal is the address for a memory mapped write to the user logic from the host.
m_axib_awid [C_M_AXI_ID_WIDTH-1:0]	O	Master write address ID.
m_axib_awlen[7:0]	O	Master write address length.

Signal Name	I/O	Description
m_axib_awsize[2:0]	O	Master write address size.
m_axib_awburst[1:0]	O	Master write address burst type.
m_axib_awprot[2:0]	O	Master write protection type.
m_axib_awvalid	O	The assertion of this signal means there is a valid write request to the address on m_axib_araddr.
m_axib_awready	I	Master write address ready.
m_axib_awlock	O	Master write lock type.
m_axib_awcache[3:0]	O	Master write memory type.
m_axib_awuser[28:0]	O	Master write user bits. m_axib_aruser[10:0] = reserved m_axib_aruser[11] = bridge traffic m_axib_aruser[15:12] = bar id m_axib_aruser[18:16] = reserved m_axib_aruser[30:19] = function number m_axib_aruser[31] = reserved m_axib_aruser[39:32] = bus number m_axib_aruser[42:40] = vf group m_axib_aruser[54:43] = vfg offset

Table: AXI4 Memory Mapped Master Bridge Write Interface Port Descriptions

Signal Name	I/O	Description
m_axib_wdata [C_M_AXI_DATA_WIDTH-1:0]	O	Master write data.
m_axib_wuser [C_M_AXI_DATA_WIDTH/8-1:0]	O	m_axib_wuser [C_M_AXI_DATA_WIDTH/8-1:0] = write data odd parity, per byte.
m_axib_wlast	O	Master write last.
m_axib_wstrb [C_M_AXI_DATA_WIDTH/8-1:0]	O	Master write strobe.
m_axib_wvalid	O	Master write valid.
m_axib_wready	I	Master write ready.

Table: AXI4 Memory Mapped Master Bridge Write Response Interface Port Descriptions

Signal Name	I/O	Description
m_axib_bvalid	I	Master write response valid.
m_axib_bresp[1:0]	I	Master write response.
m_axib_bid [C_M_AXI_ID_WIDTH-1:0]	I	Master write response ID.
m_axib_bready	O	Master response ready.

AXI Bridge Slave Ports

Table: AXI4 Bridge Slave Write Address Interface Port Descriptions

Port Name	I/O	Description
s_axib_awid [C_S_AXI_ID_WIDTH-1:0]	I	Slave write address ID.
s_axib_awaddr [C_S_AXI_ADDR_WIDTH-1:0]	I	Slave write address.
s_axib_awuser[7:0]	I	s_axib_awuser[7:0] indicates function_number.
s_axib_awregion[3:0]	I	Slave write region decode.
s_axib_awlen[7:0]	I	Slave write burst length.
s_axib_awsize[2:0]	I	Slave write burst size.
s_axib_awburst[1:0]	I	Slave write burst type. Only the INCR burst type is supported.
s_axib_awvalid	I	Slave address write valid.
s_axib_awready	O	Slave address write ready.

Table: AXI4 Bridge Slave Write Interface Port Descriptions

Port Name	I/O	Description
s_axib_wdata [C_S_AXI_DATA_WIDTH-1:0]	I	Slave write data.
s_axib_wstrb [C_S_AXI_DATA_WIDTH/8-1:0]	I	Slave write strobe.

Port Name	I/O	Description
s_axib_wlast	I	Slave write last.
s_axib_wvalid	I	Slave write valid.
s_axib_wready	O	Slave write ready.
s_axib_wuser [C_S_AXI_DATA_WIDTH/8-1:0]	I	s_axib_wuser [C_S_AXI_DATA_WIDTH/8-1:0] = write data odd parity, per byte.

The s_axib_wstrb signal can be used to facilitate data alignment to an address boundary. s_axib_wstrb can be equal to 0 in the beginning of a valid data cycle and will appropriately calculate an offset to the given address. However, the valid data identified by s_axib_wstrb must be continuous from the first byte enable to the last byte enable.

Table: AXI4 Bridge Slave Write Response Interface Port Descriptions

Port Name	I/O	Description
s_axib_bid [C_S_AXI_ID_WIDTH-1:0]	O	Slave response ID.
s_axib_bresp[1:0]	O	Slave write response.
s_axib_bvalid	O	Slave write response valid.
s_axib_bready	I	Slave response ready.

Table: AXI4 Bridge Slave Read Address Interface Port Descriptions

Port Name	I/O	Description
s_axib_arid [C_S_AXI_ID_WIDTH-1:0]	I	Slave read address ID.
s_axib_araddr [C_S_AXI_ADDR_WIDTH-1:0]	I	Slave read address.
s_axib_arregion[3:0]	I	Slave read region decode.
s_axib_arlen[7:0]	I	Slave read burst length.
s_axib_arsize[2:0]	I	Slave read burst size.
s_axib_arburst[1:0]	I	Slave read burst type. Only the INCR burst type is supported.
s_axib_arvalid	I	Slave read address valid.

Port Name	I/O	Description
s_axib_arready	O	Slave read address ready.

Table: AXI4 Bridge Slave Read Interface Port Descriptions

Port Name	I/O	Description
s_axib_rid [C_S_AXI_ID_WIDTH-1:0]	O	Slave read ID tag.
s_axib_rdata [C_S_AXI_ID_WIDTH-1:0]	O	Slave read data.
s_axib_ruser [C_S_AXI_DATA_WIDTH/8-1:0]	O	s_axib_aruser[C_S_AXI_ID_WIDTH/8-1:0] = read data odd parity, per byte.
s_axib_rresp[1:0]	O	Slave read response.
s_axib_rlast	O	Slave read last.
s_axib_rvalid	O	Slave read valid.
s_axib_rready	I	Slave read ready.

AXI4-Lite Master Ports

Table: Config AXI4-Lite Memory Mapped Write Master Interface Port Descriptions

Signal Name	I/O	Description
m_axil_awaddr[31:0]	O	This signal is the address for a memory mapped write to the user logic from the host.
m_axil_awprot[2:0]	O	Protection type.
m_axil_awvalid	O	The assertion of this signal means there is a valid write request to the address on m_axil_awaddr.
m_axil_awready	I	Master write address ready.
m_axil_awuser [54:0]	O	m_axil_awuser[11:0] = reserved m_axil_awuser[15:12] = bar id m_axil_awuser[18:16] = reserved m_axil_awuser[30:19] = function number m_axil_awuser[31] = reserved m_axil_awuser[39:32] = bus number m_axil_awuser[42:40] = vf group

Signal Name	I/O	Description
		m_axil_awuser[54:43] = vfg offset
m_axil_wdata[31:0]	O	Master write data.
m_axil_wstrb[3:0]	O	Master write strobe.
m_axil_wvalid	O	Master write valid.
m_axil_wready	I	Master write ready.
m_axil_bvalid	I	Master response valid.
m_axil_bresp[1:0]	I	
m_axil_bready	O	Master response valid.

Table: Config AXI4-Lite Memory Mapped Read Master Interface Port Descriptions

Signal Name	I/O	Description
m_axil_araddr[31:0]	O	This signal is the address for a memory mapped read to the user logic from the host.
m_axil_aruser[54:0]	O	m_axil_aruser[11:0] = reserved m_axil_aruser[15:12] = bar id m_axil_aruser[18:16] = reserved m_axil_aruser[30:19] = function number m_axil_aruser[31] = reserved m_axil_aruser[39:32] = bus number m_axil_aruser[42:40] = vf group m_axil_aruser[54:43] = vfg offset
m_axil_arprot[2:0]	O	Protection type.
m_axil_arvalid	O	The assertion of this signal means there is a valid read request to the address on m_axil_araddr.
m_axil_arready	I	Master read address ready.
m_axil_rdata[31:0]	I	Master read data.
m_axil_rresp[1:0]	I	Master read response.
m_axil_rvalid	I	Master read valid.
m_axil_rready	O	Master read ready.

AXI4-Lite Slave Ports

AXI4-Lite Slave ports can be used to access QDMA Queue space registers ([QDMA_TRQ_SEL_QUEUE_PF \(0x18000\)](#) and [QDMA_TRQ_SEL_QUEUE_VF \(0x3000\)](#)).

Table: Config AXI4-Lite Memory Mapped Write Slave Interface Signals

Signal Name	I/O	Description
s_axil_awaddr[31:0]	I	This signal is the address for a memory mapped write to the DMA Queue Space register or Bridge register from the user logic. s_axil_awaddr[31:28]: 4'b0001 - QDMA Queue Space register 4'b0000 - ECAM access (0x0000 : 0x0E00)
s_axil_awvalid	I	The assertion of this signal means there is a valid write request to the address on s_axil_awaddr.
s_axil_awuser[12:0]	I	[12:8]: Reserved [7:0]: Function number
s_axil_awprot[2:0]	I	Protection type. This port is not being used.
s_axil_awready	O	Slave write address ready.
s_axil_wdata[31:0]	I	Slave write data.
s_axil_wstrb[3:0]	I	Slave write strobe.
s_axil_wvalid	I	Slave write valid.
s_axil_wready	O	Slave write ready.
s_axil_bvalid	O	Slave write response valid.
s_axil_bresp[1:0]	O	Slave write response.
s_axil_bready	I	Save response ready.

Table: Config AXI4-Lite Memory Mapped Read Slave Interface Signals

Signal Name	I/O	Description
s_axil_araddr[31:0]	I	This signal is the address for a memory mapped read to the DMA Queue Space register or Bridge register from the user logic. s_axil_araddr[31:28]: 4'b0001 - QDMA Queue Space register 4'b0000 - ECAM access (0x0000 : 0x0E00)
s_axil_arprot[2:0]	I	Protection type. This port is not being used.

Signal Name	I/O	Description
s_axil_arvalid	I	The assertion of this signal means there is a valid read request to the address on s_axil_araddr.
s_axil_aruser[12:0]	I	[12:8]: Reserved [7:0]: Function number
s_axil_arready	O	Slave read address ready.
s_axil_rdata[31:0]	O	Slave read data.
s_axil_rresp[1:0]	O	Slave read response.
s_axil_rvalid	O	Slave read valid.
s_axil_rready	I	Slave read ready.

AXI4-Lite Slave CSR Ports

Table: Config AXI4-Lite Memory Mapped Write Slave CSR Interface Signals

Signal Name	I/O	Description									
s_axil_csr_awaddr[31:0]	I	<p>This signal is the address for a memory mapped write to the DMA from the user logic.</p> <table border="1"> <thead> <tr> <th>Mode</th><th>Bit[28]</th><th>Bit[15]</th></tr> </thead> <tbody> <tr> <td>QDMA CSR register</td><td>0</td><td>1</td></tr> <tr> <td>Bridge register</td><td>0</td><td>0</td></tr> </tbody> </table> <p>Other combinations of Bit-28 and Bit-25 are not valid.</p>	Mode	Bit[28]	Bit[15]	QDMA CSR register	0	1	Bridge register	0	0
Mode	Bit[28]	Bit[15]									
QDMA CSR register	0	1									
Bridge register	0	0									
s_axil_csr_awvalid	I	The assertion of this signal means there is a valid write request to the address on s_axil_csr_awaddr.									
s_axil_csr_awprot[2:0]	I	Protection type. This port is not being used.									
s_axil_csr_awready	O	Slave write address ready.									
s_axil_csr_wdata[31:0]	I	Slave write data.									
s_axil_csr_wstrb[3:0]	I	Slave write strobe.									
s_axil_csr_wvalid	I	Slave write valid.									

Signal Name	I/O	Description
s_axil_csr_wready	O	Slave write ready.
s_axil_csr_bvalid	O	Slave write response valid.
s_axil_csr_bresp[1:0]	O	Slave write response.
s_axil_csr_bready	I	Save response ready.

Table: Config AXI4-Lite Memory Mapped Read Slave CSR Interface Signals

Signal Name	I/O	Description		
s_axil_csr_araddr[31:0]	I	This signal is the address for a memory mapped read to the DMA from the user logic.		
		Mode	Bit[28]	Bit[15]
		QDMA CSR register	0	1
		Bridge register	0	0
		Other combinations of Bit-28 and Bit-25 are not valid.		
s_axil_csr_arprot[2:0]	I	Protection type. This port is not being used.		
s_axil_csr_arvalid	I	The assertion of this signal means there is a valid read request to the address on s_axil_csr_araddr.		
s_axil_csr_arready	O	Slave read address ready.		
s_axil_csr_rdata[31:0]	O	Slave read data.		
s_axil_csr_rresp[1:0]	O	Slave read response.		
s_axil_csr_rvalid	O	Slave read valid.		
s_axil_csr_rready	I	Slave read ready.		

AXI4 Memory Mapped DMA Ports

Table: AXI4 Memory Mapped DMA Read Address Interface Signals

Signal Name	I/O	Description
m_axi_araddr	O	This signal is the address for a memory mapped

Signal Name	I/O	Description
[C_M_AXI_ADDR_WIDTH-1:0]		read to the user logic from the DMA.
m_axi_arid [3:0]	O	Standard AXI4 description, which is found in the AXI4 Protocol Specification AMBA AXI4-Stream Protocol Specification (ARM IHI 0051A).
m_axi_aruser[31:0]	O	m_axi_aruser[18:0] = reserved m_axi_aruser[31:19] = queue number
m_axi_arlen[7:0]	O	Master read burst length.
m_axi_arsize[2:0]	O	Master read burst size.
m_axi_arprot[2:0]	O	Protection type.
m_axi_arvalid	O	The assertion of this signal means there is a valid read request to the address on m_axi_araddr.
m_axi_arready	I	Master read address ready.
m_axi_arlock	O	Lock type.
m_axi_arcache[3:0]	O	Memory type.
m_axi_arburst[1:0]	O	Master read burst type.

Table: AXI4 Memory Mapped DMA Read Interface Signals

Signal Name	I/O	Description
m_axi_rdata [C_M_AXI_DATA_WIDTH-1:0]	I	Master read data.
m_axi_rid [3:0]	I	Master read ID.
m_axi_rresp[1:0]	I	Master read response.
m_axi_rlast	I	Master read last.
m_axi_rvalid	I	Master read valid.
m_axi_rready	O	Master read ready.
m_axi_ruser [C_M_AXI_DATA_WIDTH/8-1:0]	I	Master read odd data parity, per byte. This port is enabled only in Data Protection mode.

Table: AXI4 Memory Mapped DMA Write Address Interface Signals

Signal Name	I/O	Description
m_axi_awaddr [C_M_AXI_ADDR_WIDTH-1:0]	O	This signal is the address for a memory mapped write to the user logic from the DMA.
m_axi_awid[3:0]	O	Master write address ID.
m_axi_awuser[31:0]	O	m_axi_awuser[18:0] = reserved m_axi_awuser[31:19] = queue number
m_axi_awlen[7:0]	O	Master write address length.
m_axi_awsize[2:0]	O	Master write address size.
m_axi_awburst[1:0]	O	Master write address burst type.
m_axi_awprot[2:0]	O	Protection type.
m_axi_awvalid	O	The assertion of this signal means there is a valid write request to the address on m_axi_araddr.
m_axi_awready	I	Master write address ready.
m_axi_awlock	O	Lock type.
m_axi_awcache[3:0]	O	Memory type.

Table: AXI4 Memory Mapped DMA Write Interface Signals

Signal Name	I/O	Description
m_axi_wdata [C_M_AXI_DATA_WIDTH-1:0]	O	Master write data.
m_axi_wlast	O	Master write last.
m_axi_wstrb[31:0]	O	Master write strobe.
m_axi_wvalid	O	Master write valid.
m_axi_wready	I	Master write ready.
m_axi_wuser [C_M_AXI_DATA_WIDTH/8-1:0]	O	Master write user. m_axi_wuser[C_M_AXI_DATA_WIDTH/8-1:0] = write data odd parity, per byte. This port is enabled only in Data Protection mode.

Table: AXI4 Memory Mapped DMA Write Response Interface Signals

Signal Name	I/O	Description
m_axi_bvalid	I	Master write response valid.
m_axi_bresp[1:0]	I	Master write response.
m_axi_bid[3:0]	I	Master response ID.
m_axi_bready	O	Master response ready.

AXI4-Stream H2C Ports

Table: AXI4-Stream H2C Port Descriptions

Port Name	I/O	Description
m_axis_h2c_tdata [AXI_DATA_WIDTH-1:0]	O	Data output for H2C AXI4-Stream.
m_axis_h2c_tcrc [31:0]	O	32-bit CRC value for that beat. IEEE 802.3 CRC-32 Polynomial
m_axis_h2c_tuser_qid[10:0]	O	Queue ID
m_axis_h2c_tuser_port_id[2:0]	O	Port ID
m_axis_h2c_tuser_err	O	If set, indicates the packet has an error. The error could come from the PCIe, or the error could be in the DMA transfer. AMD recommends that you look at the error registers and context for details. When the DMA first detects the error, the error bit is set to 1. And the error bit is set for the remainder of that packet. The error bit is reset to 0 for the next packet if there are no errors in that packet.
m_axis_h2c_tuser_mdata[31:0]	O	Metadata In internal mode, QDMA passes the lower 32 bits of the H2C AXI4-Stream descriptor on this field.
m_axis_h2c_tuser_mty[5:0]	O	The number of bytes that are invalid on the last beat of the transaction. This field is 0 for a 64B transfer.
m_axis_h2c_tuser_zero_byte	O	When set, it indicates that the current beat is an empty beat (zero bytes are being transferred).
m_axis_h2c_tvalid	O	Valid

Port Name	I/O	Description
m_axis_h2c_tlast	O	Indicates that this is the last cycle of the packet transfer.
m_axis_h2c_tready	I	Ready

AXI4-Stream C2H Ports

Table: AXI4-Stream C2H Port Descriptions

Port Name	I/O	Description
s_axis_c2h_tdata [AXI_DATA_WIDTH-1:0]	I	Supported data widths 128 bits, 256 bits, and 512 bits. Every C2H data packet has a corresponding C2H completion packet.
s_axis_c2h_tcrc [31:0]	I	32 bit CRC value for that beat. IEEE 802.3 CRC-32 Polynomial IP samples CRC value only when s_axis_c2h_tlast is asserted.
s_axis_c2h_ctrl_len [15:0]	I	Length of the packet. For ZERO byte write, the length is 0. C2H stream packet data length is limited to 31 * c2h buffer size. In older versions (such as 2018.3), C2H stream packet data length was limited to 7 * C2H buffer size. ctrl_len is in bytes and should be valid in first beat of the packet.
s_axis_c2h_ctrl_qid [10:0]	I	Queue ID.
s_axis_c2h_ctrl_has_cmpt	I	1'b1: The data packet has a completion. 1'b0: The data packet doesn't have a completion.
s_axis_c2h_ctrl_marker	I	Marker message used for making sure pipeline is completely flushed. After that, you can safely do queue invalidation.
s_axis_c2h_ctrl_port_id [2:0]	I	Port ID.
s_axis_c2h_ctrl_ecc[6:0]	I	Sideband protection for C2H control signals. Output of the AMD Error Correction Code (ECC) core. ECC IP input is described below.
s_axis_c2h_mty [5:0]	I	Empty byte should be set in last beat.

Port Name	I/O	Description
s_axis_c2h_tvalid	I	Valid.
s_axis_c2h_tlast	I	Indicates that this is the last cycle of the packet transfer.
s_axis_c2h_tready	O	Ready.

To generate ECC signals for C2H control bus s_axis_c2h_ctrl_ecc[6:0], use AMD error correction code (ECC) IP. Input signals to ECC IP are listed below and you have to maintain the order that is listed below.

Input to ECC IP Using ecc_data_in[56:0]

```
assign ecc_data_in[56:0] = { 24'h0, //reserved
                           s_axis_c2h_ctrl_has_cmpt, //has compt
                           s_axis_c2h_ctrl_marker, //marker
                           s_axis_c2h_ctrl_port_id, //port_id
                           1'b0, // reserved should be set to 0.
                           s_axis_c2h_ctrl_qid, // Qid
                           s_axis_c2h_ctrl_len}; //length
```

AXI4-Stream C2H Completion Ports

Table: AXI4-Stream C2H Completion Port Descriptions

Port Name	I/O	Description
s_axis_c2h_cmpt_tdata[511:0]	I	Completion data from the user application. This contains information that is written to the completion ring in the host.
s_axis_c2h_cmpt_size [1:0]	I	00: 8B completion. 01: 16B completion. 10: 32B completion. 11: 64B completion
s_axis_c2h_cmpt_dpar [15:0]	I	Odd parity computed as bit per 32b. s_axis_c2h_cmpt_dpar[0] is parity over s_axis_c2h_cmpt_tdata[31:0]. s_axis_c2h_cmpt_dpar[1] is parity over s_axis_c2h_cmpt_tdata[63:31] and so on.
s_axis_c2h_cmpt_ctrl_qid[10:0]	I	Completion queue ID.

Port Name	I/O	Description
s_axis_c2h_cmpt_ctrl_marker	I	Marker message used for making sure pipeline is completely flushed. After that, you can safely do queue invalidation.
s_axis_c2h_cmpt_ctrl_user_trig	I	User can trigger the interrupt and the status descriptor write if they are enabled.
s_axis_c2h_cmpt_ctrl_cmpt_type[1:0]		2'b00: NO_PLD_NO_WAIT. The CMPT packet does not have a corresponding payload packet, and it does not need to wait. 2'b01: NO_PLD_BUT_WAIT. The CMPT packet does not have a corresponding payload packet; however, it still needs to wait for the payload packet to be sent before sending the CMPT packet. 2'b10: RSVD. 2'b11: HAS_PLD. The CMPT packet has a corresponding payload packe, and it needs to wait for the payload packet to be sent before sending the CMPT packet.
s_axis_c2h_cmpt_ctrl_wait_pld_pkt_id[15:0]		The data payload packet ID that the CMPT packet needs to wait for before it can be sent.
s_axis_c2h_cmpt_ctrl_port_id[2:0]	I	Port ID.
s_axis_c2h_cmpt_ctrl_col_idx[2:0]	I	Color index that defines if the user wants to have the color bit in the CMPT packet and the bit location of the color bit if present.
s_axis_c2h_cmpt_ctrl_err_idx[2:0]	I	Error index that defines if the user wants to have the error bit in the CMPT packet and the bit location of the error bit if present.
s_axis_c2h_cmpt_ctrl_no_wrb_marker		Disables CMPT packet during Marker transfer. 1'b0 : CMPT packets are sent to CMPT ring 1'b1 : CMPT packets are not sent to CMPT ring.
s_axis_c2h_cmpt_tvalid	I	Valid. s_axis_c2h_cmpt_tvalid must be asserted until s_axis_c2h_cmpt_tready is asserted.
s_axis_c2h_cmpt_tready	O	Ready.

AXI4-Stream Status Ports

Table: AXI-ST C2H Status Port Descriptions

Port Name	I/O	Description
axis_c2h_status_valid	O	Valid per descriptor.
axis_c2h_status_qid [10:0]	O	QID of the packet.
axis_c2h_status_drop	O	The QDMA Subsystem for PCIe drops the packet if it does not have enough descriptors to transfer the full packet to the host. This bit indicates if the packet was dropped or not. A packet that is not dropped is considered as having been accepted. 0: Packet is not dropped. 1: Packet is dropped.
axis_c2h_status_last	O	Last descriptor.
axis_c2h_status_cmp	O	0: Dropped packet or C2H packet with has_cmpt of 1'b0. 1: C2H packet that has completions.
axis_c2h_status_error	O	When axis_c2h_status_error is set to 1, the descriptor fetched has an error. When set to 0, there is no error.

AXI4-Stream C2H Write Completion Ports

Table: AXI-ST C2H Write Completion Port Descriptions

Port Name	I/O	Description
axis_c2h_dmawr_cmp	O	This signal is asserted when the last data payload write request of the packet gets the write completion. It is one pulse per packet.

VDM Ports

Table: VDM Port Descriptions

Port Name	I/O	Description
st_rx_msg_valid	O	Valid
st_rx_msg_data[31:0]	O	Beat 1: {REQ_ID[15:0], VDM_MSG_CODE[7:0], VDM_MSG_ROUTING[2:0], VDM_DW_LENGTH[4:0]} Beat 2: VDM Lower Header [31:0]

Port Name	I/O	Description
		or {((Payload_length=0), VDM Higher Header [31:0])} Beat 3 to Beat <n>: VDM Payload
st_rx_msg_last	O	Indicate the last beat
st_rx_msg_rdy	I	Ready. <hr/>  Note: When this interface is not used, Ready must be tied-off to 1.

RX Vendor Defined Messages are stored in shallow FIFO before they are transmitted to output ports. When there are many back to back VDM messages, FIFO overflows and these messages are dropped. It is best to repeat VDM messages at regular intervals.

Configuration Extend Interface Ports

The Configuration Extend interface allows the core to transfer configuration information with the user application when externally implemented configuration registers are implemented.

Table: Configuration Extend Interface Port Descriptions

Port Name	I/O	Width	Description
cfg_ext_read_received	O	1	Configuration Extend Read Received The core asserts this output when it has received a configuration read request from the link. Set when PCI Express Extended Configuration Space Enable is selected in the user defined configuration Capabilities tab in in the AMD Vivado™ IDE.

Port Name	I/O	Width	Description
			<ul style="list-style-type: none"> • All received configuration reads with cfg_ext_register_number in the range of 0xb0-0xbf is considered to be PCIe Legacy Extended Configuration Space. • All received configuration reads with cfg_ext_register_number in the range of 0x120-13F is considered to be PCIe Extended Configuration Space. • All received configuration reads regardless of their address will be indicated by 1 cycle assertion of cfg_ext_read_received. Valid data is driven on cfg_ext_register_number and cfg_ext_function_number. • Only received configuration reads within the two aforementioned ranges need to be responded by the user application outside of the IP.
cfg_ext_write_received	O	1	<p>Configuration Extend Write Received The core asserts this output when it has received a configuration write request from the link. Set when PCI Express Extended Configuration Space Enable is selected in Capabilities tab in the Vivado IDE.</p> <ul style="list-style-type: none"> • Data corresponding to all received configuration writes with cfg_ext_register_number in the range 0xb0-0xbf is presented on cfg_ext_register_number, cfg_ext_function_number, cfg_ext_write_data and cfg_ext_write_byte_enable. • All received configuration writes with cfg_ext_register_number in the range 0x120-13F is presented on cfg_ext_register_number, cfg_ext_function_number, cfg_ext_write_data and cfg_ext_write_byte_enable.

Port Name	I/O	Width	Description
cfg_ext_register_number	O	10	Configuration Extend Register Number The 10-bit address of the configuration register being read or written. The data is valid when cfg_ext_read_received or cfg_ext_write_received is High.
cfg_ext_function_number	O	8	Configuration Extend Function Number. The 8-bit function number corresponding to the configuration read or write request. The data is valid when cfg_ext_read_received or cfg_ext_write_received is High.
cfg_ext_write_data	O	32	Configuration Extend Write Data Data being written into a configuration register. This output is valid when cfg_ext_write_received is High.
cfg_ext_write_byte_enable	O	4	Configuration Extend Write Byte Enable Byte enables for a configuration write transaction.
cfg_ext_read_data	I	32	Configuration Extend Read Data You can provide data from an externally implemented configuration register to the core through this bus. The core samples this data on the next positive edge of the clock after it sets cfg_ext_read_received High, if you have set cfg_ext_read_data_valid.

Port Name	I/O	Width	Description
cfg_ext_read_data_valid	I	1	Configuration Extend Read Data Valid The user application asserts this input to the core to supply data from an externally implemented configuration register. The core samples this input data on the next positive edge of the clock after it sets cfg_ext_read_received High. The core expects the assertions of this signal within 262144 ('h4_0000) clock cycles of user clock after receiving the read request on cfg_ext_read_received signal. If no response is received by this time, the core will send auto-response with 'h0 payload, and the user application must discard the response and terminate that particular request immediately

FLR Ports

Table: FLR Port Descriptions

Port Names	I/O	Description
usr_flr_fnc [7:0]	O	Function The function number of the FLR status change.
usr_flr_set	O	Set Asserted for 1 cycle indicating that the FLR status of the function indicated on usr_flr_fnc[7:0] is active.
usr_flr_done_fnc [7:0]	I	Done Function The function for which FLR has been completed by user logic.
usr_flr_done_vld	I	Done Valid Assert for one cycle to signal that FLR for the function on usr_flr_done_fnc[7:0] has been completed.

QDMA Descriptor Bypass Input Ports

Table: QDMA H2C-Streaming Bypass Input Port Descriptions

Port Name	I/O	Description
-----------	-----	-------------

Port Name	I/O	Description
h2c_byp_in_st_addr [63:0]	I	64-bit starting address of the DMA transfer.
h2c_byp_in_st_len [15:0]	I	The number of bytes to transfer.
h2c_byp_in_st_sop	I	Indicates start of packet. Set for the first descriptor. Reset for the rest of the descriptors.
h2c_byp_in_st_eop	I	Indicates end of packet. Set for the last descriptor. Reset for the rest of the descriptors.
h2c_byp_in_st_sdi	I	<p>H2C Bypass In Status Descriptor/Interrupt If set, it is treated as an indication from the user application to the QDMA to send the status descriptor to host, and to generate an interrupt to host when the QDMA has fetched the last byte of the data associated with this descriptor. The QDMA honors the request to generate an interrupt only if interrupts are enabled in the H2C SW context for this QID and armed by the driver. This can only be set for an EOP descriptor.</p> <p>QDMA hangs if the last descriptor without h2c_byp_in_st_sdi has an error. This results in a missing writeback and hw_ctxt.dsc_pend bit that are asserted indefinitely. The workaround is to send a zero length descriptor to trigger the Completion (CMPT) Status.</p> <hr/> <p>⌚ Recommended: For performance reasons, AMD recommends that this port be asserted once in 32 or 64 descriptors and assert at the last descriptor if there are no more descriptors left.</p>
h2c_byp_in_st_mrkr_req	I	<p>H2C Bypass In Marker Request When set, the descriptor passes through the H2C Engine pipeline and once completed, produces a marker response on the Queue Status port interface. This can only be set for an EOP descriptor.</p>
h2c_byp_in_st_no_dma	I	<p>H2C Bypass In No DMA When sending in a descriptor through this interface with this signal asserted, it informs the QDMA to not send any PCIe requests for this descriptor. Because no PCIe request is sent out, no corresponding DMA data is issued on the H2C</p>

Port Name	I/O	Description
		<p>Streaming output interface.</p> <p>This is typically used in conjunction with h2c_byp_in_st_sdi to cause Status Descriptor/ Interrupt when the user logic is out of the actual descriptors and still wants to drive the h2c_byp_in_st_sdi signal.</p> <p>If h2c_byp_in_st_mrkr_req and h2c_byp_in_st_sdi are reset when sending in a no-DMA descriptor, the descriptor is treated as a NOP and is completely consumed inside the QDMA without any interface activity.</p> <p>If h2c_byp_in_st_no_dma is set, then both h2c_byp_in_st_sop and h2c_byp_in_st_eop must be set.</p> <p>If h2c_byp_in_st_no_dma is set, the QDMA ignores the address and length fields of this interface.</p>
h2c_byp_in_st_qid [10:0]	I	The QID associated with the H2C descriptor ring.
h2c_byp_in_st_error	I	This bit can be set to indicate an error for the queue. The descriptor is not processed. Context is updated to reflect an error in the queue.
h2c_byp_in_st_func [7:0]	I	PCIe function ID
h2c_byp_in_st_cidx [15:0]	I	The CIDX that is used for the status descriptor update and/or interrupt (aggregation mode). Generally the CIDX should be left unchanged from when it was received from the descriptor bypass output interface.
h2c_byp_in_st_port_id [2:0]	I	QDMA port ID
h2c_byp_in_st_vld	I	Valid. High indicates descriptor is valid, one pulse for one descriptor.
h2c_byp_in_st_rdy	O	Ready to take in descriptor

Table: QDMA H2C-MM Descriptor Bypass Input Port Descriptions

Port Name	I/O	Description
h2c_byp_in_mm_radr[63:0]	I	The read address for the DMA data.
h2c_byp_in_mm_wadr[63:0]	I	The write address for the dma data.

Port Name	I/O	Description
h2c_byp_in_mm_no_dma	I	<p>H2C Bypass In No DMA</p> <p>When sending in a descriptor through this interface with this signal asserted, this signal informs the QDMA to not send any PCIe requests for this descriptor. Because no PCIe request is sent out, no corresponding DMA data is issued on the H2C MM output interface.</p> <p>This is typically used with h2c_byp_in_mm_sdi to cause Status Descriptor/Interrupt when the user logic is out of the actual descriptors and still wants to drive the h2c_byp_in_mm_sdi signal.</p> <p>If h2c_byp_in_mm_mrkr_req and h2c_byp_in_mm_sdi are reset when sending in a no-DMA descriptor, the descriptor is treated as a No Operation (NOP) and is completely consumed inside the QDMA without any interface activity.</p> <p>If h2c_byp_in_mm_no_dma is set, the QDMA ignores the address. The length field should be set to 0.</p>
h2c_byp_in_mm_len[27:0]	I	<p>The DMA data length.</p> <p>The upper 12 bits must be tied to 0. Thus only the lower 16 bits of this field can be used for specifying the length.</p>
h2c_byp_in_mm_sdi	I	<p>H2C-MM Bypass In Status Descriptor/Interrupt</p> <p>If set, it is treated as an indication from the User to QDMA to send the status descriptor to host and generate an interrupt to host when the QDMA has fetched the last byte of the data associated with this descriptor. The QDMA honors the request to generate an interrupt only if interrupts are enabled in the H2C ring context for this QID and armed by the driver.</p> <p>QDMA hangs if the last descriptor without h2c_byp_in_mm_sdi has an error. This results in a missing writeback and hw_ctxt.dsc_pend bit is asserted indefinitely. The workaround is to send a zero length descriptor to trigger the Completion (CMPT) Status.</p> <hr/> <p> Recommended: For performance reasons, AMD recommends that this port be asserted once in 32</p>

Port Name	I/O	Description
		or 64 descriptors and be asserted at the last descriptor if there are no more descriptors left.
h2c_byp_in_mm_mrkr_req	I	H2C-MM Bypass In Marker Request Indication from the user that the QDMA must send a completion status to the User once the QDMA has completed the data transfer of this descriptor.
h2c_byp_in_mm_qid [10:0]	I	The QID associated with the H2C descriptor ring.
h2c_byp_in_mm_error	I	This bit can be set to indicate an error for the queue. The descriptor is not processed. Context is updated to reflect an error in the queue.
h2c_byp_in_mm_func [7:0]	I	PCIe function ID
h2c_byp_in_mm_cidx [15:0]	I	The CIDX that is used for the status descriptor update and/or interrupt (aggregation mode). Generally the CIDX should be left unchanged from when it was received from the descriptor bypass output interface.
h2c_byp_in_mm_port_id [2:0]	I	QDMA port ID
h2c_byp_in_mm_vld	I	Valid. High indicates descriptor is valid, one pulse for one descriptor.
h2c_byp_in_mm_rdy	O	Ready to take in descriptor

Table: QDMA C2H-Streaming Bypass Input Port Descriptions ¹

Port Name	I/O	Description
c2h_byp_in_st_csh_addr [63:0]	I	64 bit address where DMA writes data.
c2h_byp_in_st_csh_qid [10:0]	I	The QID associated with the C2H descriptor ring.
c2h_byp_in_st_csh_error	I	This bit can be set to indicate an error for the queue. The descriptor is not processed. Context is updated to reflect an error in the queue.
c2h_byp_in_st_csh_func [7:0]	I	PCIe function ID

Port Name	I/O	Description
c2h_byp_in_st_csh_port_id[2:0]	I	QDMA port ID
c2h_byp_in_st_csh_pfch_tag[6:0]	I	Prefetch tag. The prefetch tag points to the cam that stores the active queues in the prefetch engine. In Cache Bypass mode, you must loop back c2h_byp_out_pfch_tag[6:0] to c2h_byp_in_st_csh_pfch_tag[6:0]. In Simple Bypass mode, used need to pass in the Prefetch tag value from MDMA_C2H_PFCH_BYP_TAG (0x140C) register.
c2h_byp_in_st_csh_vld	I	Valid. High indicates descriptor is valid, one pulse for one descriptor.
c2h_byp_in_st_csh_rdy	O	Ready to take in descriptor.
1. AXI4-Stream C2H Simple Bypass mode and Cache Bypass mode both use the same bypass ports, c2h_byp_in_st_csh_*.		

Table: QDMA C2H-MM Descriptor Bypass Input Port Descriptions

Port Name	I/O	Description
c2h_byp_in_mm_raddr [63:0]	I	The read address for the DMA data.
c2h_byp_in_mm_waddr[63:0]	I	The write address for the DMA data.
c2h_byp_in_mm_no_dma	I	C2H Bypass In No DMA When sending in a descriptor through this interface with this signal asserted, this signal informs the QDMA to not send any PCIe requests for this descriptor. Because no PCIe request is sent out, no corresponding DMA data is read from C2H MM interface. This is typically used with c2h_byp_in_mm_sdi to cause Status Descriptor/Interrupt when the user logic is out of the actual descriptors and still wants to drive the c2h_byp_in_mm_sdi signal. If c2h_byp_in_mm_mrkr_req and c2h_byp_in_mm_sdi are reset when sending in a no-DMA descriptor, the descriptor is treated as a NOP and is completely consumed inside the QDMA without any interface activity.

Port Name	I/O	Description
		If c2h_byp_in_mm_no_dma is set, the QDMA ignores the address. The length field should be set to 0.
c2h_byp_in_mm_len[27:0]	I	The DMA data length
c2h_byp_in_mm_sdi	I	C2H Bypass In Status Descriptor/Interrupt If set, it is treated as an indication from the User to QDMA to send the status descriptor to host, and generate an interrupt to host when the QDMA has fetched the last byte of the data associated with this descriptor. The QDMA honors the request to generate an interrupt only if interrupts are enabled in the C2H ring context for this QID and armed by the driver. 🔧 Recommended: For performance reasons, AMD recommends that this port be asserted once in 32 or 64 descriptors, and asserted at the last descriptor if there are no more descriptors left.
c2h_byp_in_mm_mrkr_req	I	C2H Bypass In Marker Request Indication from the User that the QDMA must send a completion status to the User once the QDMA has completed the data transfer of this descriptor.
c2h_byp_in_mm_qid [10:0]	I	The QID associated with the C2H descriptor ring
c2h_byp_in_mm_error	I	This bit can be set to indicate an error for the queue. The descriptor is not processed. Context is updated to reflect an error in the queue.
c2h_byp_in_mm_func [7:0]	I	PCIe function ID
c2h_byp_in_mm_cidx [15:0]	I	The User must echo the CIDX from the descriptor that it received on the bypass-out interface.
c2h_byp_in_mm_port_id[2:0]	I	QDMA port ID
c2h_byp_in_mm_vld	I	Valid. High indicates descriptor is valid, one pulse for one descriptor.
c2h_byp_in_mm_rdy	O	Ready to take in descriptor.

QDMA Descriptor Bypass Output Ports

Table: QDMA H2C Descriptor Bypass Output Port Descriptions

Port Name	I/O	Description
h2c_byp_out_dsc [255:0]	0	The H2C descriptor fetched from the host. For H2C AXI-MM, the subsystem uses all 256 bits, and the structure of the bits are the same as this table . For H2C AXI-ST, the subsystem uses [127:0] bits, and the structure of the bits are the same as this table .
h2c_byp_out_st_mm	0	Indicates whether this is a streaming data descriptor or memory-mapped descriptor. 0: streaming 1: memory-mapped
h2c_byp_out_dsc_sz [1:0]	0	Descriptor size. This field indicates size of the descriptor. 0: 8B 1: 16B 2: 32B 3: 64B - 64B descriptors will be transferred with two valid/ready cycles. The first cycle has the least significant 32 bytes. The second cycle has the most significant 32 bytes. CIDX and other queue information is valid only on the second beat of a 64B descriptor .
h2c_byp_out_qid [10:0]	0	The QID associated with the H2C descriptor ring.
h2c_byp_out_error	0	Indicates that an error was encountered in descriptor fetch or execution of a previous descriptor.
h2c_byp_out_func [7:0]	0	PCIe function ID
h2c_byp_out_cidx [15:0]	0	H2C Bypass Out Consumer Index The ring index of the descriptor fetched. The User must echo this field back to QDMA when submitting the descriptor on the bypass-in interface.
h2c_byp_out_port_id [2:0]	0	QDMA port ID
h2c_byp_out_fmt[2:0]	0	Format The encoding for this field is as follows.

Port Name	I/O	Description
		0x0: Standard descriptor 0x1 - 0x7: Reserved
h2c_byp_out_vld	O	Valid. High indicates descriptor is valid, one pulse for one descriptor.
h2c_byp_out_rdy	I	Ready. When this interface is not used, Ready must be tied-off to 1.

Table: QDMA C2H Descriptor Bypass Output Port Descriptions

Port Name	I/O	Description
c2h_byp_out_dsc [255:0]	O	The C2H descriptor fetched from the host. For C2H AXI-MM, the subsystem uses all 256 bits, and the structure of the bits is the same as this table . For C2H AXI-ST, the subsystem uses [63:0] bits, and the structure of the bits is the same as this table . The remaining bits are ignored.
c2h_byp_out_st_mm	O	Indicates whether this is a streaming data descriptor or memory-mapped descriptor. 0: streaming 1: memory-mapped
c2h_byp_out_dsc_sz [1:0]	O	Descriptor size. This field indicates the amount of valid descriptor information on h2c_byp_out_dsc. 0: 8B 1: 16B 2: 32B 3: 64B - 64B descriptors will be transferred with two valid/ready cycles. The first cycle has the least significant 32 bytes. The second cycle has the most significant 32 bytes. CIDX and other queue information is valid only on the second beat of a 64B descriptor.
c2h_byp_out_qid [10:0]	O	The QID associated with the H2C descriptor ring.
c2h_byp_out_error	O	Indicates that an error was encountered in descriptor fetch or execution of a previous descriptor.
c2h_byp_out_func [7:0]	O	PCIe function ID.

Port Name	I/O	Description
c2h_byp_out_cidx [15:0]	O	C2H Bypass Out Consumer Index The ring index of the descriptor fetched. The User must echo this field back to QDMA when submitting the descriptor on the bypass-in interface.
c2h_byp_out_port_id [2:0]	O	QDMA port ID
c2h_byp_out_pfch_tag[6:0]	O	Prefetch tag. The prefetch tag points to the cam that stores the active queues in prefetch engine
c2h_byp_out_fmt[2:0]	O	Format The encoding for this field is as follows. 0x0 : Standard descriptor 0x1 - 0x7 : Reserved
c2h_byp_out_vld	O	Valid. High indicates descriptor is valid, one pulse for one descriptor.
c2h_byp_out_rdy	I	Ready. When this interface is not used, Ready must be tied-off to 1.

It is common for h2c_byp_out_vld or c2h_byp_out_vld to be asserted with the CIDX value; this occurs when the Descriptor bypass mode option is not set in the context programming selection. You must set the Descriptor bypass mode during QDMA IP core customization in the AMD Vivado™ IDE to see descriptor bypass output ports. When Descriptor bypass option is selected in the AMD Vivado™ IDE but the descriptor bypass bit is not set in context programming, you will see valid signals getting asserted with CIDX updates.

QDMA Descriptor Credit Input Ports

Table: QDMA Descriptor Credit Input Port Descriptions

Port Name	I/O	Description
dsc_crdt_in_vld	I	Valid. When asserted the user must be presenting valid data on the bus and maintain the bus values until both valid and ready are asserted on the same cycle.
dsc_crdt_in_rdy	O	Ready. Assertion of this signal indicates the DMA is ready to accept data from this bus.
dsc_crdt_in_dir	I	Indicates whether credits are for H2C or C2H descriptor ring. 0: H2C

Port Name	I/O	Description
		1: C2H
dsc_crdt_in_fence	I	If the fence bit is set, the credits are not coalesced, and the queue is guaranteed to generate a descriptor fetch before subsequent credit updates are processed. The fence bit should only be set for a queue that is enabled, and has both descriptors and credits available, otherwise a hang condition might occur.
dsc_crdt_in_qid [10:0]	I	The QID associated with the descriptor ring for the credits are being added.
dsc_crdt_in_crdt [15:0]	I	The number of descriptor credits that the user application is giving to QDMA Subsystem for PCIe to fetch descriptors from the host.

QDMA Traffic Manager Credit Output Ports

Table: QDMA TM Credit Output Port Descriptions

Port Name	I/O	Description
tm_dsc_sts_vld	O	Valid. Indicates valid data on the output bus. Valid data on the bus is held until tm_dsc_sts_rdy is asserted by the user.
tm_dsc_sts_rdy	I	Ready. Assertion indicates that the user logic is ready to accept the data on this bus. When this interface is not used, Ready must be tied-off to 1. Note: When this interface is not used, Ready must be tied-off to 1.
tm_dsc_sts_byp	O	Shows the bypass bit in the SW descriptor context
tm_dsc_sts_dir	O	Indicates whether the status update is for a H2C or C2H descriptor ring. 0: H2C 1: C2H
tm_dsc_sts_mm	O	Indicates whether the status update is for a streaming or memory-mapped queue. 0: streaming 1: memory-mapped
tm_dsc_sts_qid [10:0]	O	The QID of the ring

Port Name	I/O	Description
tm_dsc_sts_avl [15:0]	O	If tm_dsc_sts_qinv is set, this is the number of credits available in the descriptor engine. If tm_dsc_sts_qinv is not set this is the number of new descriptors that have been posted to the ring since the last time this update was sent.
tm_dsc_sts_qinv	O	If set, it indicates that the queue has been invalidated. This is used by the user application to reconcile the credit accounting between the user application and QDMA.
tm_dsc_sts_qen	O	The current queue enable status.
tm_dsc_sts_irq_arm	O	If set, it indicates that the driver is ready to accept interrupts
tm_dsc_sts_error	O	Set to 1 if the PIDX update is rolled over the current CIDX of associated queue.
tm_dsc_sts_pidx[15:0]	O	PIDX of the Queue
tm_dsc_sts_port_id [2:0]	O	The port id associated with the queue from the queue context.

User Interrupts

Table: User Interrupts Port Descriptions

Port Name	I/O	Description
usr_irq_in_vld	I	Valid An assertion indicates that an interrupt associated with the vector, function, and pending fields on the bus should be generated to PCIe. Once asserted, Usr_irq_in_vld must remain high until usr_irq_out_ack is asserted by the DMA.
usr_irq_in_vec [11:0]	I	Vector The MSIX vector to be sent. Vector number is the index for MSI_X table. Vector 0 is the first vector for that function. Each function has 8 vectors.
usr_irq_in_fnc [7:0]	I	Function The function of the vector to be sent.

Port Name	I/O	Description
usr_irq_out_ack	O	Interrupt Acknowledge An assertion of the acknowledge bit indicates that the interrupt was transmitted on the link the user logic must wait for this pulse before signaling another interrupt condition.
usr_irq_out_fail	O	Interrupt Fail An assertion of fail indicates that the interrupt request is aborted before transmission on the link. The interrupt fail can happen for several reasons, for details, see the following MSIX Interrupt Options section.

 **Note:** Maximum eight vectors are allowed per function.

MSIX Interrupt Options

The following table describes the possible scenarios and outcomes:

Table: MSIX Interrupt Options

MSIX Enable (MSIX Enable Mask)	MSIX Mapable (MSIX Mapable Mask)	MSIX Vector (MSIX Table Vector)	MSIX Fail Output	Interrupt at Host
0	X	X	usr_irq_fail	Not Received
1	0	0	usr_irq_ack	Received
1	0	1	usr_irq_ack with PBA bit set	Not Received
1	1	0	usr_irq_ack with PBA bit set	Not Received
1	1	1	usr_irq_ack with PBA bit set	Not Received

Queue Status Ports

All marker responses are sent out through this interface. If the IP is configured in an internal mode (no descriptor bypass), there are some marker responses sent out on this interface for packet completion. You can ignore this interface or set the `mrkr_dis` signal bit in the context settings.

Table: Queue Status Ports

Port Name	I/O	Description
qsts_out_op[7:0]	0	Opcode This indicates the type of packet being issued. Encoding of this field is as follows. 0x0: CMPT Marker Response 0x1: H2C-ST Marker Response 0x2: C2H-MM Marker Response 0x3: H2C-MM Marker Response 0x4-0xff: reserved
qsts_out_data[63:0]	0	The data field for the individual opcodes are defined in the tables Table 1 and Table 2 .
qsts_out_port_id[2:0]	0	Port ID
qsts_out_qid[11:0]	0	Queue ID
qsts_out_vld	0	Queue status valid
qsts_out_rdy	I	Queue status ready. Ready must be tied to 1 so status output will not be blocked. Even if this interface is not used, the ready port must be tied to 1.

Queue Status Data (qsts_out_data)

Table: Description for CMPT Marker Response

qsts_out_data	Field	Description
[1:0]	err	Error code reported by the CMPT Engine. 0: No error 1: SW gave bad Completion CIDX update 2: Descriptor error received while processing the C2H packet 3: Completion dropped by the C2H Engine because Completion Ring was full
[2]	retry_marker_req	An Interrupt could not be generated in spite of being enabled. This happens when an Interrupt is already outstanding on the queue when the marker request was received. The user logic must wait and retry the marker request again if an Interrupt is desired to be sent.

qsts_out_data	Field	Description
[25:3]	marker_cookie	<p>When the CMPT Engine sends a marker to the Queues status port interface, it sends the lower 23 bits of the CMPT as part of the marker response on the Queues status port interface. Thus the user logic can place a 23-bit value in the CMPT when making the marker request and it will receive the same 23 bits with the marker response. When the marker is generated as a result of an error that the CMPT Engine encountered (as opposed to a marker request made by the user logic), then this 23-bit field is not valid.</p> <hr/> <p>☞ Note: Even if the user has enabled stamping of error and/or color bits in the CMPT writes to the host, the marker_cookie does not contain them. It is exactly the lower 23 bits of the CMPT that the user logic provided to the QDMA when making the marker request.</p>
[26]	is_mrkr_rsp	This bit will be set to 1 if the marker response is based on marker request. If this bit is set to '0' marker response is based on errors.
[63:27]	rsv	Reserved

Table: Description of AXI_ST H2C, AXI-MM H2C and C2H Marker Response

qsts_out_data	Field	Description
[15:0]	cidx	<p>For a queue in the bypass mode, this is the CIDX of the descriptor that contains the marker request.</p> <p>For a queue in an internal mode, this is the CIDX of the descriptor that resulted in the marker response issued.</p>
[16]	err	Indicates that an error has occurred on this queue. This error could be in the queue context or a data error. It is expected that appropriate error registers

qsts_out_data	Field	Description
		are read to determine the exact error.
[63:17]	rsv	Reserved

Register Space

This section provides register space information for the QDMA Subsystem for PCIe. In register space descriptions, configuration register attributes are defined as follows:

NA

Reserved

RO

Read-Only - Register bits are read-only and cannot be altered by the software.

RW

Read-Write - Register bits are read-write and are permitted to be either Set or Cleared by the software to the desired state.

RW1C

Write-1-to-clear-status - Register bits indicate status when read. A Set bit indicates a status event which is Cleared by writing a 1b. Writing a 0b to RW1C bits has no effect.

W1C

Non-readable-write-1-to-clear-status - Register will return 0 when read. Writing 1b Clears the status for that bit index. Writing a 0b to W1C bits has no effect.

W1S

Non-readable-write-1-to-set - Register will return 0 when read. Writing 1b Sets the control set for that bit index. Writing a 0b to W1S bits has no effect.

QDMA PF Address Register Space

All the physical function (PF) registers are listed in the qdma_v5_1_pf_registers.csv available in the [Register Reference File](#).

★ Tip: When you generate the IP in default mode, not all registers are exposed. For example, debug registers are missing. Refer to the qdma_v5_1_pf_registers.csv file to identify the debug registers. To expose all registers, use the following tcl command during IP generation:

`set_property CONFIG.debug_mode {DEBUG_REG_ONLY} [get_ips qdma_0]`

Table: QDMA PF Address Register Space

Register Name	Base (Hex)	Byte Size (Dec)	Register List and Details
QDMA_CSR	0x0000	9216	QDMA Configuration Space Register (CSR) found in

Register Name	Base (Hex)	Byte Size (Dec)	Register List and Details
			qdma_v5_1_pf_registers.csv. For QDMA MM register, see QDMA MM Register Offset .
QDMA_TRQ_SEL_QUEUE_PF	0x18000	32768	Also found in QDMA_TRQ_SEL_QUEUE_PF (0x18000) .
QDMA_PF_MAILBOX	0x22400	16384	Also found in QDMA_PF_MAILBOX (0x22400) .
QDMA_TRQ_MSIX	0x30000	32768	Also found in QDMA_TRQ_MSIX (0x30000) .
QDMA_TRQ_MSIX_PBA	0x34000	256	QDMA_TRQ_MSIX_PBA (0x34000)

QDMA_CSR (0x0000)

QDMA Configuration Space Register (CSR) descriptions are accessible in qdma_v5_1_pf_registers.csv available in the [Register Reference File](#).

QDMA MM Register Offset

QDMA_C2H_MM_* QDMA+H2C_MM_* registers represent two engines namely MM0 and MM1. In the register .csv file, it is represented as an array [1:0]. Following are the register offsets for MM0 and MM1:

- QDMA_C2H_MM0 start at 0x1000
- QDMA_C2H_MM1 start at 0x1100
- QDMA_H2C_MM0 start at 0x1200
- QDMA_H2C_MM1 start at 0x1300

QDMA_TRQ_SEL_QUEUE_PF (0x18000)

Table: QDMA_TRQ_SEL_QUEUE_PF (0x18000) Register Space

Register	Address	Description
QDMA_DMAP_SEL_INT_CIDX[2048] (0x18000)	0x18000-0x1cff0	Interrupt Ring Consumer Index (CIDX)
QDMA_DMAP_SEL_H2C_DSC_PIDX[2048]	0x18004-	H2C Descriptor Producer index

Register	Address	Description
(0x18004)	0x1CFF4	(PIDX)
QDMA_DMAP_SEL_C2H_DSC_PIDX[2048] (0x18008)	0x18008- 0x1cff8	C2H Descriptor Producer Index (PIDX)
QDMA_DMAP_SEL_CMPT_CIDX[2048] (0x1800C)	0x1800C- 0x1cffc	C2H Completion Consumer Index (CIDX)

There are 2048 Queues, each Queue will have more than four registers. All these registers can be dynamically updated at any time. This set of registers can be accessed based on the Queue number.

- Queue number is absolute *Qnumber* [0 to 2047].
- Interrupt CIDX address = 0x18000 + *Qnumber**16
- H2C PIDX address = 0x18004 + *Qnumber**16
- C2H PIDX address = 0x18008 + *Qnumber**16
- Write Back CIDX address = 0x1800C + *Qnumber**16

For Queue 0:

- 0x18000 correspond to QDMA_DMAP_SEL_INT_CIDX
- 0x18004 correspond to QDMA_DMAP_SEL_H2C_DSC_PIDX
- 0x18008 correspond to QDMA_DMAP_SEL_C2H_DSC_PIDX
- 0x1800C correspond to QDMA_DMAP_SEL_CMPT_CIDX

For Queue 1:

- 0x18010 correspond to QDMA_DMAP_SEL_INT_CIDX
- 0x18014 correspond to QDMA_DMAP_SEL_H2C_DSC_PIDX
- 0x18018 correspond to QDMA_DMAP_SEL_C2H_DSC_PIDX
- 0x1801C correspond to QDMA_DMAP_SEL_CMPT_CIDX

For Queue 2:

- 0x18020 correspond to QDMA_DMAP_SEL_INT_CIDX
- 0x18024 correspond to QDMA_DMAP_SEL_H2C_DSC_PIDX
- 0x18028 correspond to QDMA_DMAP_SEL_C2H_DSC_PIDX
- 0x1802C correspond to QDMA_DMAP_SEL_CMPT_CIDX

QDMA_DMAP_SEL_INT_CIDX[2048] (0x18000)

Table: QDMA_DMAP_SEL_INT_CIDX[2048] (0x18000)

Bit	Default	Access Type	Field	Description
[31:24]	0	NA	Reserved	Reserved

Bit	Default	Access Type	Field	Description
[23:16]	0	RW	ring_idx	Ring index of the Interrupt Aggregation Ring
[15:0]	0	RW	sw_cidx	Software Consumer index (CIDX)

QDMA_DMAP_SEL_H2C_DSC_PIDX[2048] (0x18004)

Table: QDMA_DMAP_SEL_H2C_DSC_PIDX[2048] (0x18004)

Bit	Default	Access Type	Field	Description
[31:17]	0	NA	Reserved	Reserved
[16]	0	RW	irq_arm	Interrupt arm. Set this bit to 1 for next interrupt generation.
[15:0]	0	RW	h2c_pidx	H2C Producer Index

QDMA_DMAP_SEL_C2H_DSC_PIDX[2048] (0x18008)

Table: QDMA_DMAP_SEL_C2H_DSC_PIDX[2048] (0x18008)

Bit	Default	Access Type	Field	Description
[31:17]	0	NA	Reserved	Reserved
[16]	0	RW	irq_arm	Interrupt arm. Set this bit to 1 for next interrupt generation.
[15:0]	0	RW	c2h_pidx	C2H Producer Index

QDMA_DMAP_SEL_CMPT_CIDX[2048] (0x1800C)

Table: QDMA_DMAP_SEL_CMPT_CIDX[2048] (0x1800C)

Bit	Default	Access Type	Field	Description
[31:29]	0	NA	Reserved	Reserved
[28]	0	RW	irq_en_wrb	Interrupt arm. Set this bit to 1 for next interrupt generation.
[27]	0	RW	en_sts_desc_wrb	Enable Status Descriptor for CMPT
[26:24]	0	RW	trigger_mode	Interrupt and Status Descriptor Trigger Mode: 0x0: Disabled

Bit	Default	Access Type	Field	Description
				0x1: Every 0x2: User_Count 0x3: User 0x4: User_Timer 0x5: User_Timer_Count
[23:20]	0	RW	c2h_timer_cnt_index	Index to QDMA_C2H_TIMER_CNT
[19:16]	0	RW	c2h_count_threshholdIndex	Index to QDMA_C2H_CNT_TH
[15:0]	0	RW	wrb_cidx	CMPT Consumer Index (CIDX)

 **Note:** CMPT CIDX is written from the host. It is recommended to read the updated CMPT CIDX from the Q context.

QDMA_PF_MAILBOX (0x22400)

Table: QDMA_PF_MAILBOX (0x22400) Register Space

Register	Address	Description
Function Status Register (0x22400)	0x22400	Status bits
Function Command Register (0x22404)	0x22404	Command register bits
Function Interrupt Vector Register (0x22408)	0x22408	Interrupt vector register
Target Function Register (0x2240C)	0x2240C	Target Function register
Function Interrupt Vector Register (0x22410)	0x22410	Interrupt Control Register
RTL Version Register (0x22414)	0x22414	RTL Version Register
PF Acknowledgment Registers (0x22420-0x2243C)	0x22420-0x2243C	PF acknowledge
FLR Control/Status Register (0x22500)	0x22500	FLR control and status
Incoming Message Memory (0x22C00-0x22C7C)	0x22C00-0x22C7C	Incoming message (128 bytes)
Outgoing Message Memory (0x23000-0x2307C)	0x23000-0x2307C	Outgoing message (128 bytes)

Mailbox Addressing

PF addressing

$\text{Addr} = \text{PF_Bar_offset} + \text{CSR_addr}$

VF addressing

$\text{Addr} = \text{VF_Bar_offset} + \text{VF_Start_offset} + \text{VF_offset} + \text{CSR_addr}$

Function Status Register (0x22400)

Table: Function Status Register (0x22400)

Bit	Default	Access Type	Field	Description
[31:12]	0	NA	Reserved	Reserved
[11:4]	0	RO	cur_src_fn	This field is for PF use only. The source function number of the message on the top of the incoming request queue.
[2]	0	RO	ack_status	This field is for PF use only. The status bit will be set when any bit in the acknowledgment status register is asserted.
[1]	0	RO	o_msg_status	For VF: The status bit will be set when VF driver write msg_send to its command register. When The associated PF driver send acknowledgment to this VF, the hardware clear this field. The VF driver is not allow to update any content in its outgoing mailbox memory (OMM) while o_msg_status is asserted. Any illegal write to the OMM will be discarded (optionally, this can cause an error in the AXI Lite response channel). For PF: The field indicated the message status of the target FN which is specified in the <i>Target FN Register</i> . The status bit will be set when PF driver sends msg_send command. When the corresponding function driver send acknowledgment by sending msg_rcv, the hardware clear this field. The PF

Bit	Default	Access Type	Field	Description
				driver is not allow to update any content in its outgoing mailbox memory (OMM) while o_msg_status(target_fn_id) is asserted. Any illegal write to the OMM will be discarded (optionally, case an error in the AXI4L response channel).
[0]	0	RO	i_msg_status	For VF: When asserted, a message in the VF's incoming Mailbox memory is pending for process. The field will be cleared once the VF driver write msg_rcv to its command register. For PF: When asserted, the messages in the incoming Mailbox memory are pending for process. The field will be cleared only when the event queue is empty.

Function Command Register (0x22404)

Table: Function Command Register (0x22404)

Bit	Default	Access Type	Field	Description
[31:3]	0	NA	Reserved	Reserved
[2]	0	RO	Reserved	Reserved
[1]	0	RW	msg_rcv	For VF: VF marks the message in its Incoming Mailbox Memory as received. Hardware asserts the acknowledgement bit of the associated PF. For PF: PF marks the message send by target_fn as received. The hardware will refresh the i_msg_status of the PF, and clear the o_msg_status of the target_fn.
[0]	0	RW	msg_send	For VF: VF marks the current message in its own Outgoing Mailbox as valid. For PF:

Bit	Default	Access Type	Field	Description
				<ul style="list-style-type: none"> • Current target_fn_id belongs to a VF: PF finished writing a message into the Incoming Mailbox memory of the VF with target_fn_id. The hardware sets the i_msg_status field of the target FN's status register. • Current target_fn_id belongs to a PF: PF finished writing a message into its own outgoing Mailbox memory. Hardware will push the message to the event queue of the PF with target_fn_id.

Function Interrupt Vector Register (0x22408)

Table: Function Interrupt Vector Register (0x22408)

Bit	Default	Access Type	Field	Description
[31:5]	0	NA	Reserved	Reserved
[4:0]	0	RW	int_vect	5-bit interrupt vector assigned by the driver.

Target Function Register (0x2240C)

Table: Target Function Register (0x2240C)

Bit	Default	Access Type	Field	Description
[31:8]	0	NA	Reserved	Reserved
[7:0]	0	RW	target_fn_id	This field is for PF use only. The FN number which the current operation is targeting at.

Function Interrupt Vector Register (0x22410)

Table: Function Interrupt Vector Register (0x22410)

Bit	Default	Access Type	Field	Description

Bit	Default	Access Type	Field	Description
[31:1]	0	NA	Reserved	Reserved
[0]	0	RW	int_en	Interrupt enable.

RTL Version Register (0x22414)

Table: RTL Version Register (0x22414)

Bit	Default	Access Type	Field	Description
[31:16]	0x1fd3	RO		QDMA ID
[15:0]	0	RO		Vivado versions 0x0100 : QDMA 3.0 Vivado version 2019.1 0x0201 : QDMA3.1 Vivado version 2019.2 Patch 0x0010 : QDMA 4.0 Vivado version 2020.1 0x0020 : QDMA 5.0 Vivado version 2023.1

PF Acknowledgment Registers (0x22420-0x2243C)

Table: PF Acknowledgment Registers (0x22420-0x2243C)

Register	Addr	Default	Access Type	Field	Width	Description
Ack0	0x22420	0	RW		32	Acknowledgment from FN 31~0
Ack1	0x22424	0	RW		32	Acknowledgment from FN 63~32
Ack2	0x22428	0	RW		32	Acknowledgment from FN 95~64
Ack3	0x2242C	0	RW		32	Acknowledgment from FN 127~96
Ack4	0x22430	0	RW		32	Acknowledgment from FN 159~128
Ack5	0x22434	0	RW		32	Acknowledgment from FN 191~160

Register	Addr	Default	Access Type	Field	Width	Description
Ack6	0x22438	0	RW		32	Acknowledgment from FN 223~192
Ack7	0x2243C	0	RW		32	Acknowledgment from FN 255~224

FLR Control/Status Register (0x22500)

Table: FLR Control/Status Register (0x22500)

Bit	Default	Access Type	Field	Description
[31:1]	0	NA	Reserved	Reserved
[0]	0	RW	Flr_status	Software write 1 to initiate the Function Level Reset (FLR) for the associated function. The field is kept asserted during the FLR process. After the FLR is done, the hardware de-asserts this field.

Incoming Message Memory (0x22C00-0x22C7C)

Table: Incoming Message Memory (0x22C00-0x22C7C)

Register	Addr	Default	Access Type	Field	Width	Description
i_msg_i	0x22C00 + i*4	0	RW		32	The <i>i</i> th word of the incoming message (0 ≤ <i>i</i> < 128).

Outgoing Message Memory (0x23000-0x2307C)

Table: Outgoing Message Memory (0x23000-0x2307C)

Register	Addr	Default	Access Type	Field	Width	Description
o_msg_i	0x23000 + i *4	0	RW		32	The <i>i</i> th word of the outgoing message (0 ≤ <i>i</i> < 128).

QDMA_TRQ_MSIX (0x30000)

Table: QDMA_TRQ_MSIX (0x30000)

Byte Offset	Bit	Default	Access Type	Field	Description

Byte Offset	Bit	Default	Access	Type	Field	Description
0x30000	[31:0]	0	RW	addr		MSI-X vector0 message lower address. MSIX_Vector0_Address[63:32]
0x30004	[31:0]	0	RW	addr		MSI-X vector0 message upper address. MSIX_Vector0_Address[63:32]
0x30008	[31:0]	0	RW	data		MSIX_Vector0_Data[31:0] MSI-X vector0 message data.
0x3000C	[31:0]	0	RW	control		MSIX_Vector0_Control[31:0] MSI-X vector0 control. Bit Position: 31:1: Reserved. 0: Mask. When set to 1, this MSI-X vector is not used to generate a message. When reset to 0, this MSI-X vector is used to generate a message.

 **Note:** The table above represents one MSI-X table entry 0. There are 2K MSI-X table entries for the QDMA.

QDMA_TRQ_MSIX_PBA (0x34000)

MSIX pending bit array (PBA) offset is 0x34000. Each MSIX vector has one bit of PBA. For 2K vectors there are total 2K bit 256 bytes of PBA.

QDMA VF Address Register Space

All the virtual function (VF) registers are listed in the qdma_v5_1_vf_registers.csv available in the [Register Reference File](#).

Table: QDMA VF Address Register Space

Target Name	Base (Hex)	Byte Size (Dec)	Notes
QDMA_TRQ_SEL_QUEUE_VF (0x3000)	00003000	4096	VF Direct QCSR (16B per Queue, up to max of 256Queue per function)
QDMA_TRQ_MSIX_VF (0x4000)	00004000	4096	Space for 32 MSIX vectors and PBA

Target Name	Base (Hex)	Byte Size (Dec)	Notes
QDMA_TRQ_MSIX_VF_PBA (0x4800)	00004800	32	MSIX PBA
QDMA_VF_MAILBOX (0x5000)	00005000	8192	Mailbox address space

QDMA_TRQ_SEL_QUEUE_VF (0x3000)

VF functions can access direct update registers per queue with offset (0x3000). The description for this register space is the same as [QDMA_TRQ_SEL_QUEUE_PF \(0x18000\)](#).

This set of registers can be accessed based on Queue number. Queue number is relative Qnumber for that VF.

- Interrupt CIDX address = 0x3000 + Qnumber*16
- H2C PIDX address = 0x3004 + Qnumber*16
- C2H PIDX address = 0x3008 + Qnumber*16
- Completion CIDX address = 0x300C + Qnumber*16

For Queue 0:

- 0x3000 correspond to QDMA_DMAP_SEL_INT_CIDX
- 0x3004 correspond to QDMA_DMAP_SEL_H2C_DSC_PIDX
- 0x3008 correspond to QDMA_DMAP_SEL_C2H_DSC_PIDX
- 0x300C correspond to QDMA_DMAP_SEL_WRB_CIDX

For Queue 1:

- 0x3010 correspond to QDMA_DMAP_SEL_INT_CIDX
- 0x3014 correspond to QDMA_DMAP_SEL_H2C_DSC_PIDX
- 0x3018 correspond to QDMA_DMAP_SEL_C2H_DSC_PIDX
- 0x301C correspond to QDMA_DMAP_SEL_WRB_CIDX

QDMA_TRQ_MSIX_VF (0x4000)

VF functions can access the MSIX table with offset (0x0000) from that function. The description for this register space is the same as [QDMA_TRQ_MSIX \(0x30000\)](#).

QDMA_VF_MAILBOX (0x5000)

Table: QDMA_VF_MAILBOX (0x05000) Register Space

Registers (Address)	Address	Description
---------------------	---------	-------------

Registers (Address)	Address	Description
Function Status Register (0x5000)	0x5000	Status register bits
Function Command Register (0x5004)	0x5004	Command register bits
Function Interrupt Vector Register (0x5008)	0x5008	Interrupt vector register
Target Function Register (0x500C)	0x500C	Target Function register
Function Interrupt Control Register (0x5010)	0x5010	Interrupt Control Register
RTL Version Register (0x5014)	0x5014	RTL Version Register
Incoming Message Memory (0x5800-0x587C)	0x5800-0x587C	Incoming message (128 bytes)
Outgoing Message Memory (0x5C00-0x5C7C)	0x5C00-0x5C7C	Outgoing message (128 bytes)

Function Status Register (0x5000)

Table: Function Status Register (0x5000)

Bit Index	Default	Access Type	Field	Description
[31:12]	0	NA	Reserved	Reserved
[11:4]	0	RO	cur_src_fn	This field is for PF use only. The source function number of the message on the top of the incoming request queue.
[2]	0	RO	ack_status	This field is for PF use only. The status bit will be set when any bit in the acknowledgement status register is asserted.
[1]	0	RO	o_msg_status	For VF: The status bit will be set when VF driver write msg_send to its command register. When the associated PF driver sends

Bit Index	Default	Access Type	Field	Description
				<p>acknowledgement to this VF, the hardware clears this field. The VF driver is not allow to update any content in its outgoing mailbox memory (OMM) while o_msg_status is asserted. Any illegal writes to the OMM are discarded (optionally, case an error in the AXI4-Lite response channel).</p> <p>For PF: The field indicated the message status of the target FN which is specified in the Target FN Register. The status bit is set when PF driver sends the msg_send command. When the corresponding function driver sends acknowledgement through msg_rcv, the hardware clears this field. The PF driver is not allow to update any content in its outgoing mailbox memory (OMM) while o_msg_status(target_fn_id) is asserted. Any illegal writes to the OMM are discarded (optionally, case an error in the AXI4L response channel).</p>
[0]	0	RO	i_msg_status	<p>For VF: When asserted, a message in the VF's incoming Mailbox memory is pending for process. The field is cleared after the VF driver writes msg_rcv to its command register.</p> <p>For PF: When asserted, the messages in the incoming Mailbox memory are pending for process. The field is cleared only when the event queue is empty.</p>

Function Command Register (0x5004)

Table: Function Command Register (0x5004)

Bit Index	Default	Access Type	Field	Description
[31:3]	0	NA	Reserved	Reserved
[2]	0	RO	Reserved	Reserved

Bit Index	Default	Access Type	Field	Description
[1]	0	RW	msg_rcv	<p>For VF: VF marks the message in its Incoming Mailbox Memory as received. The hardware asserts the acknowledgement bit of the associated PF.</p> <p>For PF: PF marks the message send by target_fn as received. The hardware refreshes the i_msg_status of the PF, and clears the o_msg_status of the target_fn.</p>
[0]	0	RW	msg_send	<p>For VF: VF marks the current message in its own Outgoing Mailbox as valid.</p> <p>For PF:</p> <p>Current target_fn_id belongs to a VF: PF finished writing a message into the Incoming Mailbox memory of the VF with target_fn_id. The hardware sets the i_msg_status field of the target FN's status register.</p> <p>Current target_fn_id belongs to a PF: PF finished writing a message into its own outgoing Mailbox memory. The hardware pushes the message to the event queue of the PF with target_fn_id.</p>

Function Interrupt Vector Register (0x5008)

Table: Function Interrupt Vector Register (0x5008)

Bit Index	Default	Access Type	Field	Description
[31:5]	0	NA	Reserved	Reserved
[4:0]	0	RW	int_vect	5-bit interrupt vector assigned by the driver software.

Target Function Register (0x500C)

Table: Target Function Register (0x500C)

Bit Index	Default	Access Type	Field	Description
[31:8]	0	NA	Reserved	Reserved

Bit Index	Default	Access Type	Field	Description
[7:0]	0	RW	target_fn_id	This field is for PF use only. The FN number that the current operation is targeting.

Function Interrupt Control Register (0x5010)

Table: Function Interrupt Control Register (0x5010)

Bit Index	Default	Access Type	Field	Description
[31:1]	0	NA	res	Reserved
[0]	0	RW	int_en	Interrupt enable.

RTL Version Register (0x5014)

Table: RTL Version Register (0x5014)

Bit	Default	Access Type	Field	Description
[31:16]	0x1fd3	RO		QDMA ID
[15:0]	0	RO		Vivado versions 0x0100: QDMA 3.0 Vivado version 2019.1 0x0201: QDMA 3.1 Vivado version 2019.2 patch 0x0010 : QDMA 4.0 Vivado version 2020.1

Incoming Message Memory (0x5800-0x587C)

Table: Incoming Message Memory (0x5800-0x587C)

Register	Addr	Default	Access Type	Field	Width	Description
i_msg_i	0x5800 + i*4	0	RW		32	The <i>i</i> th word of the incoming message (<i>i</i> < 128).

Outgoing Message Memory (0x5C00-0x5C7C)

Table: Outgoing Message Memory (0x5C00-0x5C7C)

Register	Addr	Default	Access Type	Field	Width	Description

Register	Addr	Default	Access Type	Field	Width	Description
o_msg_i	0x5C00 + i *4	0	RW		32	The <i>i</i> th word of the outgoing message (<i>i</i> < 128).

AXI4-Lite Slave CSR Register Space

The Bridge register space and DMA register space are accessible through the AXI4-Lite Slave CSR interface. This interface is accessible only when `csr_prog_done` port is set to 1. You must wait until `csr_prog_done` port is set.

Table: AXI4-Lite Slave CSR Register Space

Register Space	AXI4-Lite Slave CSR Interface	Details
Bridge registers	AXI4-Lite Slave CSR Address <i>bit [15]</i> is set to 0	Found in qdma_v5_1_bridge_registers.csv available in the Register Reference File .
DMA registers	AXI4-Lite Slave CSR Address <i>bit [15]</i> is set to 1	Described in QDMA PF Address Register Space and QDMA VF Address Register Space . ☞ Note: Through this interface, only the DMA CSR register can be accessed. The DMA Queue space register can only be accessed through AXI4-Lite Slave.

Bridge Register Space

Bridge register addresses start at 0xE00.

QDMA Bridge register descriptions are found in `qdma_v5_1_bridge_registers.csv` available in the [Register Reference File](#).

DMA Register Space

The DMA register space is described in the following sections:

- [QDMA PF Address Register Space](#)
- [QDMA VF Address Register Space](#)

AXI4-Lite Slave Register Space

DMA queue space registers can be accessed through the AXI4-Lite Slave interface.

QDMA Queue space PF register addresses and QDMA Queue space VF register addresses are described in [QDMA_TRQ_SEL_QUEUE_PF \(0x18000\)](#) and [QDMA_TRQ_SEL_QUEUE_VF \(0x3000\)](#).

 **Note:** Through this interface, only the DMA Queue space registers can be accessed. DMA CSR register can be accessed only through AXI4-Lite Slave CSR interface.

In Bridge mode (Bit[28] = 0) address 0x00 to 0xE00 are directed to the PCIe Core configuration register space (ECAM).

Bridge Register space

In Bridge mode (Bit[28] = 0) address 0x00 to 0xE00 are directed to the PCIe core configuration register space (ECAM).

DMA Register space

DMA queue space registers can be accessed through the AXI4-Lite Slave interface.

QDMA Queue space PF register addresses and QDMA Queue space VF register addresses are described in [QDMA_TRQ_SEL_QUEUE_PF \(0x18000\)](#) and [QDMA_TRQ_SEL_QUEUE_VF \(0x3000\)](#).

 **Note:** Through this interface, only the DMA Queue space registers can be accessed. DMA CSR register can be accessed only through AXI4-Lite Slave CSR interface.

Designing with the Subsystem

General Design Guidelines

Use the Example Design

Each instance of the subsystem created by the AMD Vivado™ design tool is delivered with an example design that can be implemented in a device and then simulated. This design can be used as a starting point for your own design or can be used to sanity-check your application in the event of difficulty. See the Example Design content for information about using and customizing the example designs for the subsystem.

Registering Signals

To simplify timing and increase system performance in a programmable device design, keep all inputs and outputs registered between the user application and the subsystem. This means that all inputs and outputs from the user application need to come from, or connect to, a flip-flop. While registering signals might not be possible for all paths, it simplifies timing analysis and makes it easier for the AMD tools to place and route the design.

Recognize Timing Critical Signals


The constraints provided with the example design identify the critical signals and timing constraints that need to be applied.

Make Only Allowed Modifications

You must not modify the subsystem. Any modifications can have adverse effects on system timing and protocol compliance. Supported user configurations of the subsystem can only be made by selecting the options in the customization IP dialog box when the subsystem is generated.

Clocking

Figure: Clocking



PCIe® clocks (pipe_clk, core_clk, user_clk, and mcap_clk) are all driven by bufg_gt sourced from txoutclk pin. These clocks are derived clock from gt refclk0 through a CPLL. In an application where QPLL is used, QPLL is only provided to the GT PCS/ PMA block while txoutclk continues to be derived from a CPLL. All user interface signals of the IP are timed with respect to the same clock (user_clk) which can have a frequency of 62.5, 125, or 250 MHz depending on the link speed and width configured. The QDMA Subsystem for PCIe and the user logic primarily work on user_clk.

Tandem Configuration

Tandem Configuration is available for the QDMA Subsystem for PCI Express for specific AMD UltraScale+™ devices. The two-stage methodology to quickly configure the IP to meet the configuration time requirements in the PCI Express Specification is also available for other AMD IP, including the UltraScale+ Devices Integrated Block for PCI Express and the DMA/Bridge Subsystem for PCI Express. For more information on the Tandem Configuration solution, see *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* (PG213).

The QDMA Subsystem for PCI Express also allows you to enable the MCAP for delivery of partial bitstreams for the Dynamic Function eXchange (DFX) solution. This feature uses a standard initial configuration of the device (NOT Tandem Configuration), followed by delivery of partial bitstreams over the PCIe link to dynamically update portions of the programmable logic. The Vivado DFX design flow is used to create and manage these partial bitstreams; the PCIe end point remains in the static portion of the design.

Configuring the Subsystem for Tandem Configuration or Dynamic Function eXchange

To enable Tandem PROM, Tandem PCIe, or DFX over PCIe capabilities for UltraScale+ devices, select the appropriate IP catalog option when customizing the core. In the Basic tab:

1. Change the Mode to Advanced.
2. Change the Tandem Configuration or Dynamic Function eXchange option according to your particular case:

Tandem Configuration


Load the single two-stage bitstream from flash, with the goal of meeting the 100 ms PCIe enumeration requirement.

Tandem PCIe

Load the first stage bitstream from flash, with the goal of meeting the 100 ms PCIe enumeration requirement. Then load the second stage bitstream over the PCIe link to complete configuration.

DFX over PCIe

This is a standard configuration followed by DFX, using the PCIe / MCAP as the delivery path of partial bitstreams.



The QDMA Subsystem for PCIe Express does not support the pre-defined Tandem PCIe with field updates solution. If your design requires both Tandem Configuration to meet the 100 ms enumeration goal and DFX to enable dynamic reconfiguration of modules during device operation, select Tandem PCIe during IP customization and follow the DFX design flow. This allows you to build your own dynamic design structure using one or more Reconfigurable Partitions.

Supported Devices

Tandem Configuration is available as a production solution for many AMD UltraScale+™ devices. All AMD Virtex™ UltraScale+™ devices along with the two largest AMD Kintex™ UltraScale+™ and single largest AMD Zynq™ UltraScale+™ MPSoCs support Tandem PROM and Tandem PCIe options, as well as the DFX over PCIe feature. AMD Zynq™ UltraScale+™ RFSoC and the KU19P do not have PCIe block locations with the MCAP resource so support for these devices is limited to Tandem PROM only; at this point device support is limited to the three devices listed in the table. All remaining devices, including all AMD Artix™ UltraScale+™, do not offer Tandem or DFX over PCIe capability. Tandem with Field updates is not supported for the QDMA IP for any device.

 **Note:** The Tandem Configuration solution requires a -2 or faster speed grade to meet timing goals for the QDMA IP.

Tandem Configuration supports the configurations described in the following table. Tandem Configuration for PCIe enabled AMD Spartan™ UltraScale+™ devices is not supported at this point.

Table: Tandem PCIe Supported Configurations

Device	Description			
HDL	Verilog Only			
PCIe Configuration	All configurations (max: X16Gen3 or X8Gen4)			
AMD Reference Board Support	KCU116 Evaluation Board for AMD Kintex™ UltraScale+™ FPGA VCU118 Evaluation Board for AMD Virtex™ UltraScale+™ FPGA			
Device Support	Part ¹	PCIe Block Location	Tandem Configuration ²	DFX over PCIe
AMD Artix™ UltraScale+™	AU7P	PCIE40E4_X0Y0	Not supported	Not supported
	AU10P	PCIE40E4_X0Y0	Not supported	Not supported
	AU15P	PCIE40E4_X0Y0	Not supported	Not supported
	AU20P	PCIE40E4_X0Y0	Not supported	Not supported
	AU25P	PCIE40E4_X0Y0	Not supported	Not supported

Device	Description			
AMD Kintex™ UltraScale+™	KU11P	PCIE40E4_X1Y0	Supported	Supported
	KU15P	PCIE40E4_X1Y0	Supported	Supported
	KU19P ³	PCIE4CE4_X0Y0	Tandem PROM only	Not supported
AMD Virtex™ UltraScale+™	VU3P	PCIE40E4_X1Y0	Supported	Supported
	VU5P	PCIE40E4_X1Y0	Supported	Supported
	VU7P	PCIE40E4_X1Y0	Supported	Supported
	VU9P	PCIE40E4_X1Y2	Supported	Supported
	VU11P	PCIE40E4_X0Y0	Supported	Supported
	VU13P	PCIE40E4_X0Y1	Supported	Supported
	VU19P	PCIE4CE4_X0Y2	Supported	Supported
	VU23P	PCIE4CE4_X0Y0	Supported	Supported
	VU27P	PCIE40E4_X0Y0	Supported	Supported
	VU29P	PCIE40E4_X0Y0	Supported	Supported
	VU31P	PCIE4CE4_X1Y0	Supported	Supported
	VU33P	PCIE4CE4_X1Y0	Supported	Supported
	VU35P	PCIE4CE4_X1Y0	Supported	Supported
	VU37P	PCIE4CE4_X1Y0	Supported	Supported
	VU45P	PCIE4CE4_X1Y0	Supported	Supported
	VU47P	PCIE4CE4_X1Y0	Supported	Supported
	VU57P	PCIE4CE4_X1Y0	Supported	Supported
AMD Zynq™ UltraScale+™ MPSoC	ZU3T	PCIE40E4_X0Y0	Not supported	Not supported
	ZU4CG/EG/EV	PCIE40E4_X0Y1	Not supported	Not supported
	ZU5CG/EG/EV	PCIE40E4_X0Y1	Not supported	Not supported
	ZU7CG/EG/EV	PCIE40E4_X0Y1	Not supported	Not supported
	ZU11EG	PCIE40E4_X1Y0	Not supported	Not supported

Device	Description			
	ZU17EG	PCIE40E4_X1Y0	Not supported	Not supported
	ZU19EG	PCIE40E4_X1Y0	Supported	Supported
AMD Zynq™ UltraScale+™ RFSoC 3	ZU21DR	PCIE40E4_X0Y0	Tandem PROM only	Not supported
	ZU25DR	PCIE40E4_X0Y0	Tandem PROM only	Not supported
	ZU27DR	PCIE40E4_X0Y0	Tandem PROM only	Not supported
	ZU28DR	PCIE40E4_X0Y0	Tandem PROM only	Not supported
	ZU29DR	PCIE40E4_X0Y0	Tandem PROM only	Not supported
	ZU39DR	PCIE40E4_X0Y0	Tandem PROM only	Not supported
	ZU43DR	PCIE4CE4_X0Y0	Tandem PROM only	Not supported
	ZU46DR	PCIE4CE4_X0Y0	Tandem PROM only	Not supported
	ZU47DR	PCIE4CE4_X0Y0	Tandem PROM only	Not supported
	ZU48DR	PCIE4CE4_X0Y0	Tandem PROM only	Not supported
	ZU49DR	PCIE4CE4_X0Y0	Tandem PROM only	Not supported

Device	Description
	<ol style="list-style-type: none">1. Only production silicon is officially supported. Bitstream generation is disabled for all engineering sample silicon (ES1, ES2) devices.2. Supported = Tandem PROM and Tandem PCIe both available unless otherwise noted.3. All AMD Zynq RFSoC devices and AMD Kintex UltraScale+ KU19P do not have MCAP-enabled PCIe block locations. Explicit support is added to accommodate Tandem PROM configuration only for these devices. <p> Note: When customizing the PCIe IP for Zynq UltraScale+ RFSoCs, it is important to select GTY locations that align with the PCIe X0Y0 block site. This ensures that the design is routable. For more information, see Answer Record 000037787.</p>

For more information about Tandem Configuration available in other IP and devices, see *UltraScale+ Devices Integrated Block for PCI Express LogicORE IP Product Guide* ([PG213](#)) and *DMA/Bridge Subsystem for PCI Express Product Guide* ([PG195](#)). For more information about Dynamic Function eXchange, see the *Vivado Design Suite User Guide: Dynamic Function eXchange* ([UG909](#)).

Design Flow Steps

This section describes customizing and generating the subsystem, constraining the subsystem, and the simulation, synthesis, and implementation steps that are specific to this IP subsystem. More detailed information about the standard AMD Vivado™ design flows and the IP integrator can be found in the following Vivado Design Suite user guides:

- *Vivado Design Suite User Guide: Designing IP Subsystems using IP Integrator* ([UG994](#))
- *Vivado Design Suite User Guide: Designing with IP* ([UG896](#))
- *Vivado Design Suite User Guide: Getting Started* ([UG910](#))
- *Vivado Design Suite User Guide: Logic Simulation* ([UG900](#))

Customizing and Generating the Subsystem

This section includes information about using AMD tools to customize and generate the subsystem in the AMD Vivado™ Design Suite.

If you are customizing and generating the subsystem in the Vivado IP integrator, see the *Vivado Design Suite User Guide: Designing IP Subsystems using IP Integrator* ([UG994](#)) for detailed information. IP integrator might auto-compute certain configuration values when validating or generating the design. To check whether the values do change, see the description of the parameter in this chapter. To view the parameter value, run the `validate_bd_design` command in the Tcl console.

You can customize the IP for use in your design by specifying values for the various parameters

associated with the IP subsystem using the following steps:

1. Select the IP from the IP catalog.
2. Double-click the selected IP or select the Customize IP command from the toolbar or right-click menu.


For details, see the *Vivado Design Suite User Guide: Designing with IP* ([UG896](#)) and the *Vivado Design Suite User Guide: Getting Started* ([UG910](#)).

Figures in this chapter are illustrations of the Vivado IDE. The layout depicted here might vary from the current version.

Basic Tab

The Basic tab is shown in the following figure.

Figure: Basic Tab



Functional Mode

Option to select between QDMA and AXI Bridge.

Mode

Allows you to select the Basic or Advanced mode of the configuration of core.

Device /Port Type

Only PCI Express® Endpoint device mode is supported.

GT Selection/Enable GT Quad Selection

Select the Quad in which lane 0 is located.

PCIe Block Location

Selects from the available integrated blocks to enable generation of location-specific constraint files and pinouts. This selection is used in the default example design scripts. This option is not available if an AMD Development Board is selected.

Lane Width

The core requires the selection of the initial lane width. The *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG213](#)) defines the available widths and associated generated core. Wider lane width cores can train down to smaller lane widths if attached to a smaller lane-width device. Options are 4, 8, or 16 lanes.

Maximum Link Speed

The core allows you to select the Maximum Link Speed supported by the device. The *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG213](#)) defines the lane widths and link speeds supported by the device. Higher link speed cores are capable of training to a lower link speed if connected to a lower link speed capable device. The default option is Gen3.

Reference Clock Frequency

The default is 100 MHz.

Reset Source

You can choose one of:

PCIe User Reset

The user reset comes from PCIe core after the link is established. When the PCIe link goes down, the user reset is asserted and the core goes to reset mode. And when the link comes back up, the user reset is deasserted.

Phy Ready

When selected, the core is not affected by PCIE link status. This option is useful when a PCIe controller is needed out of reset, regardless of the PCIe link status. For instance, QDMA in root port AXI bridge mode is an example.

AXI Data Width

Select 128, 256 bit, or 512 bit (only for UltraScale+). The core allows you to select the Interface Width, as defined in the *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG213](#)). The default interface width set in the Customize IP dialog box is the lowest possible interface width.

AXI Clock Frequency

250 MHz depending on the lane width/speed.

DMA Interface Option

You can select one of these options:

- AXI Memory Mapped and AXI Stream with Completion
- AXI Memory Mapped only
- AXI Stream with Completion

Number of Queues (up to 2048)

Selects maximum number of queues. Options are 512 (default), 1024 and 2048.

Enable Bridge Slave Mode

Select to enable the AXI-MM Slave interface.

VDM Enable

Select to enable Vendor Define Messages.

AXI Lite Slave Interface

Select to enable the AXI4-Lite slave interface, which can access DMA queue space.

AXI Lite CSR Slave Interface

Select to enable the AXI4-Lite CSR slave interface, which can access DMA Configuration Space Register or Bridge registers.

Enable PIPE Simulation

When selected, this option enables an external third-party bus functional model (BFM) to connect to the PIPE interface of integrated block for PCIe. For details, see *PIPE Mode Simulation Using Integrated Endpoint PCI Express Block in Gen3 x8 and Gen2 x8 Configurations* ([XAPP1184](#)). Refer to these designs to connect the External PIPE Interface ports of the AMD UltraScale™ device core to third-party BFMs.

Enable pipe simulation for faster simulation. This is used only for simulation.

Tandem Configuration or Dynamic Function eXchange

You can select the Dynamic Function eXchange (DFX) over PCIe, which uses the MCAP interface. Tandem Configuration modes are not supported for the QDMA Subsystem for PCIe.

Enable GT Channel DRP Ports

Enable GT-specific DRP ports.

Enable PCIe DRP Ports

Enable PCIe-specific DRP ports.

Additional Transceiver Control and Status Ports

Select to enable any additional ports.

Capabilities Tab

The Capabilities Tab is shown in the following figure.

Figure: Capabilities Tab

The screenshot shows the 'Capabilities' tab selected in the top navigation bar. The interface includes tabs for Basic, Capabilities, PCIe : BARs, PCIe : MISC, PCIe : DMA, Debug Options, Shared Logic, and GT Settings.

SRIOV Capabilities

- SRIOV Capability
- Enable FLR
- Enable Mailbox among functions

Physical Functions

Total Physical Functions: 4

PF#	Vendor ID	Device ID	Revision ID	Subsystem Vendor ID	Subsystem ID
PF0	10EE	903F	00	10EE	0007
PF1	10EE	913F	00	10EE	0007
PF2	10EE	923F	00	10EE	0007
PF3	10EE	933F	00	10EE	0007

PF - ID Initial Values

PF#	Use Classcode Lookup Assistant	Base Class Menu	Base Class Value	Subclass Interface Menu	Subclass Value	Interface Value	Class Code
PF0	<input type="checkbox"/>	Memory controller	▼ 05	Other memory cont...	▼ 80	00	058000
PF1	<input type="checkbox"/>	Memory controller	▼ 05	Other memory cont...	▼ 80	00	058000
PF2	<input type="checkbox"/>	Memory controller	▼ 05	Other memory cont...	▼ 80	00	058000
PF3	<input type="checkbox"/>	Memory controller	▼ 05	Other memory cont...	▼ 80	00	058000

Class Code

PF#	Use Classcode Lookup Assistant	Base Class Menu	Base Class Value	Subclass Interface Menu	Subclass Value	Interface Value	Class Code
PF0	<input type="checkbox"/>	Memory controller	▼ 05	Other memory cont...	▼ 80	00	058000
PF1	<input type="checkbox"/>	Memory controller	▼ 05	Other memory cont...	▼ 80	00	058000
PF2	<input type="checkbox"/>	Memory controller	▼ 05	Other memory cont...	▼ 80	00	058000
PF3	<input type="checkbox"/>	Memory controller	▼ 05	Other memory cont...	▼ 80	00	058000

SRIOV Capability

Enables Single Root Port I/O Virtualization (SR-IOV) capabilities. The integrated block implements extended SR-IOV PCIe. When this is enabled, SR-IOV is implemented on all selected physical functions. When SR-IOV capabilities are enabled only MSI-X interrupt is supported.

Enable Mailbox among functions

This is a Mailbox system to communicate between different functions. When SR-IOV capability (above) is enabled, this option is enabled by default. Mailbox can be selected independently of the SR-IOV Capability selection.

Enable FLR

Enables the functional level reset port. When SR-IOV capability (above) is enabled, this option is enabled by default.

Physical Functions

A maximum of four Physical Functions can be enabled.

PF - ID Initial Values

Vendor ID

Identifies the manufacturer of the device or application. Valid identifiers are assigned by the PCI Special Interest Group to guarantee that each identifier is unique. The default value, 10EEh is the Vendor ID for AMD. Enter a vendor identification number here. FFFFh is reserved.

Device ID

A unique identifier for the application; the default value, which depends on the configuration selected, is 70h. This field can be any value; change this value for the application.

The Device ID parameter is evaluated based on:

- The device family: 9 for AMD UltraScale+™ , 8 for AMD UltraScale™ , and 7 for 7 series devices.
- EP or RP mode
- Link width
- Link speed

For example, Device ID 903F represents Device ID for AMD UltraScale+™ B, 3 for Gen3, and F for X16 (width).

If any of the above values are changed, the Device ID value will be re-evaluated, replacing the previous set value.

 **Recommended:** It is always recommended that the link width, speed and Device Port type be changed first and then the Device ID value. Make sure the Device ID value is set correctly before generating the IP.

Revision ID

Indicates the revision of the device or application; an extension of the Device ID. The default value is 00h; enter values appropriate for the application.

Subsystem Vendor ID

Further qualifies the manufacturer of the device or application. Enter a Subsystem Vendor ID here; the default value is 10EEh. Typically, this value is the same as Vendor ID. Setting the value to 0000h can cause compliance testing issues.

Subsystem ID

Further qualifies the manufacturer of the device or application. This value is typically the same as the Device ID; the default value depends on the lane width and link speed selected. Setting the value to 0000h can cause compliance testing issues.

Class Code

The Class Code identifies the general function of a device.

Use Classcode Lookup Assistant

If selected, the Class Code Look-up Assistant provides the Base Class, Sub-Class and Interface values for a selected general function of a device. This Look-up Assistant tool only displays the three values for a selected function. You must enter the values in Class Code for these values to be translated into device settings.

Base Class

Broadly identifies the type of function performed by the device.

Subclass

More specifically identifies the device function.

Interface

Defines a specific register-level programming interface, if any, allowing device-independent software to interface with the device.

PCIe BARs Tab

The PCIe BARs tab is shown in the following figure.

Figure: PCIe BARs Tab

Board Basic Capabilities PCIe : BARs **PCIe : MISC** PCIe : DMA Debug and Additional Options

Base Address Registers (BARs) serve two purposes. Initially, they serve as a mechanism for the device to request blocks of address space in the system memory map. After the BIOS or OS determines what addresses to assign to the device, the Base Address Registers are programmed with addresses and the device uses this information to perform address decoding.

PF0								
Bar	Type	64 bit	Prefetchable	Size	Scale	Value (Hex)	PCIe to AXI Translation	⋮
<input checked="" type="checkbox"/>	DMA	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	256	Kilobytes	FFFFFFFFFFFC000C	0x0000000000000000	⊕
<input checked="" type="checkbox"/>	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	128	Megabytes	00000000	0x0000000000000000	⊕
<input checked="" type="checkbox"/>	AXI Lite Master	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	4	Kilobytes	FFFFFFFFFFFFF00C	0x0000000000000000	⊖
<input type="checkbox"/>	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	128	Kilobytes	00000000	0x0000000000000000	⊕
<input type="checkbox"/>	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	128	Kilobytes	00000000	0x0000000000000000	⊕
<input type="checkbox"/>	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	128	Kilobytes	00000000	0x0000000000000000	⊕
<input type="checkbox"/>	Expansion ROM	<input type="checkbox"/>	<input type="checkbox"/>	4	Kilobytes	00000000	0x0000000000000000	⊕

Copy PF0

PF1								
Bar	Type	64 bit	Prefetchable	Size	Scale	Value (Hex)	PCIe to AXI Translation	⋮
<input checked="" type="checkbox"/>	DMA	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	256	Kilobytes	FFFFFFFFFFFC000C	0x0000000000000000	⊕
<input type="checkbox"/>	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	128	Megabytes	00000000	0x0000000010000000	⊕
<input checked="" type="checkbox"/>	AXI Lite Master	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	4	Kilobytes	FFFFFFFFFFFFF00C	0x0000000010000000	⊖
<input type="checkbox"/>	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	128	Kilobytes	00000000	0x0000000000000000	⊕
<input type="checkbox"/>	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	128	Kilobytes	00000000	0x0000000000000000	⊕
<input type="checkbox"/>	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	128	Kilobytes	00000000	0x0000000000000000	⊕
<input type="checkbox"/>	Expansion ROM	<input type="checkbox"/>	<input type="checkbox"/>	4	Kilobytes	00000000	0x0000000000000000	⊕

PF2								
-----	--	--	--	--	--	--	--	--

PF3								
-----	--	--	--	--	--	--	--	--

Base Address Register Overview

In Endpoint configuration, the core supports up to six 32-bit BARs or three 64-bit BARs, and the Expansion read-only memory (ROM) BAR. BARs can be one of two sizes:

32-bit BARs

The address space can be as small as 128 bytes or as large as 2 gigabytes. Used for DMA, AXI Lite Master or AXI Bridge Master.

64-bit BARs

The address space can be as small as 128 bytes or as large as 8 Exabytes. Used for DMA, AXI Lite Master or AXI Bridge Master.

All BAR register share these options.

!! Important: The DMA requires a large amount of space to support functions and queues. By default, 64-bit BAR space is selected for the DMA BAR. This applies for PF and VF bars. You must calculate your design needs first before selecting between 64-bit and 32-bit BAR space.

BAR selections are configurable. By default DMA is at BAR 0 (64 bit), AXI-Lite Master is at BAR 2 (64 bit). These selections can be changed according to user needs.

BAR

Click the checkbox to enable the BAR. Deselect the checkbox to disable the BAR.

Type

Select from DMA (by default in BAR0), AXI Lite Master (by default in BAR1, if enabled), or AXI Bridge Master (by default in BAR2, if enabled). All other BARs, you can select between AXI List Master and AXI Bridge Master. Expansion ROM can be enabled by selecting BAR6 For 64-bit BAR (default selection), **DMA** (by default in BAR0), **AXI Lite Master** (by default in BAR2, if enabled), and **AXI Bridge Master** (by default in BAR4, if enabled). Expansion ROM can be enabled by selection BAR6.

DMA

DMA by default is assigned to BAR0 space and for all PFs. DMA option can be selected in any available BAR (only one BAR can have DMA option). If you select DMA Mailbox Management rather than DMA; however, DMA Mailbox Management will not allow you to perform any DMA operations. After selecting the DMA Mailbox Management option, the host has access to the extended Mailbox space. For details about this space, see the QDMA_PF_MAILBOX (0x22400) register space.

AXI Lite Master

Select the AXI Lite Master interface option for any BAR space. The Size, scale and address translation are configurable.

AXI Bridge Master

Select the AXI Bridge Master interface option for any BAR space. The Size, scale and address translation are configurable.

Expansion ROM

When enabled, this space is accessible on the AXI4-Lite Master. This is a read-only space. The size, scale, and address translation are configurable.

Size

The available Size range depends on the 32-bit or 64-bit bar selected. The DMA requires 256 Kbytes of space, which is the fixed default selection. Other BAR size selections are available, but must be specified.

Scale

Select between Byte, Kilobytes and Megabytes.

Value

The value assigned to the BAR based on the current selections.

☞ **Note:** For best results, disable unused base address registers to conserve system resources. A base address register is disabled by deselecting unused BARs in the Customize IP dialog box.

Related Information

[QDMA_PF_MAILBOX \(0x22400\)](#)


SRIOV Config Tab

The SRIOV Config tab allows you to specify the SR-IOV capability for a physical function (PF). The information is used to construct the SR-IOV capability structure. Virtual functions do not exist on power-on. It is the function of the system software to discover and enable VFs based on system capability. The VF support is discovered by scanning the SR-IOV capability structure for each PF.

☞ **Note:** When SRIOV Capability is selected in [Capabilities Tab](#), the SRIOV Config tab appears.

The SRIOV Config Tab is shown in the following figure.

Figure: SRIOV Config Tab



General SRIOV Config

This value specifies the offset of the first PF with at least one enabled VF. When ARI is enabled, allowed value is 'd4 or 'd64, and the total number of VF in all PFs plus this field must not be greater than 256. When ARI is disabled, this field will be set to 1 to support 1PFplus 7VF non-ARI SRIOV configurations only.

Number of PFx VFs

Indicates the number of virtual functions associated to the physical function. A total of 252 virtual functions are available that can be flexibly used across the four physical functions.

VF Device ID

Indicates the 16-bit Device ID for all virtual functions associated with the physical function.

SRIOV VF BARs Tab

The SRIOV VF BARs tab is shown in the following figure.

Figure: SRIOV VF BARs Tab

SRIOV VF BARs Configuration								
Physical Function		Type	64 bit	Prefetchable	Size	Scale	Value (Hex)	PCIe to AXI Translation
PF0	Bar 1	DMA	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	32	Kilobytes	FFFFFFFFF800C	0x0000000000000000
	Bar 2	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	4	Kilobytes	00000000	0x0000000040000000
	Bar 3	AXI Lite Master	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	4	Kilobytes	FFFFFFFFF00C	0x0000000040000000
	Bar 4	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	4	Kilobytes	00000000	0x0000000000000000
	Bar 5	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	4	Kilobytes	00000000	0x0000000000000000
	Bar 6	AXI Bridge Master	<input type="checkbox"/>	<input type="checkbox"/>	4	Kilobytes	00000000	0x0000000000000000

Copy PF0

PF1

PF2

PF3

The SRIOV VF BARs tab enables you to configure the base address registers (BARs) for all virtual function (VFs) within a virtual function group (VFG). All the VFs within the same VFG share the same BASE ADDRESS Registers (BARs) configurations. Each Virtual Function supports up to six 32-bit BARs or three 64-bit BARs. Virtual Function BARs can be configured without any dependency on the settings of the associated Physical Functions BARs.

!! Important: The DMA requires a large amount of space to support functions and queues. By default, 64-bit BAR space is selected for the DMA BAR. This applies for PF and VF bars. You must calculate your design needs first before selecting between 64-bit and 32-bit BAR space.

BAR selections are configurable. By default DMA is at BAR 0 (64 bit), AXI4-Lite Master is at BAR 2 (64-bit). These selections can be changed according to user needs.

BAR

Select applicable BARs using the checkboxes.

Type

Select the relevant option:

DMA

Is fixed to BAR0 space.

AXI Lite Master

Is fixed to BAR1 space.

AXI Bridge Master

Is fixed to BAR2 space. For all other bars, select either AXI Lite Master or AXI Bridge Master.

 **Note:** The current IP supports a maximum of one DMA BAR (or a management BAR given only mailbox is required) for one VF. The other BARs can be configured as AXI Lite Master to access the assigned memory space through the AXI4-Lite bus. Virtual Function BARs do not support I/O space and must be configured to map to the appropriate memory space.

64-bit

VF BARs can be either 64-bit or 32-bit. The default is 64-bit BAR.

- 64-bit addressing is supported for the DMA BAR.
- When a BAR is set as 64 bits, it uses the next BAR for the extended address space and makes the next BAR inaccessible.

Size

The available Size range depends on the 32-bit or 64-bit BAR selected.

The Supported Page Sizes field indicates all the page sizes supported by the PF and, as required by the SR-IOV specification. Based on the Supported Page Size field, the system software sets the System Page Size field which is used to map the VF BAR memory addresses. Each VF BAR address is aligned to the system page boundary.

By default, DMA space is 32 Kbytes. With this much space allocated, the user logic can access 256 queues for a VF function.

Value

The value assigned to the BAR based on the current selections.

PCIe MISC Tab

The PCIe Miscellaneous Tab is shown in the following figure.

Figure: PCIe MISC Tab

The screenshot shows a configuration interface for MSI-X Capabilities. It includes sections for PF0, PF1, PF2, and PF3, each with an 'Enable' checkbox and an 'MSI-X Table Settings' section. The 'Table Size' for all PFs is set to 007. There is also a 'Miscellaneous' section with several checkboxes and a 'Link Status Register' section with a note about clock selection.

Physical Function	Enable MSI-X Capability Structure	Table Size
PF0	<input checked="" type="checkbox"/>	007
PF1	<input checked="" type="checkbox"/>	000..007
PF2	<input checked="" type="checkbox"/>	007
PF3	<input checked="" type="checkbox"/>	007

Miscellaneous

- Extended Tag Field
- Configuration Extended Interface
- Legacy Configuration Extended Interface
- Add the PCIe XVC-VSEC to the Example Design
- Access Control Services (ACS) Enable
- Configuration Management Interface

Link Status Register

Selects whether the device reference clock is provided by the connector (Synchronous) or generated via an onboard PLL(Asynchronous)

- Enable Slot Clock Configuration

MSI-X Capabilities

MSI-X is enabled by default.

The MSI-X settings for different physical functions can be set as required.

MSI-X Table Settings

Defines the MSI-X Table Structure.

Table Size

Specifies the MSI-X Table size. The default is 8 (8 interrupt vectors per function).

Extended Tag Field

By default for AMD UltraScale+™ devices the Extended Tab option gives 256 tags. If Extended Tag option is not selected, the DMA uses 32 tags.

Configuration Extended Interface

The PCIe extended interface can be selected for more configuration space. When Configuration Extended Interface is selected user is responsible for adding logic to extend the interface to make it work properly.

Access Control Server (ACS) Enable

ACS is selected by default.

Configuration Management Interface

This interface is used to read and write to the configuration space registers.

Link Status Register

By default, Enable Slot Clock Configuration is selected. This means that the slot configuration bit is enabled in the link status register.

AXI BARs Tab

The SRIOV VF BARs tab is shown in the following figure.

Figure: AXI BARs Tab



Slave Bridge Address Translation

No Address Translation

When this option is selected, the DMA does not do any address translation. One full 64-Bit BAR space is provided, and you are responsible for any address translation, if required.

When address translation is required by DMA, do not select this option.

AXI Bar_0 Address Translation

Aperture Base Address and *Aperture High Address* can be programmed with desired values. This provides the AXI bar size.

Address translation for higher bits (above BAR size) can be programmed as desired.

Note: In the IP integrator, you can program the Aperture Base Address and Aperture High Address parameters using the Address Editor feature.

Related Information

[Slave Bridge](#)

PCIe DMA Tab

The PCIe DMA Tab is shown in the following figure.

Figure: PCIe DMA Tab

The screenshot shows a configuration interface for SRIOV VF BARs. At the top, there are tabs: Basic, Capabilities, PCIe : BARs, SRIOV Config, SRIOV VF BARs (which is selected), PCIe : MISC, PCIe : DMA, and Debug and Additional Options.

Descriptor Bypass

Descriptor Bypass for Read/Write (H2C/C2H)

None
 Descriptor bypass and internal

C2H Stream Completion

Color bits	Error bits
Color bit position Reg0: 1	Error bit position Reg0: 2
Color bit position Reg1: 0	Error bit position Reg1: 0
Color bit position Reg2: 0	Error bit position Reg2: 0
Color bit position Reg3: 0	Error bit position Reg3: 0
Color bit position Reg4: 0	Error bit position Reg4: 0
Color bit position Reg5: 0	Error bit position Reg5: 0
Color bit position Reg6: 0	Error bit position Reg6: 0

Performance options

Prefetch cache depth: 16
CMPT Coalesce Max buffer: 16

Data Protection

None
 Data Protection

Descriptor Bypass for Read/Write (H2C/C2H)

Two options to select from.

Note: In this mode (Internal mode) DMA will not bypass any H2C or C2H descriptors.

Descriptor bypass and Internal

In this mode descriptor ports for bypass out and bypass in are both enabled. Based on the context settings H2C or C2H descriptors can be sent out on descriptor bypass out. User can send in descriptors on Descriptor bypass in ports.

C2H Stream Completion

C2H Stream Completion Color bits

Completion Color bit position in completion entry. There are seven registers available to program, from bit 0 to 511 (for 64 bytes completion). You can program the bits, and generate a BIT file. During the DMA transfer, the input pins `s_axis_c2h_cmpt_ctrl_color_idx[2:0]` determine which Color bit position to use. Default bit position 1 is selected in register 0.

C2H Stream Completion Error bits

Completion Error bit position in completion entry. There are seven registers available to program, from bit 0 to 511 (for 64 bytes completion). You can program the bits, and generate a BIT file. During a DMA transfer, the input pins `s_axis_c2h_cmpt_ctrl_err_idx[2:0]` determine which Error bit position to use. Default bit position 2 is selected in register 0.

Performance options

Pre-fetch cache depth

The Prefetch cache supports up to 64 Queues. Select one of 16 or 64 (default 16). The Prefetch cache can support that many active queues at any given time. When one active queue finishes fetch and delivers all the descriptors for the packets of that queue, it then releases cache entry for other active queues. A larger cache size supports more active queues, but the area will also increase.

CMPT Coalesce Max buffer

Completion (CMPT) Coalesce Max buffer supports up to 64 buffers. Select one of 16 or 32 (default 16). Each entry of the CMPT Coalesce Buffer coalesces multiple Completions (up to 64B) to form a single queue before writing to the host to improve bandwidth utilization. A deeper CMPT Coalesce Buffer allows coalescing within more queues, but will increase the area as a downside.

Data Protection

Parity Checking and end to end data protection. By default, data protection is not enabled.

When **Data Protection** is not enabled:

- You do not need to give any CRC/ECC values on C2H data and the control interface.
- This will not log any Error and will not drop any packet.
- User should ground the ECC and CRC ports.
- CMPT parity check is not affected by this parameter.

 **Note:** You must always give the parity on CMPT.

When **Data Protection** is enabled:

- You must send CRC/ECC values on C2H data and the control interface.
- If there is any ECC or CRC error, error bits will be logged and data packet will be sent to host.
- If error interrupt is enabled, an interrupt will be sent to host.
- FATAL error can be enabled in the QDMA_C2H_FATAL_ERR_ENABLE register.
 - QDMA_C2H_FATAL_ERR_ENABLE[0]: If this bit is set, all packets are dropped after an error occurs.
 - QDMA_C2H_FATAL_ERR_ENABLE[1]: If this bit is set, parity is inverted and an error packet is sent to PCIe.

Debug and Additional Options Tab

Following is the Debug and Additional Options tab:.

Figure: Debug and Additional Options Tab

Board Basic Capabilities PCIe : BARs SRIOV Config SRIOV VF BARs PCIe : MISC PCIe : DMA Debug and Additional Options

Debug Options

Enable JTAG Debugger.
 Enable In System IBERT.
 Add Mark Debug Utility
 Enable Descrambler for Gen3 Mode.
 Enable PCIe Debug Ports

SharedLogic Options

GT Wizard Options
Select whether GT Wizard is included in the core itself or in the example design.
 Include GT Wizard in core
 Include GT Wizard in example design

GT COMMON Options
Select whether GT COMMON is included in the core/GT Wizard itself or in the example design.
 Include GT COMMON in example design
 Include GT COMMON inside GT Wizard
 No sharing when inside GT Wizard and PCIe

GT Settings

Form factor driven Insertion loss adjustment: Add-in Card
Link Partner TX Preset: 4
Disable GT Channel LOC Constraint: false

Debug Options

Enable JTAG Debugging

This feature provides ease of debug for the following:

LTSSM state transitions

This shows all the LTSSM state transitions that are made starting from link up.

PHY Reset FSM transitions

This shows the PHY reset FSM (internal state machine that is used by the PCIe® solution IP).

Receiver Detect

This shows all the lanes that completed receiver detect successfully

For more details, see *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG213](#)).

Enable In System IBERT

This debug option is used to check and see the eye diagram of the serial link at the desired link speed. For more information on In System IBERT, refer to *In-System IBERT LogiCORE IP Product Guide* ([PG246](#)).

!! Important: This option is used mainly for hardware debug purposes. Simulations are not supported when this option is used.

Add Mark Debug Utility

Adds predefined PCIe signals to with mark_debug attribute so these signals can be added in ILA for debug purpose.

Enable Descrambler for Gen3 Mode

This debug option integrates encrypted version of the descrambler module inside the PCIe core, which are used to descramble the PIPE data to/from PCIe integrated block in Gen3 link speed mode.

Enable PCIe Debug Ports

Reserved. This feature is not supported in this version.

Shared Logic Options

GT Wizard Options

You can select include GT Wizard in the example design and then the GT Wizard IP is delivered into the example design area. You can reconfigure the IP for further testing purposes. By default, the GT Wizard IP is delivered in the PCIe IP core as a hierarchical IP and you cannot re-customize it. For signal descriptions and for other details, see the *UltraScale Architecture GTY Transceivers User Guide* ([UG578](#)) or *UltraScale Architecture GTH Transceivers User Guide* ([UG576](#)).

GT COMMON Options

This option is used to share the GT COMMON block used in the design when Gen2 (PLL Selection is QPLL1) and Gen3 link speeds are selected.

- When Include GT COMMON in example design is selected, GT common block instance is available in the support wrapper, which is inside the AMD top file and can be used either by the core or the external logic.
- When Include GT COMMON inside GT Wizard is used, GT COMMON can be shared by external logic.
- When No Sharing when inside GT Wizard and PCIe is selected, no sharing of GT COMMON block is allowed.
- When Include GT COMMON in example design and Include GT Wizard in example design are selected together, you must use the latest GT COMMON settings from the example design project of the GT Wizard IP of the same configuration. This specific option delivers static GT COMMON wrappers which have the latest settings.

GT Settings

Form factor driven Insertion loss adjustment

Indicates the transmitter to receiver insertion loss at the Nyquist frequency depending on the form factor selection. Three options are provided:

Chip-to-Chip

The value is 5 dB

Add-in Card

The value is 15 dB and is the default option

Backplane

The value is 20 dB

These insertion loss values are applied to the GT Wizard subcore.

Link Partner TX Preset

It is not recommended that you change the default value of 4. However, a preset value of 5 might work better on some systems.

Disable GT Channel LOC Constraint

Reserved. Not supported in this version.

User Parameters

Additional core customizing options are available. For details, see AR [72352](#).

Output Generation

For details, see the *Vivado Design Suite User Guide: Designing with IP* ([UG896](#)).

Constraining the Subsystem

Required Constraints

The QDMA Subsystem for PCIe subsystem requires the specification of timing and other physical implementation constraints to meet specified performance requirements for PCI Express® . These constraints are provided in a Xilinx Design Constraints (XDC) file. Pinouts and hierarchy names in the generated XDC correspond to the provided example design.

!! Important: If the example design top file is not used, copy the IBUFDS_GTE4 instance for the reference clock, IBUF Instance for sys_rst and also the location and timing constraints associated with them into your local design top.

To achieve consistent implementation results, an XDC containing these original, unmodified constraints must be used when a design is run through the AMD tools. For additional details on the definition and use of an XDC or specific constraints, see *Vivado Design Suite User Guide: Using Constraints* ([UG903](#)).

Constraints provided with the Integrated Block for PCIe® solution are tested in the hardware and provide consistent results. Constraints can be modified, but modifications should only be made with a thorough understanding of the effect of each constraint. Additionally, support is not

provided for designs that deviate from the provided constraints.

Device, Package, and Speed Grade Selections

The device selection portion of the XDC informs the implementation tools which part, package, and speed grade to target for the design.

The device selection section always contains a part selection line, but can also contain part or package-specific options. An example part selection line follows:

```
CONFIG PART = xcvu9p-flgb2104-2-i
```

Clock Frequencies

For detailed information about clock requirements, see the *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG213](#)) .

Clock Management

For detailed information about clock requirements, see the *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG213](#)) .

Clock Placement

For detailed information about clock requirements, see the *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG213](#)) .

Banking

This section is not applicable for this IP subsystem.

Transceiver Placement

This section is not applicable for this IP subsystem.

I/O Standard and Placement

This section is not applicable for this IP subsystem.

Relocating the Integrated Block Core

By default, the IP core-level constraints lock block RAMs, transceivers, and the PCIe block to the recommended location. To relocate these blocks, you must override the constraints for these

blocks in the XDC constraint file. To do so:

1. Copy the constraints for the block that needs to be overwritten from the core-level XDC constraint file.
2. Place the constraints in the user XDC constraint file.
3. Update the constraints with the new location.

The user XDC constraints are usually scoped to the top-level of the design; therefore, ensure that the cells referred by the constraints are still valid after copying and pasting them. Typically, you need to update the module path with the full hierarchy name.

 **Note:** If there are locations that need to be swapped (that is, the new location is currently being occupied by another module), there are two ways to do this:

- If there is a temporary location available, move the first module out of the way to a new temporary location first. Then, move the second module to the location that was occupied by the first module. Next, move the first module to the location of the second module. These steps can be done in XDC constraint file.
- If there is no other location available to be used as a temporary location, use the `reset_property` command from Tcl command window on the first module before relocating the second module to this location. The `reset_property` command cannot be done in the XDC constraint file and must be called from the Tcl command file or typed directly into the Tcl Console.

Simulation

For comprehensive information about AMD Vivado™ simulation components, as well as information about using supported third-party tools, see the *Vivado Design Suite User Guide: Logic Simulation (UG900)*.

Basic Simulation

Simulation models for the AXI-MM and AXI-ST options can be generated and simulated. The simple simulation model options enable you to develop complex designs.

AXI-MM Mode

The example design for the AXI4 Memory Mapped (AXI-MM) mode has 512 KB block RAM on the user side, where data can be written to the block RAM, and read from block RAM to the Host. After the Host to Card (H2C) transfer is started, the DMA reads data from the Host memory, and writes to the block RAM. After the transfer is completed, the DMA updates the write back status and generates an interrupt (if enabled). Then, the Card to Host (C2H) transfer is started, and the DMA reads data from the block RAM and writes to the Host memory. The original data is compared with the C2H write data. H2C and C2H are set up with one descriptor each, and the total transfer size is 128 bytes.

AXI-ST Mode

The example design for the AXI4-Stream (AXI-ST) mode has a data check that checks the data from the H2C transfer, and has a data generator that generates the data for C2H transfer.

After the H2C transfer is started, the DMA engine reads data from the Host memory, and writes to the user side. After the transfer is completed, the DMA updates write back status and generates an interrupt (if enabled). The data checker on the user side checks for a predefined data to be present, and the result is posted in a predefined address for the user application to read.

After the C2H transfer is started, the data generator generates predefined data and associated control signals, and sends them to the DMA. The DMA transfers data to the Host, updates the completion (CMPT) ring entry/status, and generates an interrupt (if enabled).

H2C and C2H are set up with one descriptor each, and the total transfer size is 128 bytes.

Related Information

[Reference Software Driver Flow](#)

PIPE Mode Simulation

The QDMA Subsystem for PCIe supports the PIPE mode simulation where the PIPE interface of the core is connected to the PIPE interface of the link partner. This mode increases the simulation speed.

Use the Enable PIPE Simulation option on the Basic tab of the Customize IP dialog box to enable PIPE mode simulation in the current AMD Vivado™ Design Suite solution example design, in either Endpoint mode or Root Port mode. The External PIPE Interface signals are generated at the core boundary for access to the external device. Enabling this feature also provides the necessary hooks to use third-party PCI Express® VIPs/BFMs instead of the Root Port model provided with the example design.

The tables below describe the PIPE bus signals available at the top level of the core and their corresponding mapping inside the EP core (pcie_top) PIPE signals.

Table: In Commands and Endpoint PIPE Signal Mappings

In Commands	Endpoint PIPE Signals Mapping
common_commands_in[25:0]	not used

Table: Out Commands and Endpoint PIPE Signal Mappings

Out Commands	Endpoint PIPE Signals Mapping
common_commands_out[0]	pipe_clk ¹
common_commands_out[2:1]	pipe_tx_rate_gt ²
common_commands_out[3]	pipe_tx_rcvr_det_gt
common_commands_out[6:4]	pipe_tx_margin_gt

Out Commands	Endpoint PIPE Signals Mapping
common_commands_out[7]	pipe_tx_swing_gt
common_commands_out[8]	pipe_tx_reset_gt
common_commands_out[9]	pipe_tx_deemph_gt
common_commands_out[16:10]	not used ³

1. pipe_clk is an output clock based on the core configuration. For Gen1 rate, pipe_clk is 125 MHz. For Gen2 and Gen3, pipe_clk is 250 MHz.
 2. pipe_tx_rate_gt indicates the pipe rate (2'b00-Gen1, 2'b01-Gen2, and 2'b10-Gen3).
 3. The functionality of this port has been deprecated and it can be left unconnected.

Table: Input Bus With Endpoint PIPE Signal Mapping

Input Bus	Endpoint PIPE Signal Mapping
pipe_rx_0_sigs[31:0]	pipe_rx0_data_gt
pipe_rx_0_sigs[33:32]	pipe_rx0_char_is_k_gt
pipe_rx_0_sigs[34]	pipe_rx0_elec_idle_gt
pipe_rx_0_sigs[35]	pipe_rx0_data_valid_gt
pipe_rx_0_sigs[36]	pipe_rx0_start_block_gt
pipe_rx_0_sigs[38:37]	pipe_rx0_syncheader_gt
pipe_rx_0_sigs[83:39]	not used

Table: Output Bus with Endpoint PIPE Signal Mapping

Output Bus	Endpoint PIPE Signals Mapping
pipe_tx_0_sigs[31: 0]	pipe_tx0_data_gt
pipe_tx_0_sigs[33:32]	pipe_tx0_char_is_k_gt
pipe_tx_0_sigs[34]	pipe_tx0_elec_idle_gt
pipe_tx_0_sigs[35]	pipe_tx0_data_valid_gt
pipe_tx_0_sigs[36]	pipe_tx0_start_block_gt
pipe_tx_0_sigs[38:37]	pipe_tx0_syncheader_gt
pipe_tx_0_sigs[39]	pipe_tx0_polarity_gt

Output Bus	Endpoint PIPE Signals Mapping
pipe_tx_0_sigs[41:40]	pipe_tx0_powerdown_gt
pipe_tx_0_sigs[69:42]	not used ¹

1. The functionality of this port has been deprecated and it can be left unconnected.

Synthesis and Implementation

For details about synthesis and implementation, see the *Vivado Design Suite User Guide: Designing with IP* ([UG896](#)).

Example Design

This chapter contains information about the example designs provided in the AMD Vivado™ Design Suite.

Available Example Designs


The example designs are as follows:

- [AXI Memory Mapped and AXI4-Stream With Completion Default Example Design](#)
- [AXI Memory Mapped Example Design](#)
- [AXI Stream with Completion Example Design](#)
- [AXI Stream Loopback Example Design](#)
- [Example Design with Descriptor Bypass In/Out Loopback](#)
- [AXI Stream Performance Example Design](#)

AXI Memory Mapped and AXI4-Stream With Completion Default Example Design

The following is an example design generated when the DMA Interface Selection option is set to AXI Memory Mapped and AXI4-Stream with Completion option in the Basic tab.

Figure: Default Example Design




The generated example design provides blocks to interface with the AXI Memory Mapped and AXI4-Stream interfaces.

- The AXI MM interface is connected to 512 KB of block RAM.
- The AXI4-Stream interface is connected to custom data generator and data checker module.
- The CMPT interface is connected to the Completion block generator.
- The data generator and checker works only with predefined pattern, which is a 16-bit incremental pattern starting with 0. This data file is included in driver package.

The pattern generator and checker can be controlled using the registers found in the [Example Design Registers](#). These registers can only be controlled through the AXI4-Lite Master interface. To test the QDMA Subsystem for PCIe's AXI4-Stream interface, ensure that the AXI4-Lite Master interface is present.

AXI Memory Mapped Example Design

Figure: AXI Memory Map Example Design




The example design above is generated when the DMA Interface Selection option is set to AXI-MM only in the Basic tab. In this mode, the AXI MM interface is connected to a 512 KB block RAM. The diagram above shows that AXI4-Lite Master is connected to a 4 KB block RAM. For Host to Card (H2C) transfers, the DMA reads data from the Host and writes to the block RAM. For Card to Host (C2H) transfers, the DMA reads data from the block RAM and writes to the Host.

memory.

AXI Stream with Completion Example Design

Figure: AXI4-Stream Example Design




The example design above is generated when the DMA Interface Selection option is set to AXI Stream with Completion in the Basic tab. In this mode, the AXI-ST H2C interface is connected to a data checker, and the AXI-ST C2H interface is connected to data generator and CMPT interface is connected to Completion generator module. The diagram shows AXI4-Lite Master is connected to the 4 KB block RAM and the User Control logic. The software can control data checker and data generator through the AXI4-Lite Master interface. The data generator and checker work only with a predefined pattern, which is a 16-bit incremental pattern starting with 0. This data file is included in the driver package.

The pattern generator and checker can be controlled using the registers found in the [Example Design Registers](#).

AXI Stream Loopback Example Design

Figure: AXI4-Stream Loopback Example Design




The example design above is generated when the DMA Interface Selection option is set to AXI

Stream with Completion in the Basic tab. In this mode, the AXI-ST H2C interface is connected to a data checker, and the AXI-ST C2H interface is connected to data generator and CMPT interface is connected to Completion generator module. But this example design can also be used as a streaming loopback design.

Set the Example design register [C2H_CONTROL_REG \(0x008\)](#) bit[0] to 1 to turn this example design into a streaming loopback design. The example design then takes H2C streaming packets and loops them back to the C2H Streaming interface. Completion packets are generated from the loopback design.

Example Design with Descriptor Bypass In/Out Loopback

Figure: AXI Memory Map and Descriptor Bypass Example Design



The example design above is generated when Descriptor Bypass for Read (H2C) and Descriptor Bypass for Write (C2H) options are selected in the PCIe DMA tab. These options can be selected with any of the DMA Interface Options in the Basic tab:

- AXI Memory Mapped and AXI4-Stream with Completion
- AXI Memory Mapped only
- AXI Stream with Completion
- AXI Memory Mapped with Completion

The Descriptor Bypass in/out loopback is controlled by the AXI4-Lite Master by writing to the Example Design Register [DESCRIPTOR_BYPASS \(0x090\)](#) bit[0] and bit[1].

To enable Descriptor bypass out, proper context programming needs to be done. For details, see [Context Programming](#).

C2H Stream Simple Bypass Mode Transfer

To set up a QDMA to data transfer in simple bypass mode


1. Write the active qid to register 0x1408 (MDMA_C2H_PFCH_BYP_QID).
2. Read the tag value from 0x140C (MDMA_c2H_PFCH_BYP_TAG).

3. Write the tag value and qid that was used to fetch the tag in the example design register C3H_PREFETCH_TAG 0x24. The qid bits are [26:16] and tag bits are [6:0].
4. Set up the simple bypass descriptor loopback by writing register DESCRIPTOR_BYPASS 0x90 bits [2:0] = 3 'b100.

After the setup initial C2H stream data transfer, the prefetch tag is valid until the qid is valid. When the current qid becomes invalid, you must generate a new tag.

AXI Stream Performance Example Design

Figure: AXI4-Stream Example Design



You can get the performance example design by executing the following command in Tcl window after generating QDMA_0 IP.

```
set_property CONFIG.performance_exdes {true} [get_ips qdma_0]
```

Generate output products for QDMA IP. Then generate the example design for QDMA IP.

Example Design Registers

Table: Example Design Registers

Registers	Address	Description
C2H_ST_QID (0x000)	0x000	AXI-ST C2H Queue id
C2H_ST_LEN (0x004)	0x004	AXI-ST C2H transfer length
C2H_CONTROL_REG (0x008)	0x008	AXI-ST C2H pattern generator control
H2C_CONTROL_REG (0x00C)	0x00C	AXI-ST H2C Control
H2C_STATUS (0x010)	0x010	AXI-ST H2C Status

Registers	Address	Description
C2H_STATUS (0x018)	0x018	AXI-ST C2H Status
C2H_PACKET_COUNT (0x020)	0x020	AXI-ST C2H number of packets to transfer
C2H_COMPLETION_DATA_0 (0x030) to C2H_COMPLETION_DATA_7 (0x04C)	0x4C-0x030	AXI-ST C2H completion data
C2H_COMPLETION_SIZE (0x050)	0x050	AXI-ST completion data type
SCRATCH_REG0 (0x060)	0x060	Scratch register 0
SCRATCH_REG1 (0x064)	0x064	Scratch register 1
C2H_PACKETS_DROP (0x088)	0x088	AXI-ST C2H Packets drop count
C2H_PACKETS_ACCEPTED (0x08C)	0x08C	AXI-ST C2H Packets accepted count
DESCRIPTOR_BYPASS (0x090)	0x090	C2H and H2C descriptor bypass loopback
USER_INTERRUPT (0x094)	0x094	User interrupt, vector number, function number
USER_INTERRUPT_MASK (0x098)	0x098	User interrupt mask
USER_INTERRUPT_VECTOR (0x09C)	0x09C	User interrupt vector
DMA_CONTROL (0x0A0)	0x0A0	DMA control
VDM_MESSAGE_READ (0x0A4)	0x0A4	VDM message read

C2H_ST_QID (0x000)

Table: C2H_ST_QID (0x000)

Bit	Default	Access Type	Field	Description
[31:11]	0	NA		Reserved
[10:0]	0	RW	c2h_st_qid	AXI4-Stream C2H Queue ID

C2H_ST_LEN (0x004)

Table: C2H_ST_LEN (0x004)

Bit	Default	Access Type	Field	Description
[31:16]	0	NA		Reserved
[15:0]	0	RW	c2h_st_len	AXI4-Stream packet length

C2H_CONTROL_REG (0x008)

Table: C2H_CONTROL_REG (0x008)

Bit	Default	Access Type	Description
[31:6]	0	NA	Reserved
[5]	0	RW	C2H Stream Marker request C2H Stream Marker response will be registered at address 0x18, bit [0].
[4]	0	NA	reserved
[3]	0	RW	Disable completion. For this packet, there will not be any completion.
[2]	0	RW	Immediate data. When set, the data generator sends immediate data. This is a self-clearing bit. Write 1 to initiate transfer.
[1]	0	RW	Starts AXI-ST C2H transfer. This is a self-clearing bit. Write 1 to initiate transfer.
[0]	0	RW	Streaming loop back. When set, the data packet from H2C streaming port in the Card side is looped back to the C2H streaming ports.

For Normal C2H stream packet transfer, set address offset 0x08 to 0x2.

For C2H immediate data transfer, set address offset 0x8 to 0x4.

For C2H/H2C stream loopback, set address offset 0x8 to 0x1.

H2C_CONTROL_REG (0x00C)

Table: H2C_CONTROL_REG (0x00C)

Bit	Default	Access Type	Description
[31:30]	0	NA	Reserved
[0]	0	RW	Clear match bit for H2C transfer.

H2C_STATUS (0x010)

Table: H2C_STATUS (0x010)

Bit	Default	Access Type	Description
[31:15]	0	NA	Reserved
[14:4]	0	R	H2C transfer Queue ID
[3:1]	0	NA	Reserved
[0]	0	R	H2C transfer match

C2H_STATUS (0x018)

Table: C2H_STATUS (0x018)

Bit	Default	Access Type	Description
[31:30]	0	NA	Reserved
[0]	0	R	C2H Marker response

C2H_PACKET_COUNT (0x020)

Table: C2H_PACKET_COUNT (0x020)

Bit	Default	Access Type	Description
[31:10]	0	NA	Reserved
[9:0]	0	RW	AXI-ST C2H number of packet to transfer

C2H_PREFETCH_TAG(0x024)

Table: C2H_PREFETCH_TAG (0x024)

Bit	Default	Access Type	Description
[31:27]	0	NA	Reserved
[26:16]	0	RW	Qid for prefetch tag
[15:7]	0	NA	Reserved
[6:0]	0	RW	Prefetch tag value

C2H_COMPLETION_DATA_0 (0x030)

Table: C2H_COMPLETION_DATA_0 (0x030)

Bit	Default	Access Type	Description
[31:0]	0	NA	AXI-ST C2H Completion Data [31:0]

C2H_COMPLETION_DATA_1 (0x034)**Table: C2H_COMPLETION_DATA_1 (0x034)**

Bit	Default	Access Type	Description
[31:0]	0	NA	AXI-ST C2H Completion Data [63:32]

C2H_COMPLETION_DATA_2 (0x038)**Table: C2H_COMPLETION_DATA_2 (0x038)**

Bit	Default	Access Type	Description
[31:0]	0	NA	AXI-ST C2H Completion Data [95:64]

C2H_COMPLETION_DATA_3 (0x03C)**Table: C2H_COMPLETION_DATA_3 (0x03C)**

Bit	Default	Access Type	Description
[31:0]	0	NA	AXI-ST C2H Completion Data [127:96]

C2H_COMPLETION_DATA_4 (0x040)**Table: C2H_COMPLETION_DATA_4 (0x040)**

Bit	Default	Access Type	Description
[31:0]	0	NA	AXI-ST C2H Completion Data [159:128]

C2H_COMPLETION_DATA_5 (0x044)**Table: C2H_COMPLETION_DATA_5 (0x044)**

Bit	Default	Access Type	Description
[31:0]	0	NA	AXI-ST C2H Completion Data [191:160]

C2H_COMPLETION_DATA_6 (0x048)

Table: C2H_COMPLETION_DATA_6 (0x048)

Bit	Default	Access Type	Field	Description
[31:0]	0	NA	NA	AXI-ST C2H Completion Data [223:192]

C2H_COMPLETION_DATA_7 (0x04C)**Table: C2H_COMPLETION_DATA_7 (0x04C)**

Bit	Default	Access Type	Description
[31:0]	0	NA	AXI-ST C2H Completion Data [255:224]

C2H_COMPLETION_SIZE (0x050)**Table: C2H_COMPLETION_SIZE (0x050)**

Bit	Default	Access Type	Description
[31:13]	0	NA	Reserved
[12]	0	RW	Completion Type. 1'b1: NO_PLD_BUT_WAIT 1'b0: HAS PLD
[10:8]	0	RW	s_axis_c2h_cmpt_ctrl_err_idx[2:0] Completion Error Bit Index. 3'b000: Selects 0th register. 3'b111: No error bit is reported.
[6:4]	0	RW	s_axis_c2h_cmpt_ctrl_col_idx[2:0] Completion Color Bit Index. 3'b000: Selects 0th register. 3'b111: No color bit is reported.
[3]	0	RW	s_axis_c2h_cmpt_ctrl_user_trig Completion user trigger
[1:0]	0	RW	AXI4-Stream C2H completion data size. 00: 8 Bytes 01: 16 Bytes 10: 32 Bytes 11: 64 Bytes

Related Information

[AXI4-Stream C2H Completion Ports](#)

SCRATCH_REG0 (0x060)

Table: SCRATCH_REG0 (0x060)

Bit	Default	Access Type	Description
[31:0]	0	RW	Scratch register

SCRATCH_REG1 (0x064)**Table: SCRATCH_REG1 (0x064)**

Bit	Default	Access Type	Description
[31:0]	0	RW	Scratch register

C2H_PACKETS_DROP (0x088)**Table: C2H_PACKETS_DROP (0x088)**

Bit	Default	Access Type	Description
[31:0]	0	R	The number of AXI-ST C2H packets (descriptors) dropped per transfer

Each AXI-ST C2H transfer can contain one or more descriptors depending on transfer size and C2H buffer size. This register represents how many of the descriptors were dropped in the current transfer. This register will reset to 0 in the beginning of each transfer.

C2H_PACKETS_ACCEPTED (0x08C)**Table: C2H_PACKETS_ACCEPTED (0x08C)**

Bit	Default	Access Type	Description
[31:0]	0	R	The number of AXI-ST C2H packets (descriptors) accepted per transfer

Each AXI-ST C2H transfer can contain one or more descriptors depending on the transfer size and C2H buffer size. This register represents how many of the descriptors were accepted in the current transfer. This register will reset to 0 at the beginning of each transfer.

DESCRIPTOR_BYPASS (0x090)**Table: Descriptor Bypass (0x090)**

Bit	Default	Access Type	Field	Description
[31:3]	0	NA		Reserved
[2:1]	0	RW	c2h_dsc_bypass	C2H descriptor bypass loopback. When set, the C2H descriptor bypass-out port

Bit	Default	Access Type	Field	Description
				<p>is looped back to the C2H descriptor bypass-in port.</p> <p>2'b00: No bypass loopback.</p> <p>2'b01: C2H MM desc bypass loopback and C2H Stream cache bypass loopback.</p> <p>2'b10: C2H Stream Simple descriptor bypass loopback.</p> <p>2'b11: H2C stream 64 byte descriptors are looped back to Completion interface.</p>
[0]	0	RW	h2c_dsc_bypass	<p>H2C descriptor bypass loopback. When set, the H2C descriptor bypass-out port is looped back to the H2C descriptor bypass-in port.</p> <p>1'b1: H2C MM and H2C Stream descriptor bypass loopback</p> <p>1'b0: No descriptor loopback</p>

USER_INTERRUPT (0x094)

Table: User Interrupt (0x094)

Bit	Default	Access Type	Field	Description
[31:20]	0	NA		Reserved
[19:12]	0	RW	usr_irq_in_fun	User interrupt function number
[11:9]	0	NA		Reserved
[8:4]	0	RW	usr_irq_in_vec	User interrupt vector number
[3:1]	0	NA		Reserved
[0]	0	RW	usr_irq	User interrupt. When set, the example design generates a user interrupt.

To generate a user interrupt:


1. Write the function number at bits [19:12]. This corresponds to the function that generates the `usr_irq_in_fnc` user interrupt.
2. Write MSI-X Vector number at bits [8:4]. This corresponds to the entry in the MSI-X table that is set up for `usr_irq_in_vec` user interrupt.
3. Write 1 to bit [0] to generate user interrupt. This bit clears itself after `usr_irq_out_ack`

from the DMA is generated.

All three above steps can be done at the same time, with a single write.

The user interrupt timing diagram is shown below:

Figure: Interrupt



USER_INTERRUPT_MASK (0x098)

Table: User Interrupt Mask (0x098)

Bit	Default	Access Type	Description
[31:0]	0	RW	User Interrupt Mask

USER_INTERRUPT_VECTOR (0x09C)

Table: User Interrupt Vector (0x09C)

Bit	Default	Access Type	Description
[31:0]	0	RW	User Interrupt Vector

The user_interrupt_mask[31:0] and user_interrupt_vector[31:0] registers are provided as an example design for user interrupt aggregation that can generate a user interrupt for a function. The user_interrupt_mask[31:0] is anded (bitwise and) with user_interrupt_vector[31:0] and a user interrupt is generated. The user_interrupt_vector[31:0] is clear on read register.

To generate a user interrupt:

1. Write the function number at user_interrupt[19:12]. This corresponds to which function generates the usr_irq_in_fnc user interrupt.
2. Write the MSI-X Vector number at user_interrupt[8:4]. This corresponds to which entry in MSI-X table is set up for the usr_irq_in_vec user interrupt.
3. Write mask value in the user_interrupt_mask[31:0] register.
4. Write the interrupt vector value in the user_interrupt_vector[31:0] register.

This generates a user interrupt to the DMA block.

There are two ways to generate user interrupt:

- Write to user_interrupt[0], or
- Write to the user_interrupt_vector[31:0] register with mask set.

DMA_CONTROL (0x0A0)

Table: DMA Control (0x0A0)

Bit	Default	Access Type	Field	Description
[31:1]		NA		Reserved
[0]	0	RW	gen_qdma_reset	When soft_reset is set, generates a soft reset to the DMA block. This bit is cleared after 100 cycles.

Writing a 1 to DMA_control[0] generates a soft reset on soft_reset_n (active-Low). A reset is asserted for 100 cycles, and following which of the signals will be deasserted.

VDM_MESSAGE_READ (0x0A4)

Table: VDM Message Read (0x0A4)

Bit	Default	Access Type	Description
[31:0]		RO	VDM message read

Vendor Defined Message (VDM) messages, st_rx_msg_data, are stored in FIFO in the example design. A read to this register (0x0A4) will pop out one 32-bit message at a time.

Customizing and Generating the Example Design

In the Customize IP dialog box, use the default core parameter values for the IP example design. After reviewing the core parameters:

1. Right-click the component name.
2. Select Open IP Example Design.

This opens a separate example design.

Test Bench

The PCI Express® Root Port Model is a robust test bench environment that provides a test program interface that can be used with the provided Programmed Input/Output (PIO) design or with your design. The purpose of the Root Port Model is to provide a source mechanism for generating downstream PCI® Express TLP traffic to stimulate the customer design, and a destination mechanism for receiving upstream PCI® Express TLP traffic from the customer

design in a simulation environment.


Source code for the Root Port Model is included to provide the model for a starting point for your test bench. All the significant work for initializing the core configuration space, creating TLP transactions, generating TLP logs, and providing an interface for creating and verifying tests is complete. This allows you to focus on verifying the functionality of the design rather than spending time developing an Endpoint core test bench infrastructure.

The Root Port Model consists of:

- Test Programming Interface (TPI), which allows you to stimulate the Endpoint device for the PCI Express.
- Example tests that illustrate how to use the test program TPI.
- Verilog source code for all Root Port Model components, which allow you to customize the test bench.

The following figure illustrates the Root Port Model coupled with the PIO design.

Figure: Root Port Model and Top-Level Endpoint



X12468

Architecture

The Root Port Model consists of these blocks:

- dsport (Root Port)
- usrapp_tx
- usrapp_rx
- usrapp_com (Verilog only)

The usrapp_tx and usrapp_rx blocks interface with the dsport block for transmission and reception of TLPs to/from the Endpoint Design Under Test (DUT). The Endpoint DUT consists of the Endpoint for PCIe and the PIO design (displayed) or customer design.

The usrapp_tx block sends TLPs to the dsport block for transmission across the PCI Express Link to the Endpoint DUT. In turn, the Endpoint DUT device transmits TLPs across the PCI Express Link to the dsport block, which are subsequently passed to the usrapp_rx block. The dsport and core are responsible for the data link layer and physical link layer processing when communicating across the PCI Express logic. Both usrapp_tx and usrapp_rx use the usrapp_com block for shared functions, for example, TLP processing and log file outputting. Transaction sequences or test programs are initiated by the usrapp_tx block to stimulate the Endpoint device fabric interface. TLP responses from the Endpoint device are received by the usrapp_rx block. Communication between the usrapp_tx and usrapp_rx blocks allow the usrapp_tx block to verify correct behavior and act accordingly when the usrapp_rx block has received TLPs from the Endpoint device.

Scaled Simulation Timeouts

The simulation model of the core uses scaled-down times during link training to allow for the link to train in a reasonable amount of time during simulation. According to the *PCI Express Specification, rev. 3.0* (<http://www.pcisig.com/specifications>), there are various timeouts associated with the link training and status state machine (LTSSM) states. The core scales these timeouts by a factor of 256 during simulation, except in the Recovery Speed_1 LTSSM state, where the timeouts are not scaled.

Test Selection

All simulation test cases are based on the provided example designs. Simulation tasks perform writes and reads to specific example design registers for setup and checking. These simulation tasks might not work for all customer designs.

Available Tests

The following table describes the tests provided for simulation. These tests are selected based on the QDMA IP configuration. For example, if the AXI4-MM only option is selected, the qdma_mm_test0 test case is selected and will be executed during simulation.

Table: Test Case

Option	Test Name	Language	Description
--------	-----------	----------	-------------

Option	Test Name	Language	Description
AXI4-MM only	qdma_mm_test0	Verilog	<p>1. The test bench initializes the queue and performs the AXI4-MM transfer in the H2C direction.</p> <p>2. Then, the test bench initializes the queue and performs the AXI4-MM transfer in the C2H direction.</p> <p>The test bench compares the write data with the read data for correctness.</p>
AXI4-ST only	qdma_st_test0	Verilog	<p>1. The test bench initializes the queue, performs the AXI4-ST transfer in H2C direction, and then checks for data correctness.</p> <p>2. The test bench initializes the queue, performs the AXI4-ST transfer in the C2H direction, and then checks data for correctness.</p>
AXI4-MM and AXI4-ST with completion	qdma_mm_st_test0	Verilog	<p>1. The test bench initializes the queue and performs the AXI4-MM transfer in the H2C direction. Then, the test bench initializes the queue, performs AXI4-MM in C2H direction, and compares the data for correctness.</p> <p>2. The test bench initializes the queue, performs the AXI4-ST transfer in the H2C direction, and then checks data for correctness.</p> <p>3. Then, the test bench initializes the and performs AXI4-ST in the C2H direction and check data for correctness.</p> <p>This test is a combination of test cases qdma_mm_test0 and qdma_st_test0.</p>

Verilog Test Selection

You can change the test case by editing the `usr_pci_exp_usrapp_tx.v` file. Look for the "testname" assignment and modify the test case.

Waveform Dumping

For information on simulator waveform dumping, see the *Vivado Design Suite User Guide: Logic Simulation* ([UG900](#)).

Verilog Flow

The model provides a mechanism for outputting the simulation waveform to a file using the `+dump_all` command line parameter.

Output Logging

The simulation process creates three output log files. They are tx.dat, rx.dat, and error.dat. The rx.dat and tx.dat files each contain a detailed record of every TLP that was received and transmitted, respectively.

★ **Tip:** With an understanding of the expected TLP transmission during a specific test case, you can isolate a failure.

The error.dat file is used in conjunction with the expectation tasks. Only PCIe protocol errors will be listed in error.dat. DMA data mismatch or transfer errors will be printed out in the log file.

Test Description

The model provides a Test Program Interface (TPI). The TPI provides the means to create tests by invoking a series of Verilog tasks. All tests should follow these steps:

1. Perform conditional comparison of a unique test name.
2. Wait for reset and link-up.
3. Initialize the queue context for that queue.
4. Transmit packet for the queue.
5. Verify that the test succeeded.

Model Task List

PCIe-Related Tasks

For all the PCIe-related tasks, refer to *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG213](#)).

DMA Tasks

Table: DMA Tasks ¹

Name	Input(s)	Output	Description
------	----------	--------	-------------

Name	Input(s)	Output	Description
TSK_QDMA_MM_H2C_TEST	dsc_bypass irq_en	10:0 - -	This task will do a H2C AXI-MM transfer on Queue ID "qid". if "dsc_bypass" is given task will do descriptor bypass If irq_en is given, task will setup MSI-X table and it will send interrupt once transfer is completed.
TSK_QDMA_MM_C2H_TEST	dsc_bypass irq_en	10:0 - -	This task will do a C2H AXI-MM transfer on Queue ID "qid". if "dsc_bypass" is given task will do descriptor bypass If irq_en is given, task will setup MSI-X table and it will send interrupt once transfer is completed.
TSK_QDMA_ST_C2H_TEST	qdID, dsc_bypass	10:0 -	This task will do a C2H AXI-ST transfer on Queue ID "qid". if "dsc_bypass" is given task will do descriptor bypass
TSK_QDMA_ST_H2C_TEST	qdID, dsc_bypass	10:0 -	This task will do a H2C AXI-ST transfer on Queue ID "qid". if "dsc_bypass" is given task will do descriptor bypass
1. The DMA tasks in this table are recommended for testing purposes only.			

Upgrading

Changes from v3.0 to v5.1

For a list of changes in the QDMA Subsystem for PCIe from v3.0 to v5.1, see AR [75234](#).

Comparing With DMA/Bridge Subsystem for PCI Express

The table below describes the differences between the DMA/Bridge Subsystem for PCI Express® and QDMA Subsystem for PCI Express.

Table: Comparing Subsystems

	DMA/Bridge Subsystem	QDMA Subsystem
Configuration	Up to Gen3x16.	Up to Gen3x16.
Channels/Queues	Four Host to Card (H2C) channels, and four Card to Host (C2H) channels with one PF.	Up to 2K queues (All can be assigned to one PF or distributed amongst all four).
SR-IOV	Not Supported.	Supported (four PFs, and 252 VFs).
User Interface	Configured for AXI4 Memory Mapped or AXI4-Stream, but not both.	Each queue will have a context which will determine whether it goes to a AXI4 Memory Mapped or AXI4-Stream.
User Interrupts	Up to 16 user interrupts.	Interrupt aggregation per function.
Device Support	Supported for 7 Series Gen2 to AMD UltraScale+™ devices.	Only supported for UltraScale+ devices.
Interrupts	Legacy, MSI, MSI-X supported.	MSI-X Supported.
Driver Support	Linux, Windows Example Drivers.	Linux, DPDK, Windows.

GT Locations

For more information on GT Locations, see GT Locations appendix in *UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG213](#)).

Debugging

This appendix includes details about resources available on the Support website and debugging tools.

Finding Help with AMD Adaptive Computing Solutions

To help in the design and debug process when using the subsystem, the [Support web page](#) contains key resources such as product documentation, release notes, answer records, information about known issues, and links for obtaining further product support. The [Community Forums](#) are also available where members can learn, participate, share, and ask questions about AMD Adaptive Computing solutions.

Documentation

This product guide is the main document associated with the subsystem. This guide, along with documentation related to all products that aid in the design process, can be found on the [AMD Adaptive Support web page](#) or by using the AMD Adaptive Computing Documentation Navigator. Download the Documentation Navigator from the [Downloads page](#). For more information about this tool and the features available, open the online help after installation.

Debug Guide

For more information on PCIe debug, see [PCIe Debug K-Map](#).

Answer Records

Answer Records include information about commonly encountered problems, helpful information on how to resolve these problems, and any known issues with an AMD Adaptive Computing product. Answer Records are created and maintained daily to ensure that users have access to the most accurate information available.

Answer Records for this subsystem can be located by using the Search Support box on the main [AMD Adaptive Support web page](#). To maximize your search results, use keywords such as:

- Product name
- Tool message(s)
- Summary of the issue encountered

A filter search is available after results are returned to further target the results.

Master Answer Record for the Subsystem

AR [70927](#).

Technical Support

AMD Adaptive Computing provides technical support on the [Community Forums](#) for this AMD LogiCORE™ IP product when used as described in the product documentation. AMD Adaptive Computing cannot guarantee timing, functionality, or support if you do any of the following:

- Implement the solution in devices that are not defined in the documentation.
- Customize the solution beyond that allowed in the product documentation.
- Change any section of the design labeled DO NOT MODIFY.

To ask questions, navigate to the [Community Forums](#).

Debug Tools

There are many tools available to address QDMA Subsystem for PCIe design issues. It is important to know which tools are useful for debugging various situations.

Vivado Design Suite Debug Feature

The AMD Vivado™ Design Suite debug feature inserts logic analyzer and virtual I/O cores directly into your design. The debug feature also allows you to set trigger conditions to capture application and integrated block port signals in hardware. Captured signals can then be analyzed. This feature in the Vivado IDE is used for logic debugging and validation of a design running in AMD devices.

The Vivado logic analyzer is used to interact with the logic debug LogiCORE IP cores, including:

- ILA 2.0 (and later versions)
- VIO 2.0 (and later versions)

See the *Vivado Design Suite User Guide: Programming and Debugging* ([UG908](#)).

Hardware Debug

Hardware issues can range from link bring-up to problems seen after hours of testing. This section provides debug steps for common issues. The AMD Vivado™ debug feature is a valuable resource to use in hardware debug. The signal names mentioned in the following individual sections can be probed using the debug feature for debugging the specific problems.

General Checks


Ensure that all the timing constraints for the core were properly incorporated from the example design and that all constraints were met during implementation.

- Does it work in post-place and route timing simulation? If problems are seen in hardware but not in timing simulation, this could indicate a PCB issue. Ensure that all clock sources are active and clean.
- If using MMCMs in the design, ensure that all MMCMs have obtained lock by monitoring the locked port.
- If your outputs go to 0, check your licensing.

Application Software Development

Device Drivers

Figure: Device Drivers



X28798-110723


The above figure shows the usage model of Linux QDMA software drivers. The QDMA Subsystem for PCIe example design is implemented on an AMD FPGA , which is connected to an X86 host through PCI Express® .

- In the first use mode, the QDMA driver in kernel space runs on Linux, whereas the test application runs in user space.
- In the second use mode, the Data Plane Dev Kit (DPDK) is used to develop a QDMA Poll Mode Driver (PMD) running entirely in the user space, and use the UIO and VFIO kernel framework to communicate with the FPGA .

For device driver documentation click [DMA IP Drivers](#).

Linux QDMA Software Architecture (PF/VF)

Figure: Linux DMA Software Architecture



X20598-022421

The QDMA driver consists of the following three major components:

Device control tool

Creates a netlink socket for PCIe device query, queue management, reading the context of a queue, etc.

DMA tool

Is the user space application to initiate a DMA transaction. You can use standard Linux utility dd or fio, or use the example application in the driver package.

Kernel space driver

Creates the descriptors and translates the user space function into low-level command to interact with the device .

For AMD QDMA Linux driver documentation, click [here](#).

Using the Drivers

Drivers and the corresponding documentation are available at [DMA IP Drivers](#).

Note: Starting from 2022.1 release of the Linux driver for QDMA, if a design is using streaming queues, they must be explicitly enabled through API as they are not configured at module load. If a design is using tandem PCIe methodology at power-on, the enablement must occur after Stage 2 is loaded to the device.


!! Important: 8 MSI-X vectors are needed on all functions (PF/VF) for using QDMA IP driver.

Reference Software Driver Flow

AXI4 Memory Map Flow Chart


Figure: AXI4 Memory Map Flow Chart

Driver App




AXI4 Memory Mapped C2H Flow

Figure: AXI4 Memory Mapped Card to Host (C2H) Flow Diagram




AXI4 Memory Mapped H2C Flow

Figure: AXI4 Memory Mapped Host to Card (H2C) Flow Diagram



AXI4-Stream Flow Chart


Figure: AXI4-Stream Flow Chart



X20551-041521

AXI4-Stream C2H Flow


Figure: AXI4-Stream C2H Flow Diagram



X20527-041619

AXI4-Stream H2C Flow

Figure: AXI4-Stream H2C Flow Diagram



Additional Resources and Legal Notices

Finding Additional Documentation

Technical Information Portal

The AMD Technical Information Portal is an online tool that provides robust search and navigation for documentation using your web browser. To access the Technical Information Portal, go to <https://docs.amd.com>.

Documentation Navigator

Documentation Navigator (DocNav) is an installed tool that provides access to AMD Adaptive Computing documents, videos, and support resources, which you can filter and search to find information. To open DocNav:

- From the AMD Vivado™ IDE, select Help > Documentation and Tutorials.
- On Windows, click the Start button and select Xilinx Design Tools > DocNav.
- At the Linux command prompt, enter docnav.

 **Note:** For more information on DocNav, refer to the *Documentation Navigator User Guide* ([UG968](#)).

Design Hubs

AMD Design Hubs provide links to documentation organized by design tasks and other topics, which you can use to learn key concepts and address frequently asked questions. To access the Design Hubs:

- In DocNav, click the Design Hubs View tab.
- Go to the [Design Hubs](#) web page.

Support Resources

For support resources such as Answers, Documentation, Downloads, and Forums, see [Support](#).

References

These documents provide supplemental material useful with this product guide:

1. AMBA AXI4-Stream Protocol Specification ([ARM IHI 0051A](#))
2. PCI-SIG Specifications (www.pcisig.com/specifications)
3. Virtex 7 FPGA Integrated Block for PCI Express LogiCORE IP Product Guide ([PG023](#))
4. 7 Series FPGAs Integrated Block for PCI Express LogiCORE IP Product Guide ([PG054](#))
5. UltraScale Devices Gen3 Integrated Block for PCI Express LogiCORE IP Product Guide ([PG156](#))
6. AXI Bridge for PCI Express Gen3 Subsystem Product Guide ([PG194](#))
7. DMA/Bridge Subsystem for PCI Express Product Guide ([PG195](#))
8. In-System IBERT LogiCORE IP Product Guide ([PG246](#))
9. UltraScale+ Devices Integrated Block for PCI Express LogiCORE IP Product Guide ([PG213](#))
10. Vivado Design Suite: AXI Reference Guide ([UG1037](#))
11. Vivado Design Suite User Guide: Designing IP Subsystems using IP Integrator ([UG994](#))
12. UltraScale Architecture GTY Transceivers User Guide ([UG578](#))
13. UltraScale Architecture GTH Transceivers User Guide ([UG576](#))
14. Vivado Design Suite User Guide: Designing with IP ([UG896](#))

15. Vivado Design Suite User Guide: Getting Started ([UG910](#))
16. Vivado Design Suite User Guide: Logic Simulation ([UG900](#))
17. Vivado Design Suite User Guide: Using Constraints ([UG903](#))
18. Vivado Design Suite User Guide: Programming and Debugging ([UG908](#))
19. Vivado Design Suite User Guide: Dynamic Function eXchange ([UG909](#))
20. Vivado Design Suite User Guide: Release Notes, Installation, and Licensing ([UG973](#))
21. PIPE Mode Simulation Using Integrated Endpoint PCI Express Block in Gen2 x8 and Gen3 x8 Configurations ([XAPP1184](#))

Revision History

The following table shows the revision history for this document.

Section	Revision Summary
11/20/2025 Version 5.1	
Minimum Device Requirements	Updated section.
Supported Devices	Updated section.
05/29/2025 Version 5.1	
AXI4-Stream C2H Ports	Updated AXI4-Stream C2H Port Descriptions table.
AXI Memory Mapped Descriptor for H2C and C2H (32B)	Updated AXI Memory Mapped Descriptor Structure for H2C and C2H table.
QDMA_DMAP_SEL_CMPT_CIDX[2048] (0x1800C)	Updated QDMA_DMAP_SEL_CMPT_CIDX[2048] (0x1800C) table.
Supported Devices	Updated table.
12/18/2024 Version 5.0	
IP Facts	Added support for Spartan UltraScale+.
Minimum Device Requirements	Updated section.
Supported Devices	Updated table.
Traffic Manager Output Interface	Updated section.
Function Level Reset	Updated section.
05/30/2024 Version 5.0	

Section	Revision Summary
Features	Added support for 64-bit AXI width.
Performance and Resource Utilization	Updated section.
Minimum Device Requirements	Updated section.
BDF Table	Updated slave bridge examples.
Address Translation	Updated examples.
User Interrupts	Updated section.
Queue Status Ports	Added new section.
AXI Memory Mapped Descriptor for H2C and C2H (32B)	Descriptor length reduced from 28 bits to 16 bits.
10/18/2023 Version 5.0	
Performance and Resource Utilization	Updated.
Tandem Configuration	Updated.
Limitations	Updated.
RTL Version Register (0x22414)	Updated.
05/10/2023 Version 5.0	
Performance and Resource Utilization	Added QDMA performance registers table.
Interrupts	Updated.
Legacy Interrupt	Added a note.
AXI Bridge Slave Ports	Added s_axib_wstrb signal description.
User Interrupt	Updated.
Supported Devices	Updated.
Changes from v3.0 to v5.1	Updated.
10/27/2022 Version 5.0	
Features	64-bit AXI width is no longer supported.
FLR Ports	Removed usr_flr_clr port.
GT Locations	Added new appendix.

Section	Revision Summary
AXI Stream Performance Example Design	Added new section.
Tandem Configuration	Updated.
Function Map Table	Updated.
Context Programming	Updated.
05/20/2022 Version 4.0	
Slave Bridge	Updated BDF table.
AXI4-Lite Master Ports	Updated Config AXI4-Lite Memory Mapped Read Master Interface Port Descriptions table.
AXI4-Stream C2H Ports	Updated AXI4-Stream C2H Port Descriptions table.
Tandem Configuration	Updated for new device support.
Debug Guide	Added new section.
01/05/2022 Version 4.0	
General updates	Updated PDF keywords.
12/17/2021 Version 4.0	
QDMA Global Ports	Added csr_prog_done.
Queue Status Ports	Updated marker_cookie port.
04/15/2021 Version 4.0	
Slave Bridge	Added more details regarding BDF tables and address translation, with examples provided.
Tandem Configuration, and Basic Tab	Added information about for Dynamic Function eXchange (DFX) over PCIe support.
07/01/2020 Version 4.0	
C2H Stream Packet Type	Updated Marker response for QDMA 4.0 (from Queues Status ports rather than descriptor bypass out ports).
Host Profile	Added a new Host Profile Context table that needs to be programmed.

Section	Revision Summary
Register Space	<p>Updated the register CSV files.</p> <p>Updated the register address.</p> <p>Added tip to expose all debug registers.</p>
06/10/2020 Version 4.0	
Register Space	<p>Reorganized section. Some register were updated.</p>
QDMA_CSR (0x0000) and Bridge Register Space	<p>Moved register descriptions to CSV file external to product guide.</p>
Descriptor Context and Completion Context Structure	<p>Updated some context tables.</p>
Context Programming	<p>Added a new Host Profile Context table that needs to be programmed.</p>
Port Descriptions	<p>Removed ports, and added new ports.</p>
Customizing and Generating the Subsystem	<p>Updated options and descriptions for Vivado 2020.1.</p>
PCIe BARs Tab	<p>Increased QDMA bar size to 256Kbytes in PFs, and 32Kbytes in VFs.</p>
Debug and Additional Options Tab	<p>Added.</p>
Upgrading	<p>Added reference to AR for changes between core versions.</p>
11/22/2019 Version 3.0	
RTL Version Register (0x22414)	<p>Added PF RTL version register in the doc</p>
RTL Version Register (0x5014)	<p>Added VF RTL version register in the doc</p>
AXI4-Stream Status Ports	<p>Added the axis_c2h_status_error port. This port will be available starting in a 2019.2 patch release.</p>
QDMA C2H Descriptor Bypass Output Marker Response Descriptions table	<p>Added C2H Stream marker_cookie field for marker response. This feature will be available starting in a 2019.2 patch release.</p>
QDNA_GLBL2_MISC_CAP (0x134)	<p>Updated available bits and descriptions.</p>

Section	Revision Summary
VDM	Added information regarding back-to-back VDM access not being supported.
05/22/2019 Version 3.0	
Performance and Resource Utilization	Added performance details, and Performance Report answer record.
Minimum Device Requirements	Enabled Gen4 devices for QDMA.
User Parameters	Added link to AR for additional core customization options.
Capabilities Tab	Mailbox can be selected independently of SR-IOV selection.
AXI Stream Loopback Example Design	New example design added.
12/05/2018 Version 3.0	
IP Facts and Using the Drivers	Added Windows driver support.
Register Space	Added registers, and updated registers.
PCIe MISC Tab and PCIe DMA Tab	Updated for the 2018.3 release.
Example Design	Added two example designs, and updated registers.
Upgrading	Added reference to AR for changes between core versions.
09/04/2018 Version 2.0	
Port Descriptions	For tm_dsc_sts_rdy (VDM Ports) and st_rx_msg_rdy (QDMA Traffic Manager Credit Output Ports), emphasized that when this interface is not used, Ready must be tied-off to 1.
Register Space	Added a register to stall read requests from H2C Stream Engine if the amount of outstanding data exceeds a programmed threshold.
	Added a new C2H Completion interrupt trigger mode that includes user trigger, timer expiration, or count exceeding the threshold

Section	Revision Summary
06/22/2018 Version 2.0	
Overview chapter	Updated content throughout.
Port Descriptions section	Changed some table content, and some reorganization of the content.
Register Space section	Added Memory Map Register Space and AXI4-Lite Slave Register Space section.
Context Structure Definition section, and Queue Entry Structure section	Removed these sections, and moved content into the QDMA Operations section in the Overview chapter.
Design Flow Steps chapter	Updated descriptions for Basic Tab, Capabilities Tab, PCIe BARs Tab, PCIe Misc Tab, and PCIe DMA Tab.
Example Design chapter	Added two new example designs, and added example design registers.
04/17/2018 Version 1.0	
Initial Xilinx release.	

Please Read: Important Legal Notices

The information presented in this document is for informational purposes only and may contain technical inaccuracies, omissions, and typographical errors. The information contained herein is subject to change and may be rendered inaccurate for many reasons, including but not limited to product and roadmap changes, component and motherboard version changes, new model and/or product releases, product differences between differing manufacturers, software changes, BIOS flashes, firmware upgrades, or the like. Any computer system has risks of security vulnerabilities that cannot be completely prevented or mitigated. AMD assumes no obligation to update or otherwise correct or revise this information. However, AMD reserves the right to revise this information and to make changes from time to time to the content hereof without obligation of AMD to notify any person of such revisions or changes. THIS INFORMATION IS PROVIDED "AS IS." AMD MAKES NO REPRESENTATIONS OR WARRANTIES WITH RESPECT TO THE CONTENTS HEREOF AND ASSUMES NO RESPONSIBILITY FOR ANY INACCURACIES, ERRORS, OR OMISSIONS THAT MAY APPEAR IN THIS INFORMATION. AMD SPECIFICALLY DISCLAIMS ANY IMPLIED WARRANTIES OF NON-INFRINGEMENT, MERCHANTABILITY, OR FITNESS FOR ANY PARTICULAR PURPOSE. IN NO EVENT WILL AMD BE LIABLE TO ANY PERSON FOR ANY RELIANCE, DIRECT, INDIRECT, SPECIAL, OR OTHER CONSEQUENTIAL DAMAGES ARISING FROM THE USE OF ANY INFORMATION CONTAINED HEREIN, EVEN IF AMD IS EXPRESSLY ADVISED OF THE POSSIBILITY

OF SUCH DAMAGES.

AUTOMOTIVE APPLICATIONS DISCLAIMER

AUTOMOTIVE PRODUCTS (IDENTIFIED AS "XA" IN THE PART NUMBER) ARE NOT WARRANTED FOR USE IN THE DEPLOYMENT OF AIRBAGS OR FOR USE IN APPLICATIONS THAT AFFECT CONTROL OF A VEHICLE ("SAFETY APPLICATION") UNLESS THERE IS A SAFETY CONCEPT OR REDUNDANCY FEATURE CONSISTENT WITH THE ISO 26262 AUTOMOTIVE SAFETY STANDARD ("SAFETY DESIGN"). CUSTOMER SHALL, PRIOR TO USING OR DISTRIBUTING ANY SYSTEMS THAT INCORPORATE PRODUCTS, THOROUGHLY TEST SUCH SYSTEMS FOR SAFETY PURPOSES. USE OF PRODUCTS IN A SAFETY APPLICATION WITHOUT A SAFETY DESIGN IS FULLY AT THE RISK OF CUSTOMER, SUBJECT ONLY TO APPLICABLE LAWS AND REGULATIONS GOVERNING LIMITATIONS ON PRODUCT LIABILITY.

Copyright

© Copyright 2018-2025 Advanced Micro Devices, Inc. AMD, the AMD Arrow logo, Artix, Kintex, Spartan, UltraScale, UltraScale+, Virtex, Vivado, Zynq, and combinations thereof are trademarks of Advanced Micro Devices, Inc. PCI, PCIe, and PCI Express are trademarks of PCI-SIG and used under license. AMBA, AMBA Designer, Arm, ARM1176JZ-S, CoreSight, Cortex, PrimeCell, Mali, and MPCore are trademarks of Arm Limited in the US and/or elsewhere. Other product names used in this publication are for identification purposes only and may be trademarks of their respective companies.